

HARDWARE SPECULATION

Mahdi Nazm Bojnordi

Assistant Professor

School of Computing

University of Utah

Announcement

- Homework 5
 - ▣ Apparently, the easiest one so far 😊
 - ▣ Do not miss the deadline!

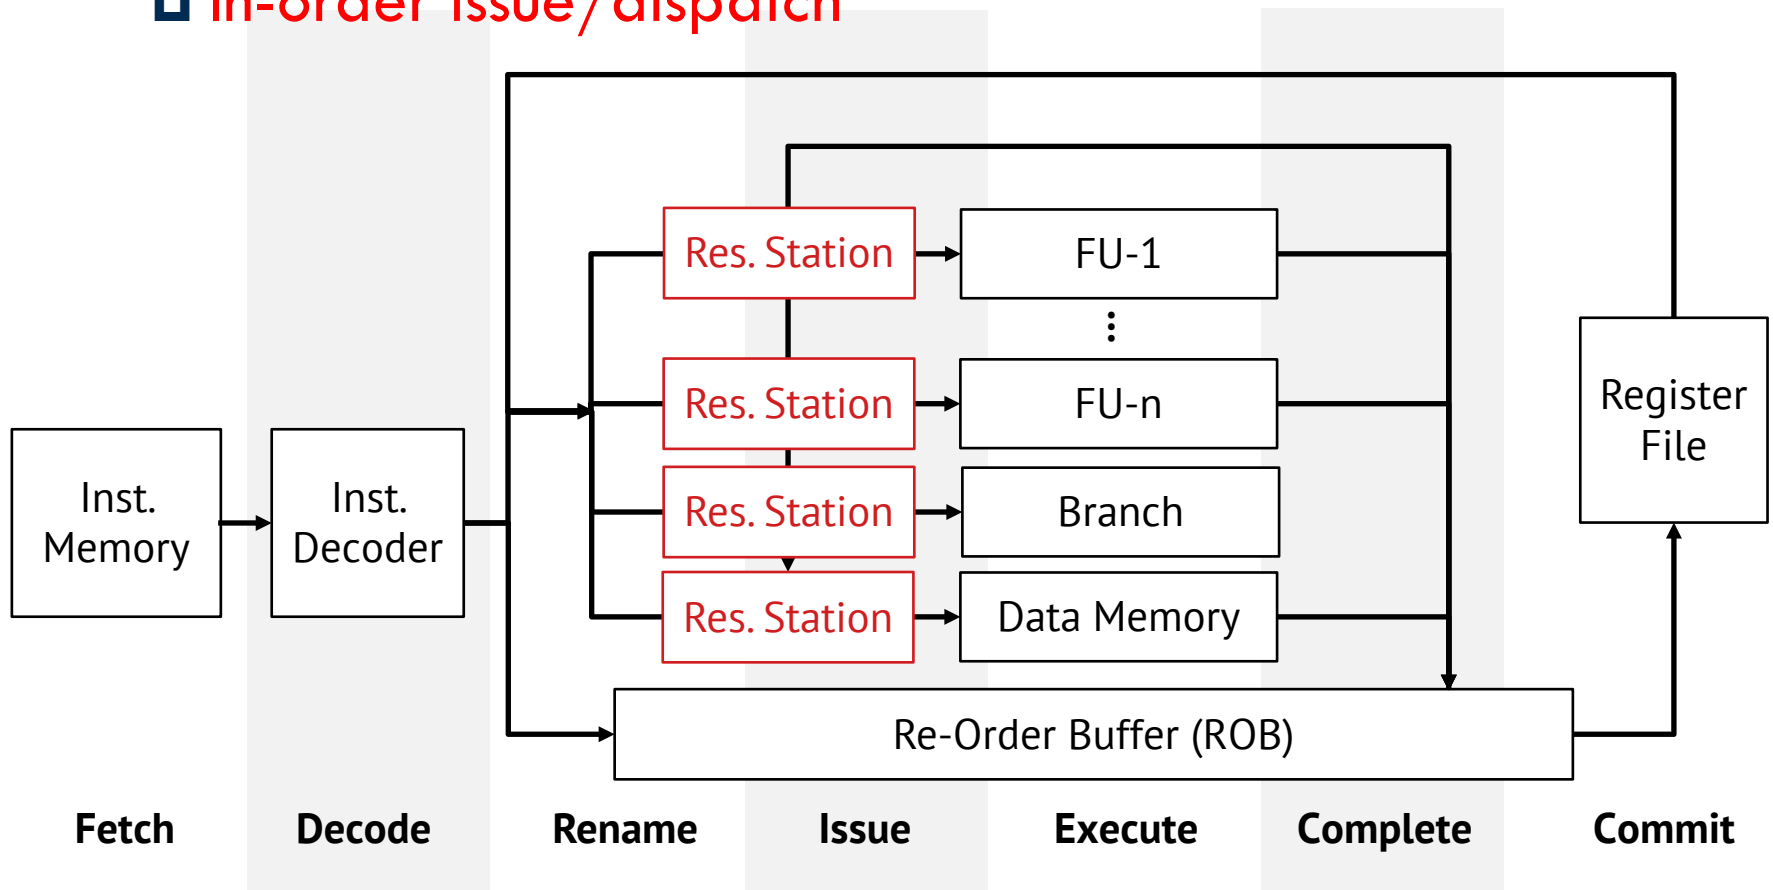
- Homework 6
 - ▣ **** Postponed till after Fall Break ****

Recall: Out-of-Order Execution

- Producer-consumer chains on the fly
 - ▣ Register renaming: remove anti-/output-dependences via register tags
 - ▣ Limited by the number of instructions in the instruction window (ROB)
- Out-of-order issue (dispatch)
 - ▣ Broadcast tags to waiting instructions
 - ▣ Wake up ready instructions and select among them
- Out-of-order execute/complete
- In-order fetch/decode and commit

Out-of-Order Pipelines

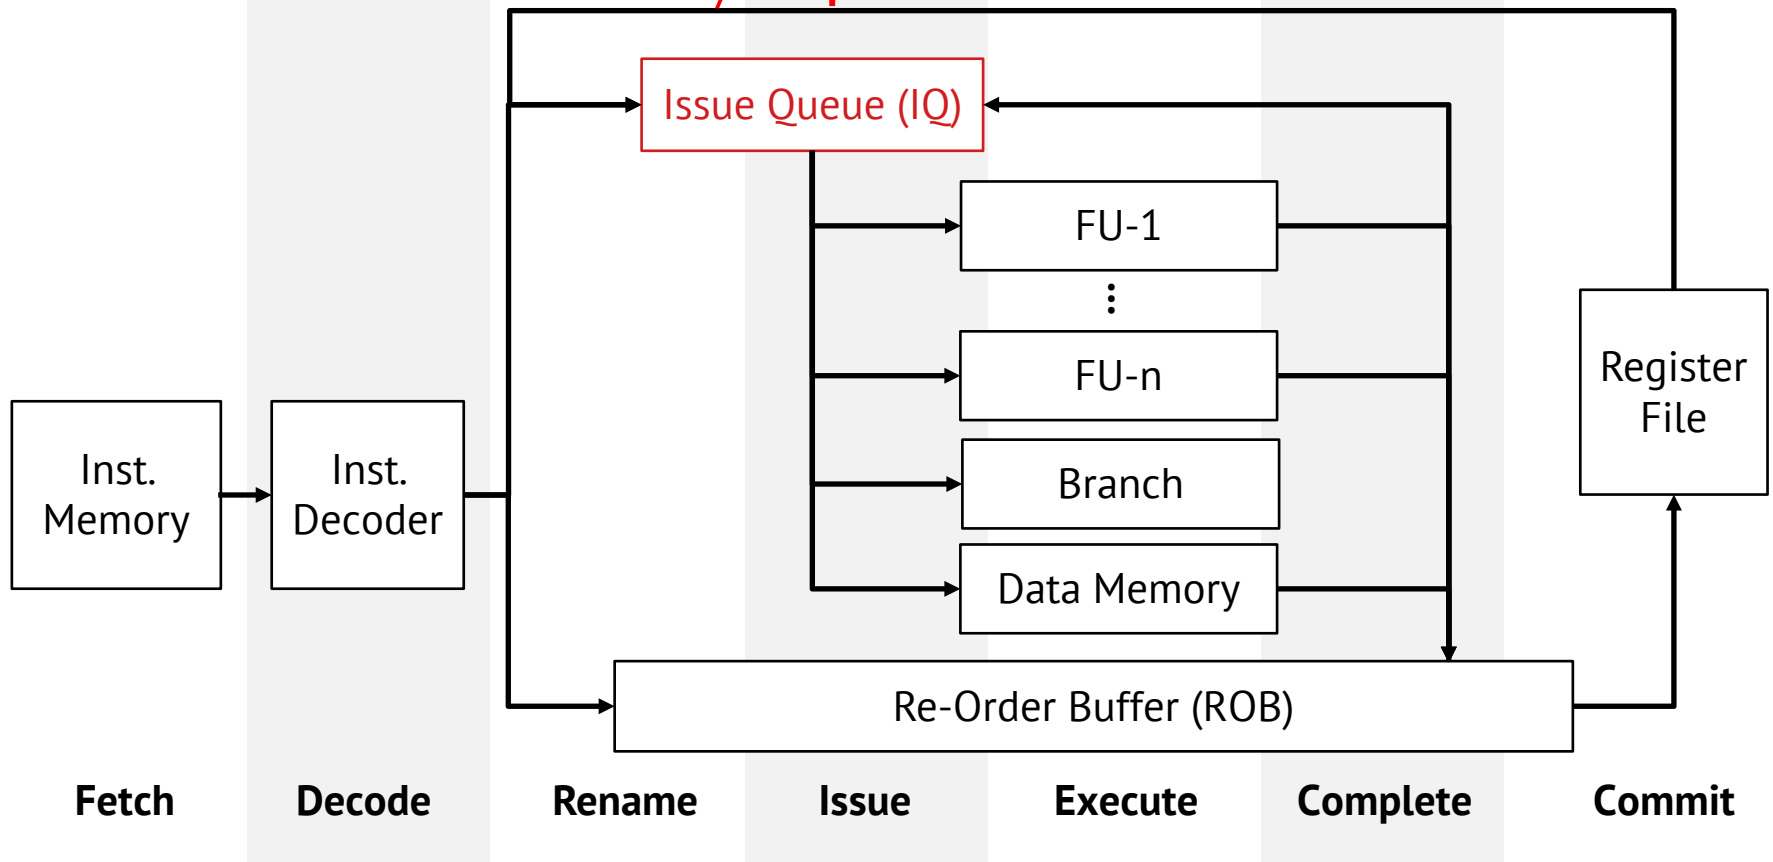
- Distributed reservation stations
 - ▣ In-order issue/dispatch



Out-of-Order Pipelines


- Out of order issue/dispatch to functional units

- Out-of-order issue/dispatch



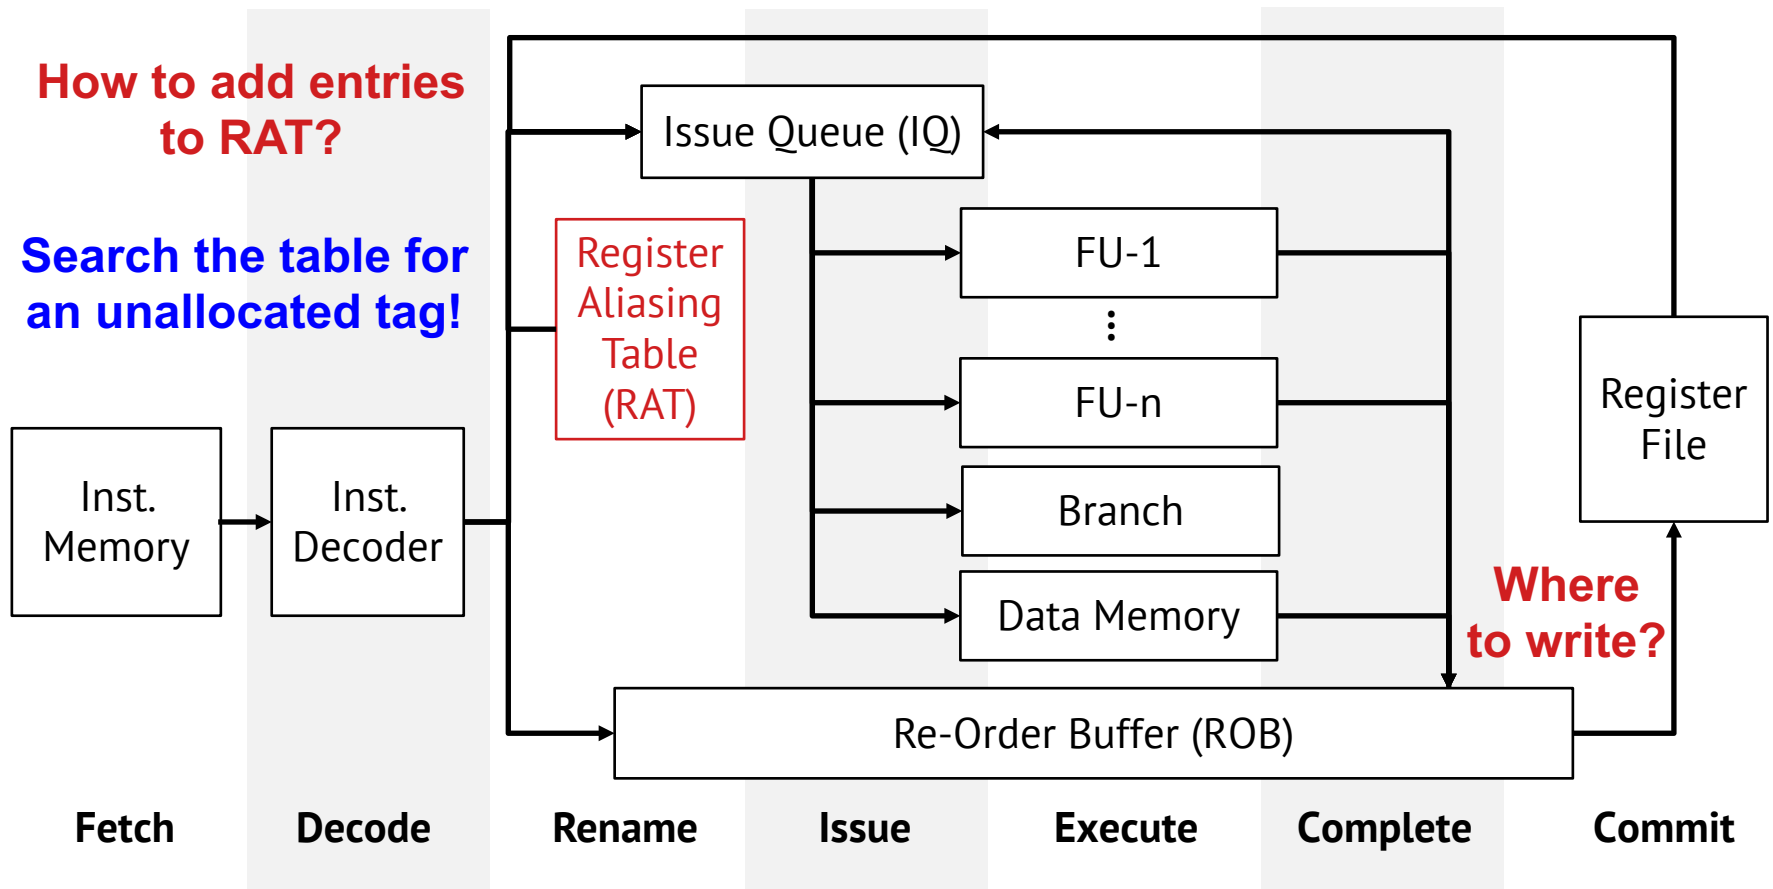
Out-of-Order Issue Queue

- Two step wakeup and select logic



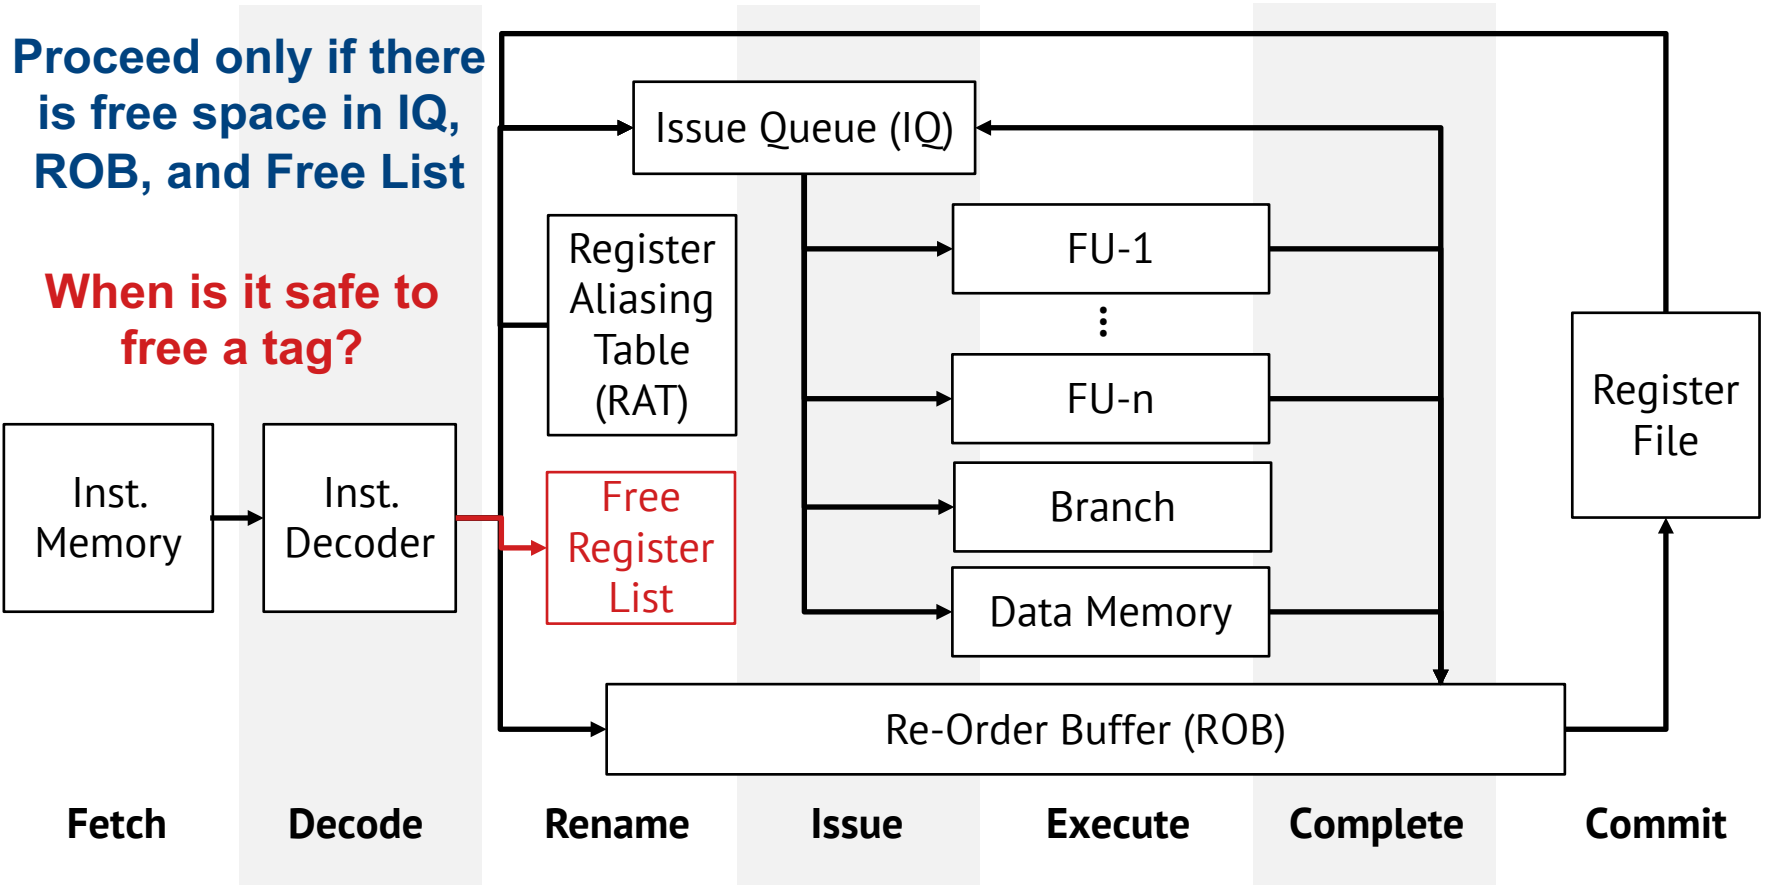
Register Renaming

- Register aliasing table for fast lookup



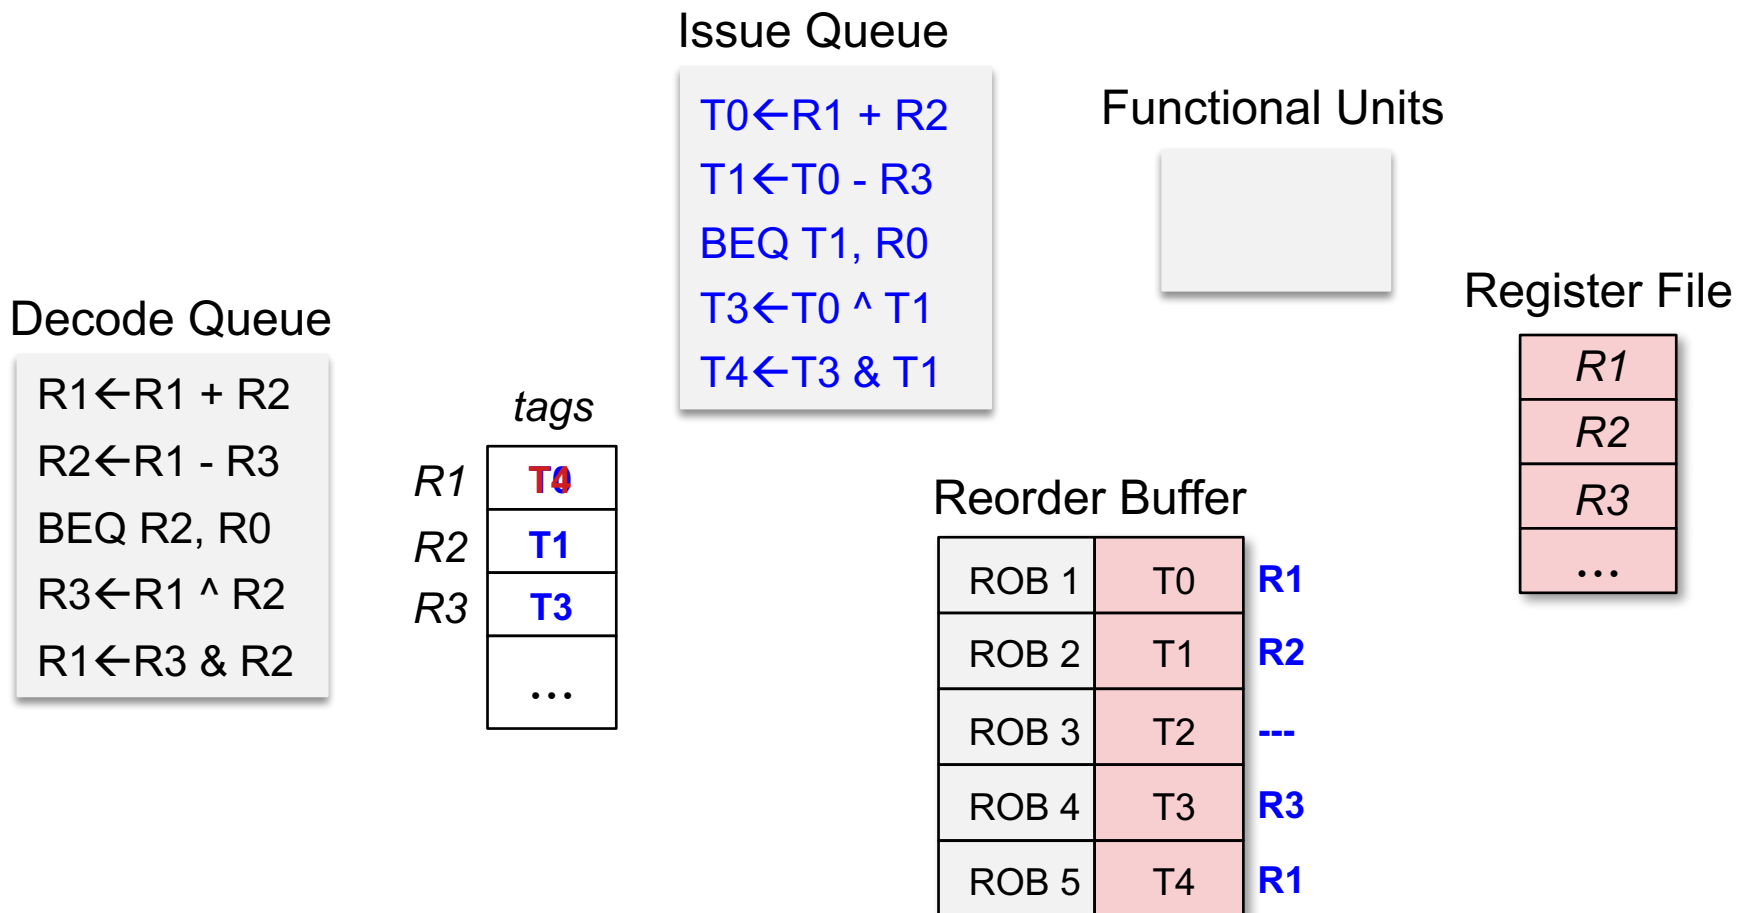
Register Renaming

- Free register list for fast register renaming



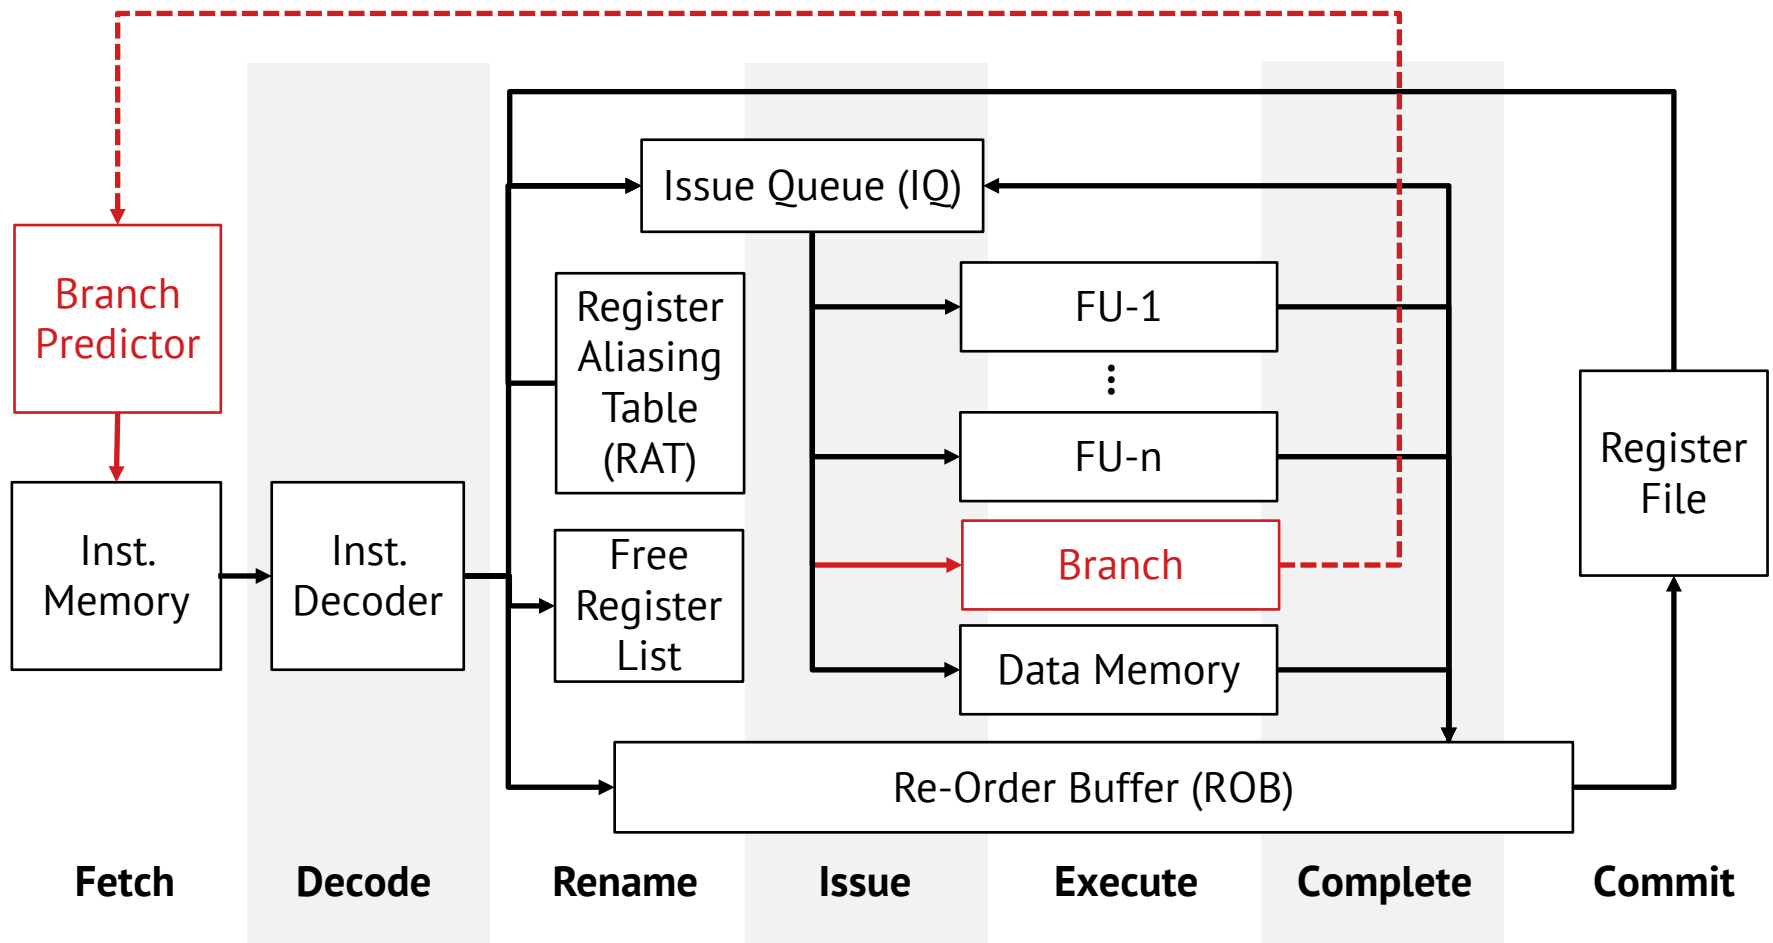
Register Renaming Example

- Where values are stored?



Branch Recovery

□ How to handle branches?



Revisit Branch Prediction

- Problem: find the average number of stall cycles caused by branches in a pipeline, where branch misprediction penalty is 20 cycles, branch predictor accuracy is 90%, and branch target buffer hit rate is 80%. Every fifth instruction is a branch; 30% of branches are actually taken.

Revisit Branch Prediction

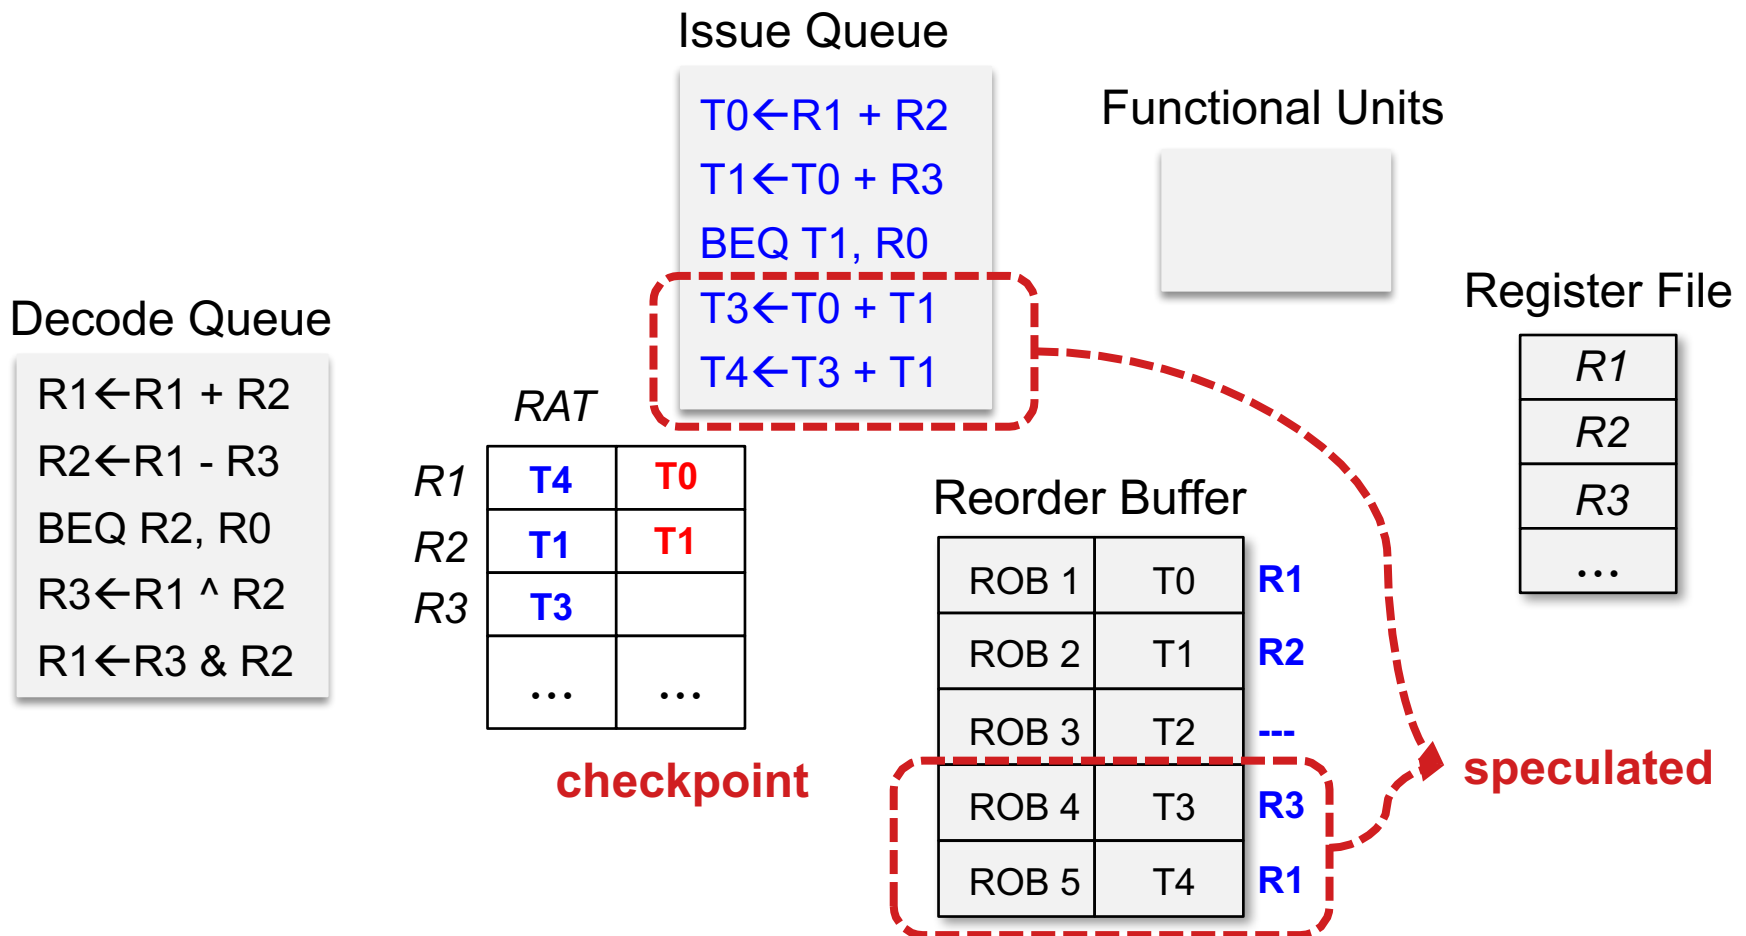
- Problem: find the average number of stall cycles caused by branches in a pipeline, where branch misprediction penalty is 20 cycles, branch predictor accuracy is 90%, and branch target buffer hit rate is 80%. Every fifth instruction is a branch; 30% of branches are actually taken.
- Average misses = $1 - (0.3 \times 0.9 \times 0.8 + 0.7 \times 0.9) = 0.151$
- Average stalls = $20 \times 0.2 \times 0.151 = 0.6$

Speculated Execution

- **Problem:** branch may significantly limit performance
 - ▣ consumer of a load or long latency instructions
- **Solution:** speculative instruction execution
 - ▣ Fetch and decode instructions speculatively
 - ▣ Issue and execute speculative instructions
 - ▣ Branch resolution
 - Nullify the impact of speculative instructions if mispredicted
 - Commit speculative instructions (writes to register file/memory) only if prediction was correct

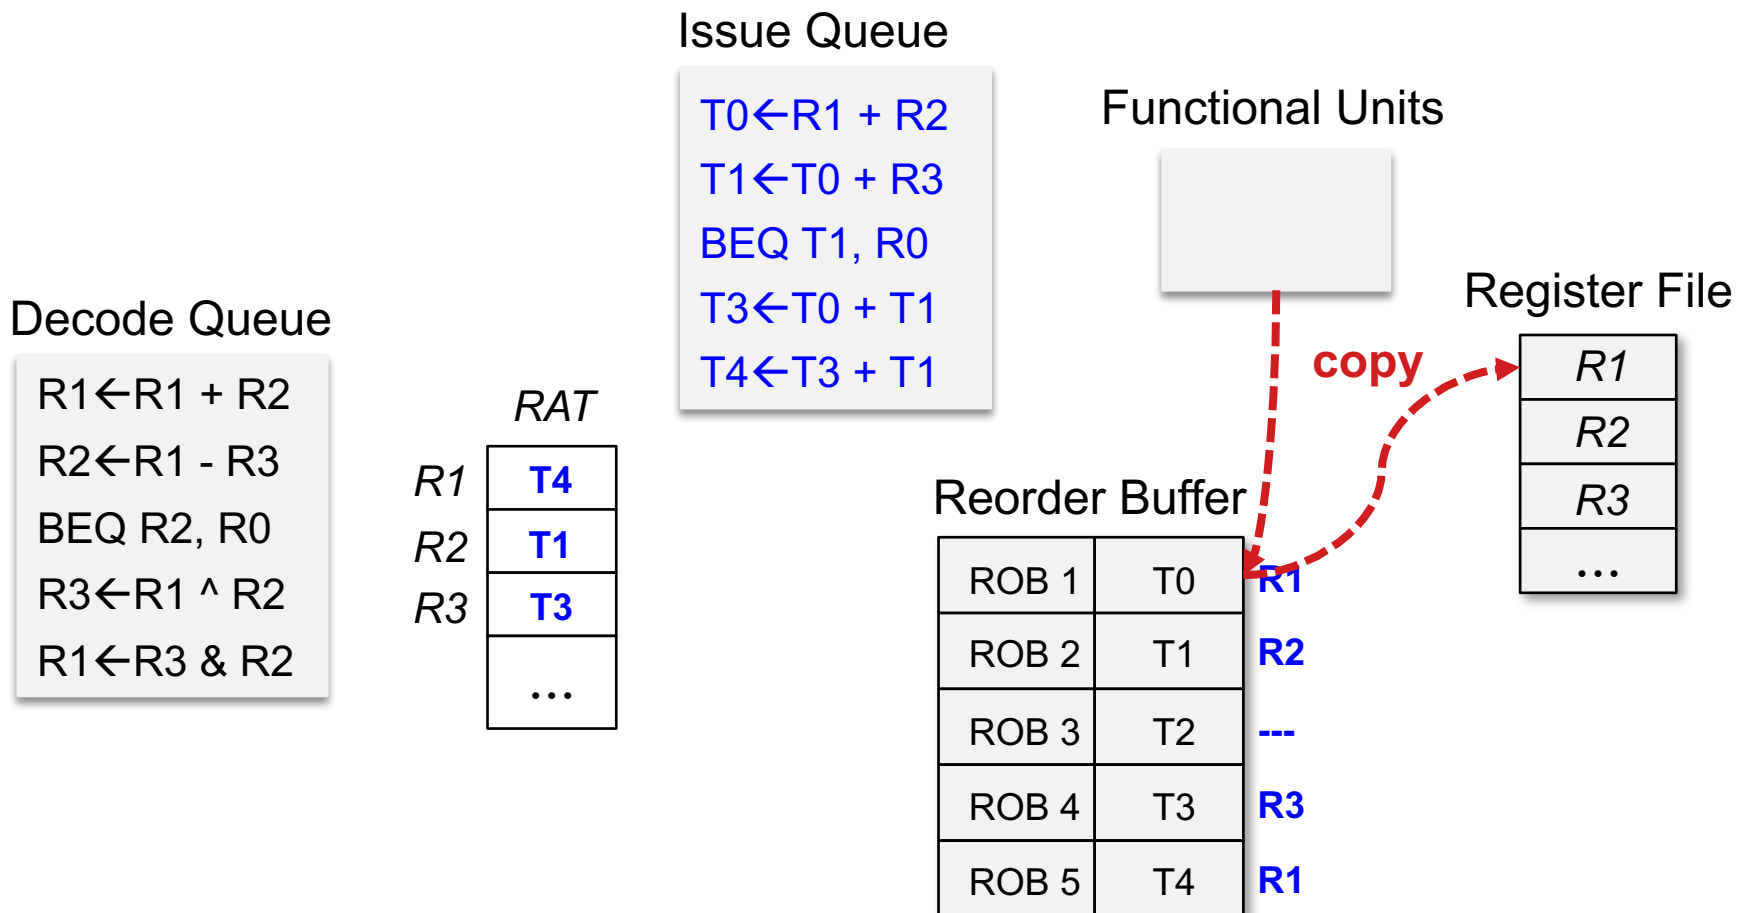
Branch Recovery

- Squash all mispredicted entries



Physical Register File

- Avoid copying register values multiple times



Physical Register File

- Avoid copying register values multiple times

Note1: only a subset of the Phy. Reg. file is committed at any time.

Decode Queue

$R1 \leftarrow R1 + R2$
 $R2 \leftarrow R1 - R3$
 $BEQ\ R2, R0$
 $R3 \leftarrow R1 \wedge R2$
 $R1 \leftarrow R3 \& R2$


Front RAT

R1	P4
R2	P1
R3	P3
	...

Issue Queue

$P0 \leftarrow R1 + R2$
 $P1 \leftarrow P0 + R3$
 $BEQ\ P1, R0$
 $P3 \leftarrow P0 + P1$
 $P4 \leftarrow P3 + P1$

Functional Units



Retire RAT

P0	R1
P1	R2
	R3
...	

Reorder Buffer

ROB 1	R1, P0
ROB 2	R2, P1
ROB 3	---
ROB 4	R3, P3
ROB 5	R1, P4

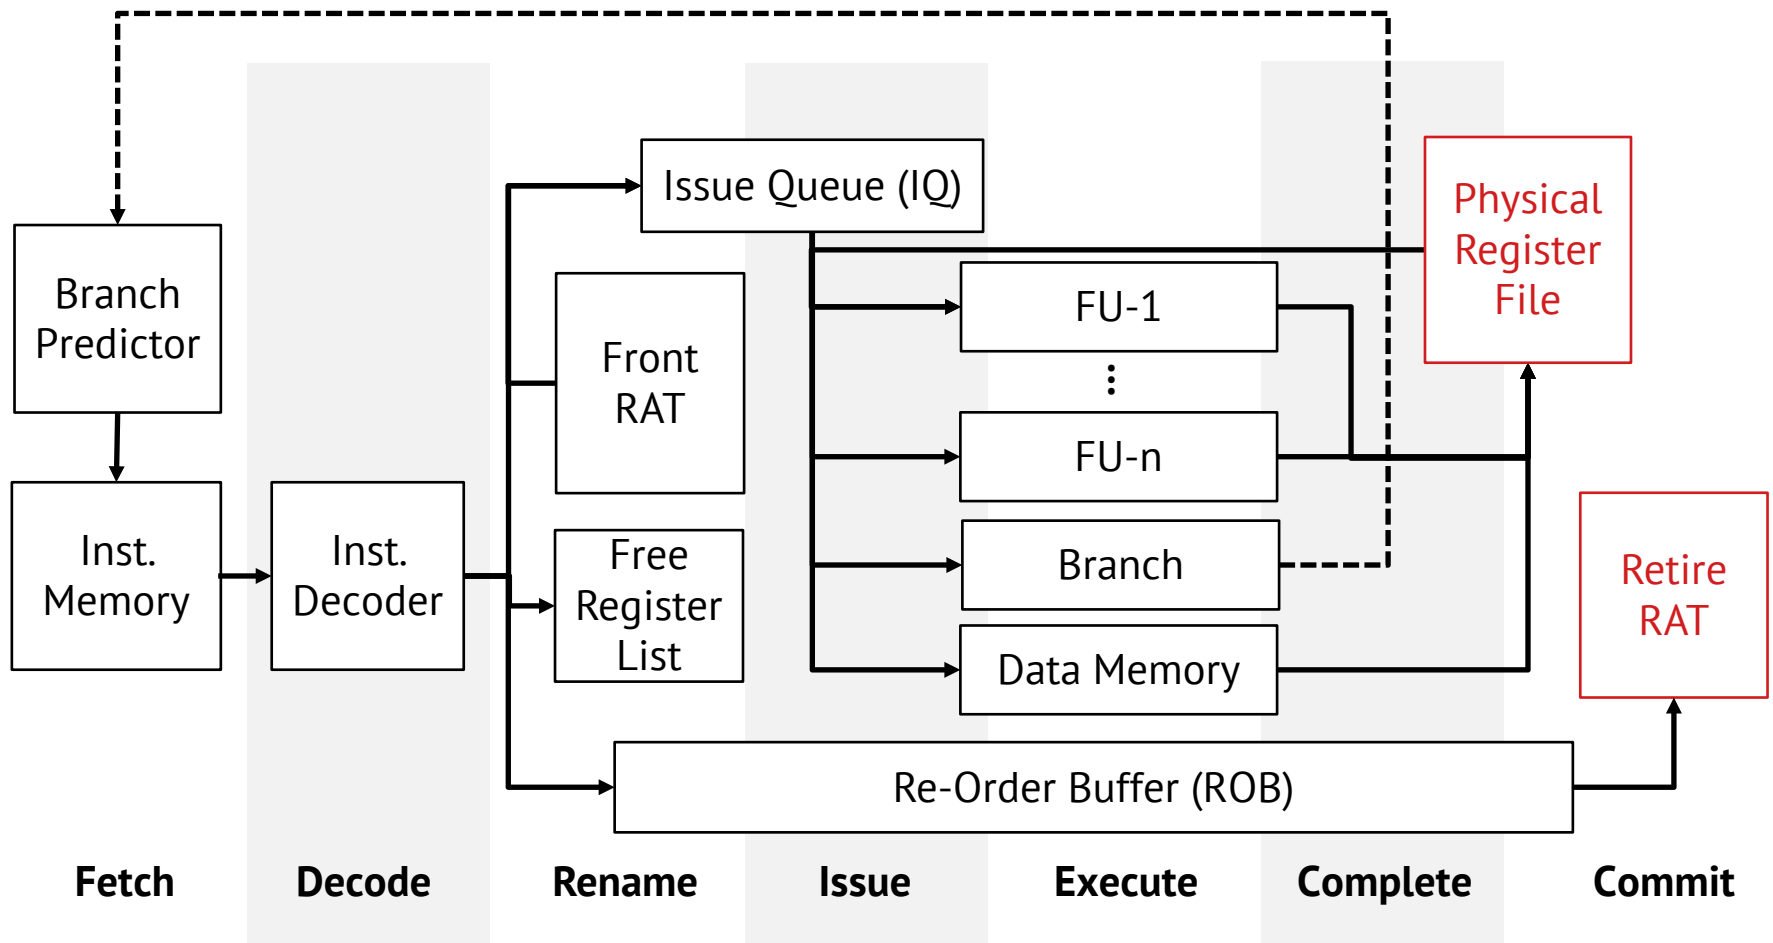
Phy. Reg. File

P0
P1
P2
P3
P4
P5
...

Note2: no need for storing values in ROB or IQ

Double RAT Architecture

- What is the size of ROB?



Physical Register Release

- Example: when is it safe to free p30 (R1)?

```
ADD R1, R2, R3
SUB R2, R1, R3
...
ADD R3, R1, R2
...
SUB R1, R3, R2
ADD R2, R1, R3
```

