

MEMORY HIERARCHY DESIGN

Mahdi Nazm Bojnordi

Assistant Professor

School of Computing

University of Utah


Overview

- Basics
 - ▣ Why use a memory hierarchy
 - ▣ Memory technologies
 - ▣ Principle of locality
- Caches
 - ▣ Concepts: block, index, tag, address, cache policies
 - ▣ Performance metric (AMAT) and optimization techniques
- Virtual Memory
 - ▣ Paging, address translation, TLB

Memory Hierarchy

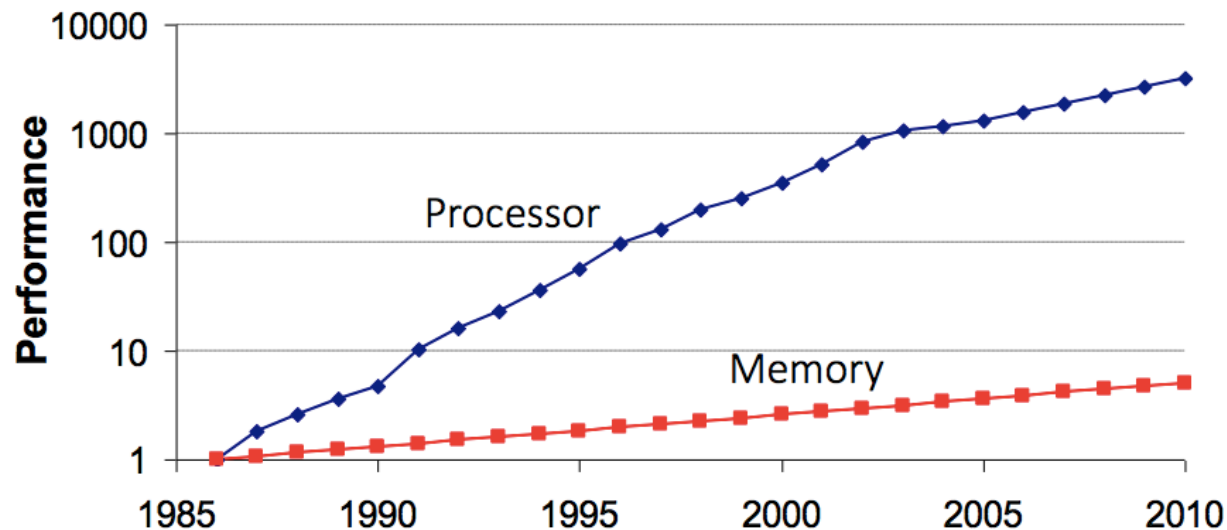
“Ideally one would desire an indefinitely large memory capacity such that any particular [...] word would be immediately available [...] We are [...] forced to recognize the possibility of constructing **a hierarchy of memories**, each of which has greater capacity than the preceding but which is less quickly accessible.”

-- *Burks, Goldstine, and von Neumann, 1946*




The Memory Wall

- Processor-memory performance gap increased over 50% per year
 - ▣ Processor performance historically improved $\sim 60\%$ per year
 - ▣ Main memory access time improves $\sim 5\%$ per year



Modern Memory Hierarchy

- Trade-off among memory speed, capacity, and cost




Memory Technology

- Random access memory (RAM) technology
 - ▣ access time same for all locations (not so true anymore)
- ▣ Static RAM (SRAM)
 - typically used for caches
 - 6T/bit; fast but – low density, high power, expensive
- ▣ Dynamic RAM (DRAM)
 - typically used for main memory
 - 1T/bit; inexpensive, high density, low power – but slow

RAM Cells


- 6T SRAM cell

- ▣ internal feedback maintains data while power on



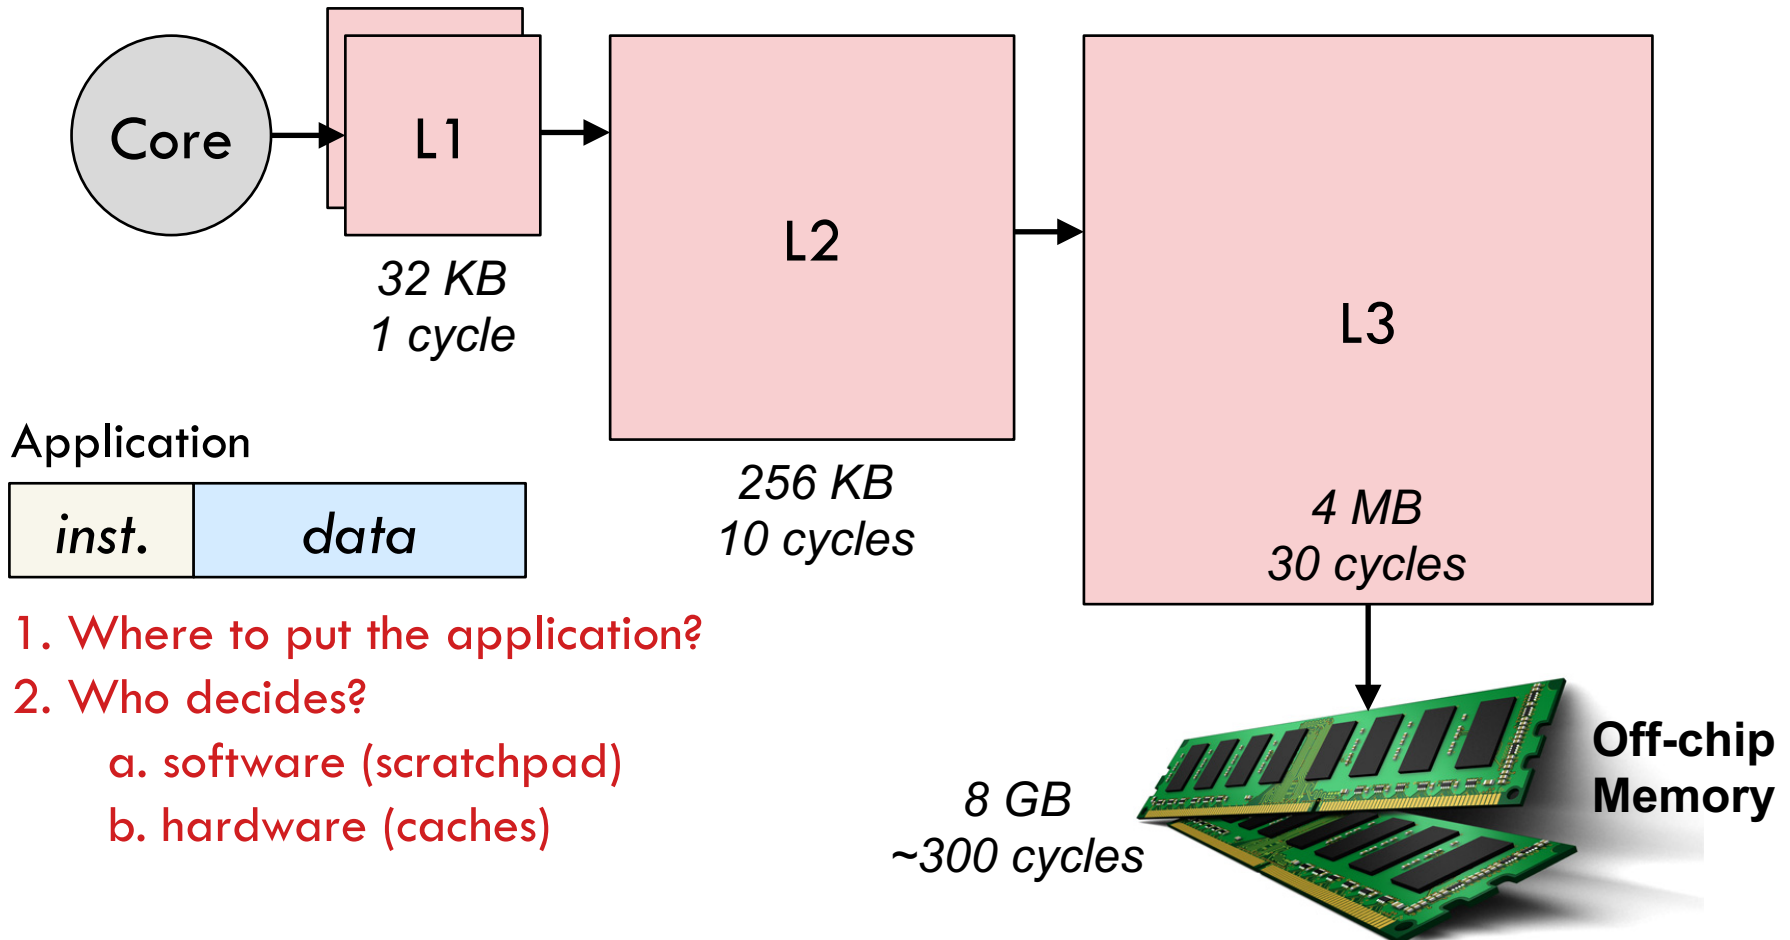
- 1T-1C DRAM cell

- ▣ needs refresh regularly to preserve data





Cache Hierarchy

□ Example three-level cache organization



Principle of Locality

- Memory references exhibit localized accesses
- Types of locality
 - ▣ *spatial*: probability of access to $A+\delta$ at time $t+\varepsilon$ highest when $\delta \rightarrow 0$
 - ▣ *temporal*: probability of accessing $A+\varepsilon$ at time $t+\delta$ highest when $\delta \rightarrow 0$



```
for (i=0; i<1000; ++i) {  
    sum = sum + a[i];  
}
```

temporal **spatial**

Red arrows point from the labels 'temporal' and 'spatial' to the 'sum' and 'a[i]' terms in the code, respectively.

Key idea: store local data in fast cache levels

Cache Terminology

- Block (cache line): unit of data access
- Hit: accessed data found at current level
 - ▣ *hit rate: fraction of accesses that finds the data*
 - ▣ *hit time: time to access data on a hit*
- Miss: accessed data NOT found at current level
 - ▣ miss rate: $1 - \text{hit rate}$
 - ▣ miss penalty: time to get block from lower level


hit time << miss penalty

Cache Performance

□ Average Memory Access Time (AMAT)

Outcome	Rate	Access Time
Hit	r_h	t_h
Miss	r_m	$t_h + t_p$

$$AMAT = r_h t_h + r_m (t_h + t_p)$$
$$r_h = 1 - r_m$$



$$AMAT = t_h + r_m t_p$$

problem: hit rate is 90%; hit time is 2 cycles;
and accessing the lower level takes 200 cycles;
find the average memory access time?

$$AMAT = 2 + 0.1 \times 200 = 22 \text{ cycles}$$

Example Problem

- Assume that the miss rate for instructions is 5%; the miss rate for data is 8%; the data references per instruction is 40%; and the miss penalty is 20 cycles; find performance relative to perfect cache with no misses
 - ▣ $\text{misses/instruction} = 0.05 + 0.08 \times 0.4 = 0.08$
 - ▣ Assuming hit time = 1
 - $\text{AMAT} = 1 + 0.08 \times 20 = 2.6$
 - $\text{Relative performance} = 1/2.6$