

ADVANCED MEMORY CONTROLLERS

Mahdi Nazm Bojnordi

Assistant Professor

School of Computing

University of Utah



The image part with
relationship ID rld2 was

CS/ECE 6810: Computer Architecture

Overview: DRAM Control Tasks

- Refresh management
 - ▣ Periodically replenish the DRAM cells (burst vs. distributed)
- Address mapping
 - ▣ Distribute the requests to destination banks (load balancing)
- Request scheduling
 - ▣ Generate a sequence of commands for memory requests
 - Reduce overheads by eliminating unnecessary commands
- Power management
 - ▣ Keep the power consumption under a cap
- Error detection/correction
 - ▣ Detect and recover corrupted data

Address Mapping


- A memory request

Type	Address	Data
------	---------	------

- Address is used to find the location in memory

- ▣ Channel, rank, bank, row, and column IDs

- Example physical address format



- A 4GB channel, 2 ranks, 4 banks/rank, 8KB page

Example Problem

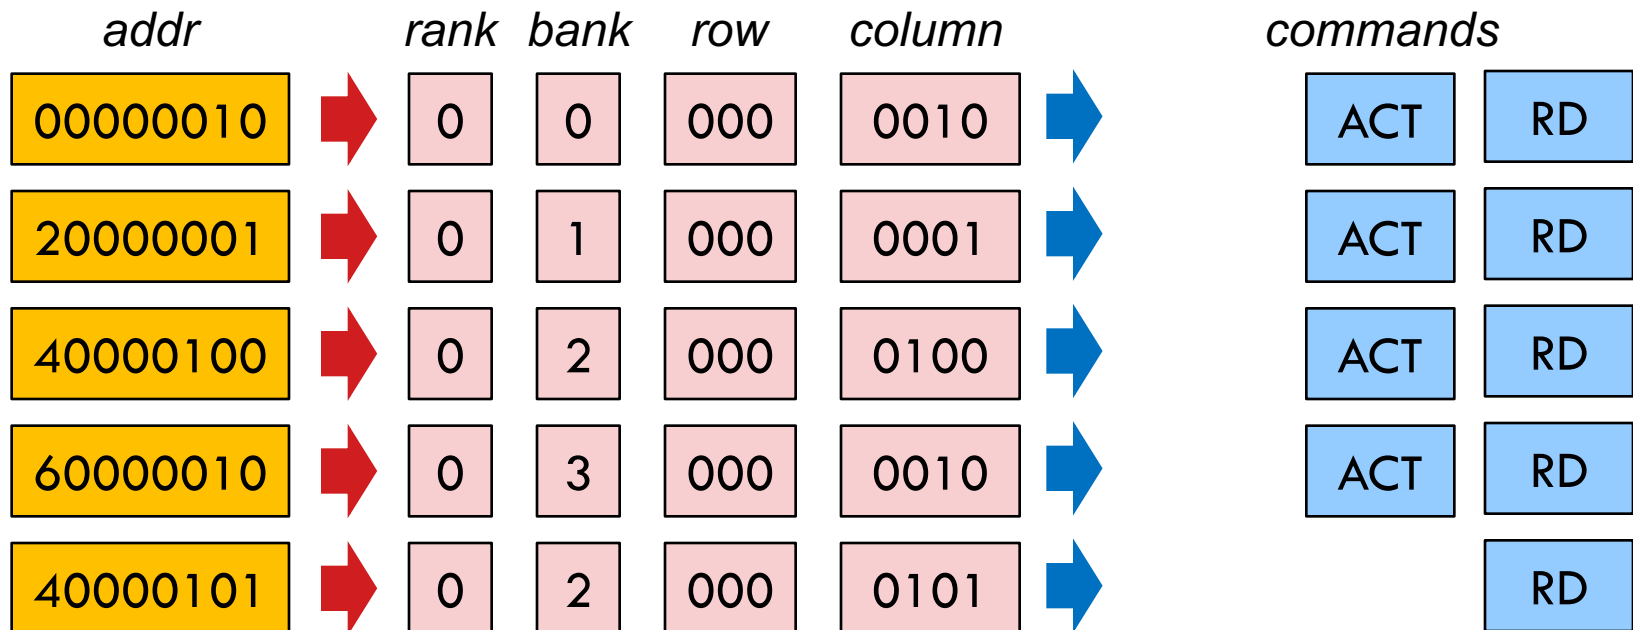
- Start with empty row buffers, find the total number of commands if all the request are served in order
 - Address= row(1 2):channel(0):rank(1):bank(3):column(1 6)

<i>addr</i>		<i>rank</i>	<i>bank</i>	<i>row</i>	<i>column</i>		<i>commands</i>		
00000010	➡	0	0	000	0010	➡		ACT	RD
20000001	➡	0	0	200	0001	➡	PRE	ACT	RD
40000100	➡	0	0	400	0100	➡	PRE	ACT	RD
60000010	➡	0	0	600	0010	➡	PRE	ACT	RD
40000101	➡	0	0	400	0101	➡	PRE	ACT	RD

Example Problem

- Find the total number of commands using the following address mapping scheme

■ $\text{Address} = \text{bank}(3):\text{rank}(1):\text{channel}(0):\text{row}(1\ 2):\text{column}(1\ 6)$



Row Buffer Management Policies

- Open-page policy
 - ▣ After access, keep page in DRAM row buffer
 - ▣ If access to different page, must close old one first
 - Good if lots of locality
- Close-page policy
 - ▣ After access, immediately close page in DRAM row buffer
 - ▣ If access to different page, old one already closed
 - Good if no locality (random access)

Command Scheduling

- Write buffering
 - ▣ Writes can wait until reads are done
- Controller queues DRAM commands
 - ▣ Usually into per-bank queues
 - ▣ Allows easily reordering ops. meant for same bank
- Two common policies
 - ▣ First-Come-First-Served (FCFS)
 - In order request scheduling
 - ▣ First-Ready First-Come-First-Served (FR-FCFS)
 - Out of order request scheduling

Command Scheduling


- First-Come-First-Served
 - ▣ Oldest request first
- First-Ready First-Come-First-Served
 - ▣ Prioritize column accesses over row changes
 - ▣ Skip over older conflicting requests
 - ▣ Find row hits (on queued requests)
 - Find oldest
 - If no conflicts with in-progress request → good
 - Otherwise (if conflicts), try next oldest

FCFS vs. FR-FCFS


- Single bank memory

- READ(B0,R0,C0) READ(B0,R1,C0) READ(B0,R0,C1)

- FCFS



- FR-FCFS




Error Detection/Correction

- ❑ Data in memory may be corrupted
 - ▣ Many reasons: leakage, alpha particles, hard errors.
- ❑ Can errors be detected?
 - ▣ **Error detection codes**: additional parity bits
- ❑ Can errors be corrected?
 - ▣ **Error correction codes**: ECC bits are added to data
- ❑ Single-Error Correction, Double-Error Detection
 - ▣ Commonly used in memory systems


ECC DIMM

- An additional DRAM chip is used for storing SECDED ECC bits for error correction



Limitations to Existing Memory Controllers

- Modern memory controllers are performance-critical and complex



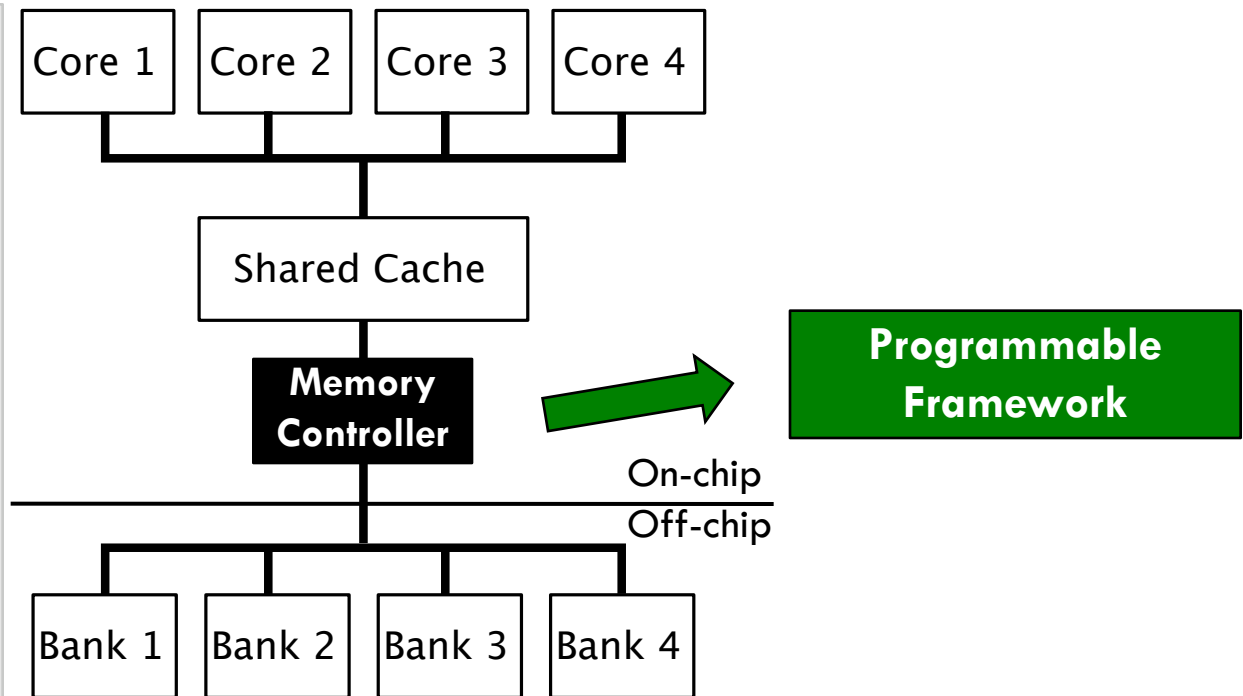
Programmable Memory Controllers

- Programmability can make a memory controller higher-performance and more flexible

Multiple performance objectives ✓


Application-specific optimizations ✓

Patches and in-field updates ✓




Design Overview

- **Key idea:** Judicious division of labor between specialized hardware and firmware
 - ▣ Request and transaction processing in firmware
 - ▣ Configurable timing validation in hardware



Request Processing

- A RISC ISA for operating on memory requests

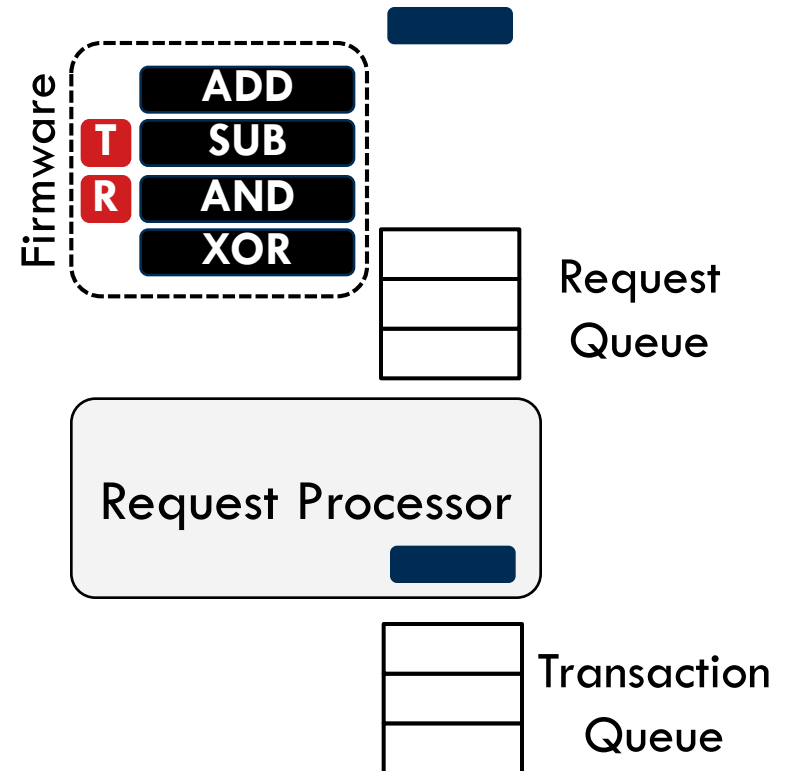


Request Processing

- Queue management with instruction flags

- ▣ R flag enqueues a request
- ▣ T flag dequeues a transaction

- An instruction can be annotated with both R and T flags if needed



Implementation

- Two five-stage pipelines and one configurable timing validation circuit

