

GRAPHICS PROCESSING UNIT

Mahdi Nazm Bojnordi

Assistant Professor

School of Computing

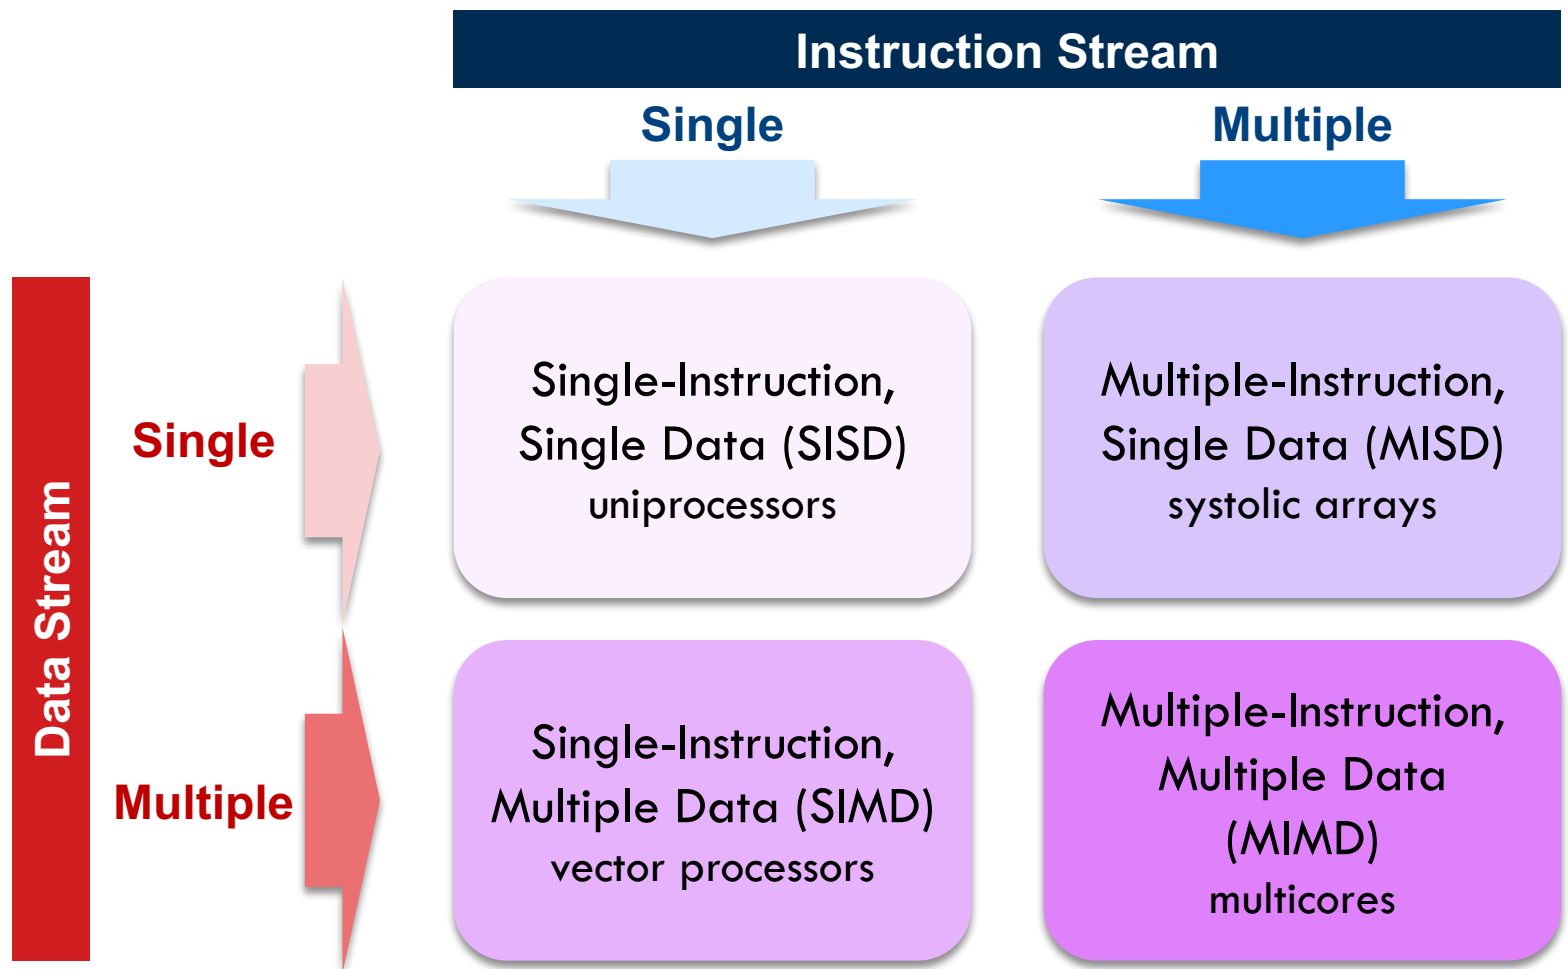
University of Utah

Overview

- Announcement
 - ▣ Homework 7 will be available tonight (due on 12/02)
- This lecture
 - ▣ Classification of parallel computers
 - ▣ Graphics processing
 - ▣ GPU architecture
 - ▣ CUDA programming model

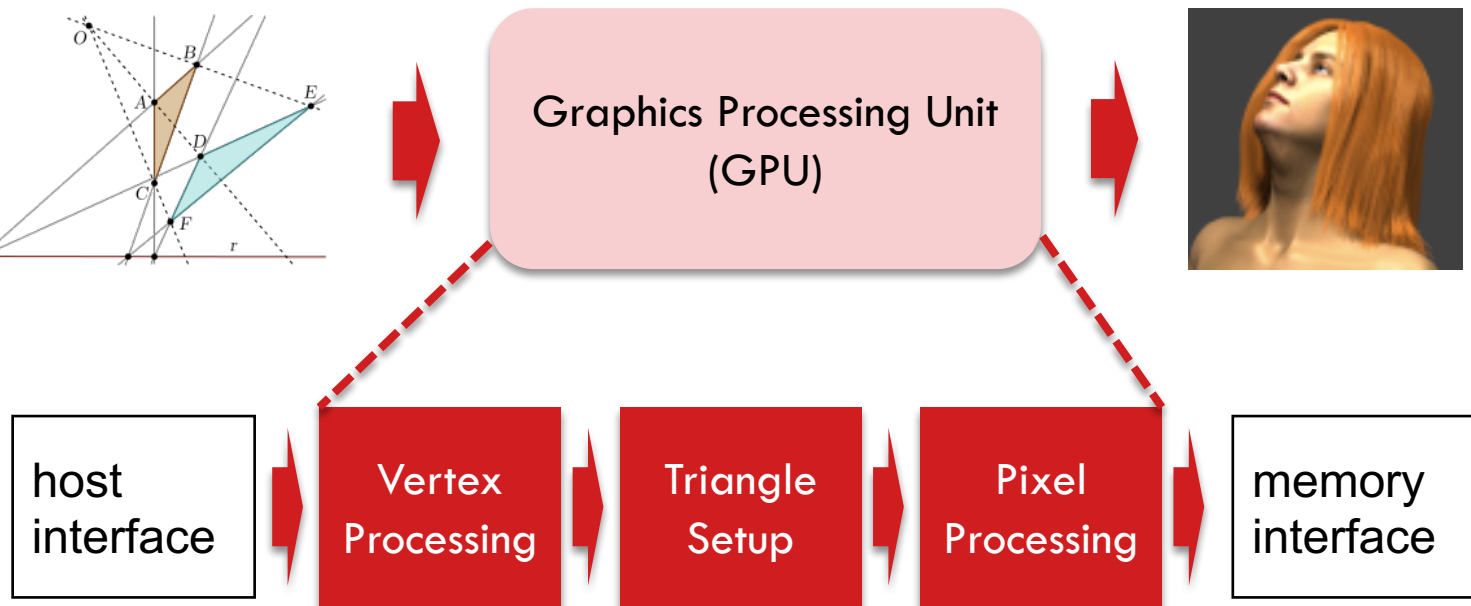

Flynn's Taxonomy

- Data vs. instruction streams



Graphics Processing Unit

- Initially developed as graphics accelerator
 - ▣ It receives geometry information from the CPU as an input and provides a picture as an output




Host Interface

- The host interface is the communication bridge between the CPU and the GPU
- It receives commands from the CPU and also pulls geometry information from system memory
- It outputs a *stream* of vertices in object space with all their associated information


Vertex Processing

- The vertex processing stage receives vertices from the host interface in object space and outputs them in screen space
- This may be a simple linear transformation, or a complex operation involving morphing effects



Pixel Processing

- ❑ Rasterize triangles to pixels
- ❑ Each fragment provided by triangle setup is fed into fragment processing as a set of attributes (position, normal, texcoord etc), which are used to compute the final color for this pixel
- ❑ The computations taking place here include texture mapping and math operations



Programming GPUs

- The programmer can write programs that are executed for every vertex as well as for every fragment
- This allows fully customizable geometry and **shading** effects that go well beyond the generic look and feel of older 3D applications




Memory Interface

- ❑ Fragment colors provided by the previous stage are written to the **framebuffer**
- ❑ Used to be the biggest bottleneck before fragment processing took over
- ❑ Before the final write occurs, some fragments are rejected by the **zbuffer**, stencil and alpha tests
- ❑ On modern GPUs, z and color are compressed to reduce framebuffer bandwidth (but not size)

Z-Buffer

- Example of 3 objects





Graphics Processing Unit

- Initially developed as graphics accelerators
 - ▣ one of the densest compute engines available now
- Many efforts to run non-graphics workloads on GPUs
 - ▣ general-purpose GPUs (GPGPUs)
- C/C++ based programming platforms
 - ▣ CUDA from NVidia and OpenCL from an industry consortium
- A heterogeneous system
 - ▣ a regular host CPU
 - ▣ a GPU that handles CUDA (may be on the same CPU chip)


Graphics Processing Unit

- Simple in-order pipelines that rely on thread-level parallelism to hide long latencies
- Many registers ($\sim 1\text{K}$) per in-order pipeline (lane) to support many active warps




The GPU Architecture

- SIMT – single instruction, multiple threads
 - ▣ GPU has many SIMT cores
- Application → many thread blocks (1 per SIMT core)
- Thread block → many warps (1 warp per SIMT core)
- Warp → many in-order pipelines (SIMD lanes)




Why GPU Computing?



Source: NVIDIA


GPU Computing

- GPU as an accelerator in scientific applications



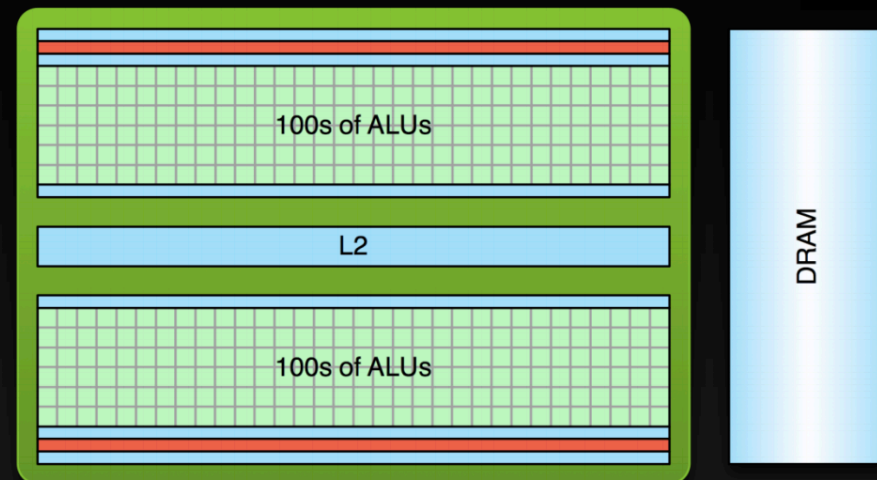
GPU Computing

□ Low latency or high throughput?



CPU

- Optimized for low-latency access to cached data sets
- Control logic for out-of-order and speculative execution

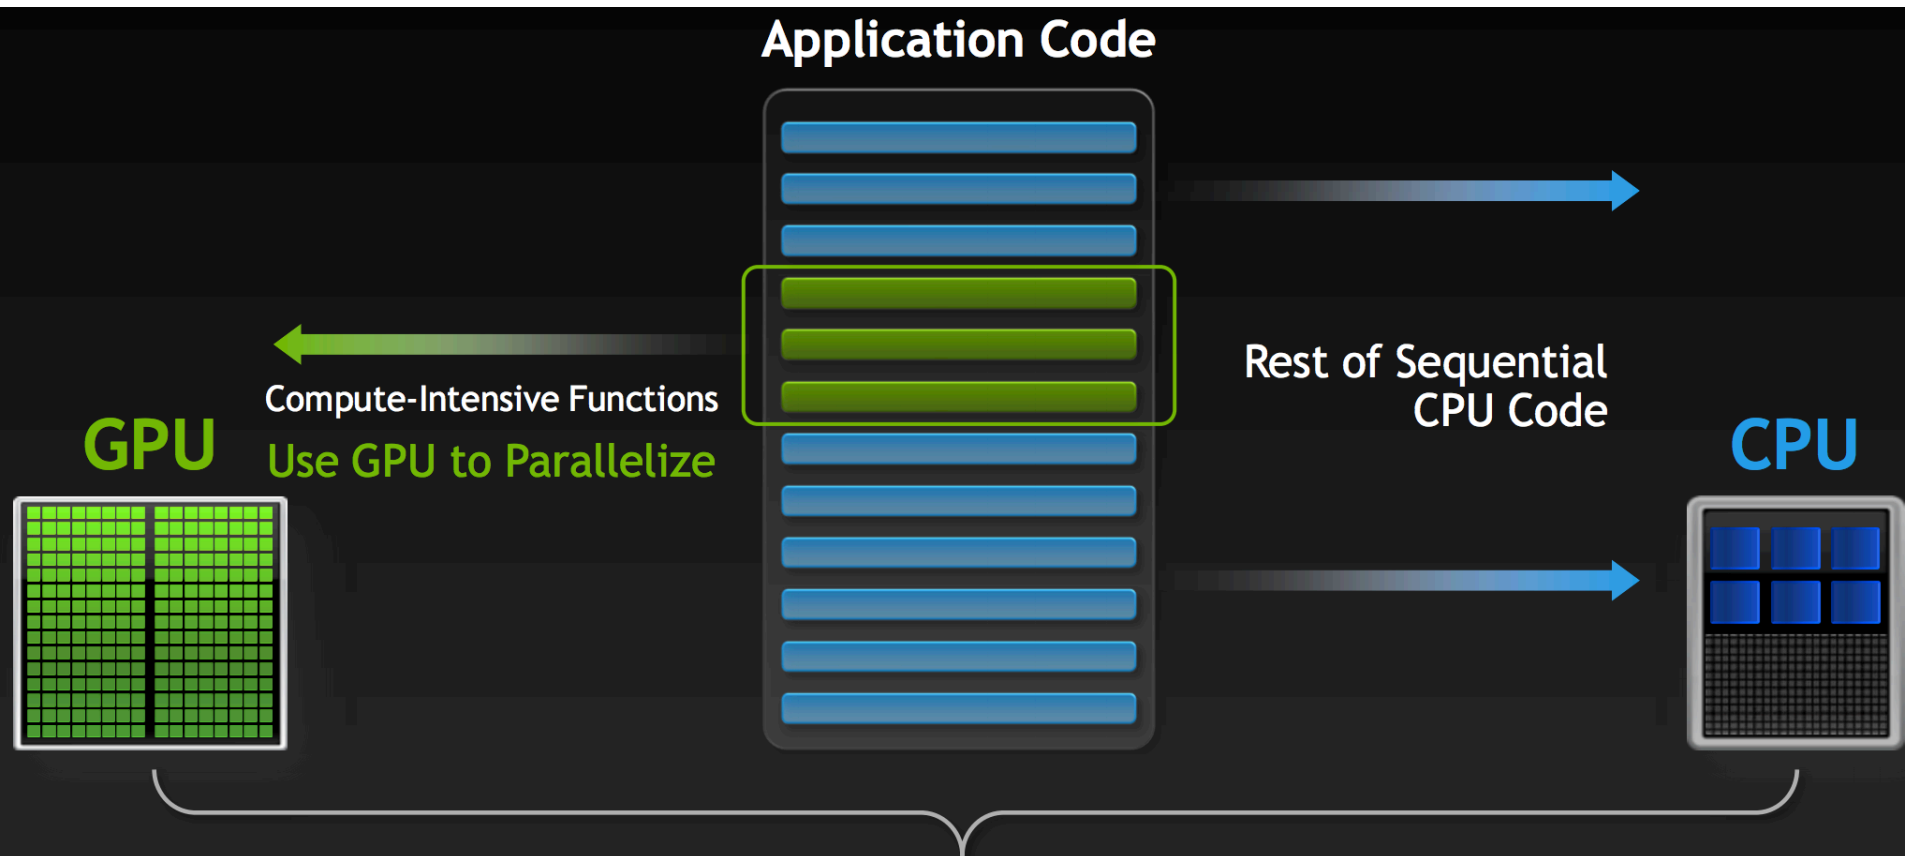


GPU

- Optimized for data-parallel, throughput computation
- Architecture tolerant of memory latency
- More transistors dedicated to computation

GPU Computing

- Low latency or high throughput



CUDA Programming Model

- Step 1: substitute library calls with equivalent CUDA library calls
 - ▣ `saxpy (...)` → `cublasSaxpy (...)`
 - single precision alpha x plus y ($z = \alpha x + y$)
- Step 2: manage data locality
 - ▣ `cudaMalloc()`, `cudaMemcpy()`, etc.
- Step 3: transfer data between CPU and GPU
 - ▣ get and set functions
- rebuild and link the CUDA-accelerated library
 - ▣ `nvcc myobj.o -l cublas`

Example: SAXPY Code

```
int N = 1 << 20;
```

```
// Perform SAXPY on 1M elements:  $y[] = a * x[] + y[]$   
saxpy(N, 2.0, x, 1, y, 1);
```

Example: CUDA Lib Calls

```
int N = 1 << 20;
```

```
// Perform SAXPY on 1M elements: d_y[] = a*d_x[] + d_y[]  
cublasSaxpy(N, 2.0, d_x, 1, d_y, 1);
```

Example: Initialize CUDA Lib

```
int N = 1 << 20;
```

```
cublasInit();
```

```
// Perform SAXPY on 1M elements: d_y[] = a*d_x[] + d_y[]
```

```
cublasSaxpy(N, 2.0, d_x, 1, d_y, 1);
```

```
cublasShutdown();
```

Example: Allocate Memory

```
int N = 1 << 20;
```

```
cublasInit();
```

```
cublasAlloc(N, sizeof(float), (void**)&d_x);
```

```
cublasAlloc(N, sizeof(float), (void*)&d_y);
```

```
// Perform SAXPY on 1M elements: d_y[] = a*d_x[] + d_y[]
```

```
cublasSaxpy(N, 2.0, d_x, 1, d_y, 1);
```

```
cublasFree(d_x);
```

```
cublasFree(d_y);
```

```
cublasShutdown();
```

Example: Transfer Data

```
int N = 1 << 20;
```

```
cublasInit();
```

```
cublasAlloc(N, sizeof(float), (void**)&d_x);
```

```
cublasAlloc(N, sizeof(float), (void*)&d_y);
```

```
cublasSetVector(N, sizeof(x[0]), x, 1, d_x, 1);
```

```
cublasSetVector(N, sizeof(y[0]), y, 1, d_y, 1);
```

```
// Perform SAXPY on 1M elements: d_y[] = α * d_x[] + d_y[]
```

```
cublasSaxpy(N, 2.0, d_x, 1, d_y, 1);
```

```
cublasGetVector(N, sizeof(y[0]), d_y, 1, y, 1);
```

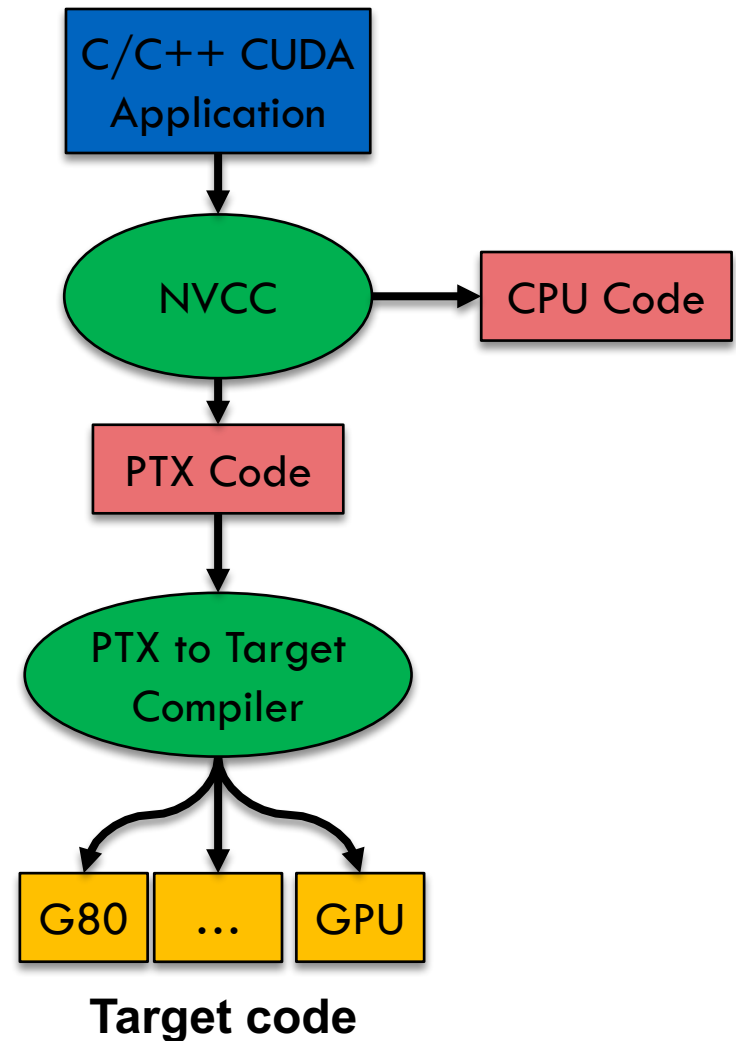
```
cublasFree(d_x);
```

```
cublasFree(d_y);
```

```
cublasShutdown();
```

Compiling CUDA

- Call nvcc
- Parallel Threads eXecution (PTX)
 - ▣ Virtual machine and ISA
- Two stage
 - ▣ 1. PTX
 - ▣ 2. device-specific binary object



Memory Hierarchy

- Throughput-oriented main memory

- ▣ Graphics DDR (GDDR)

- Wide channels: 256 bit
 - Lower clock rate than DDR

- ▣ 1.5MB shared L2

- ▣ 48KB read-only data cache

- Compiler controlled

- ▣ Wide buses

