

BRANCH PREDICTORS

Mahdi Nazm Bojnordi

Assistant Professor

School of Computing

University of Utah

Overview

- Announcements

 - ▣ Homework 2 release: Sept. 26th

- This lecture

 - ▣ Dynamic branch prediction

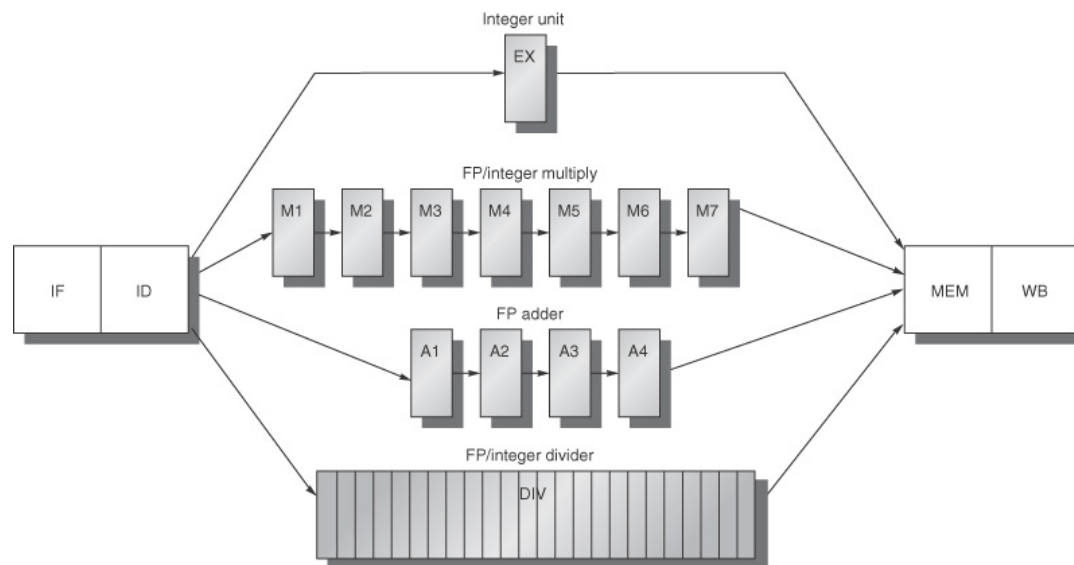
 - ▣ Counter based branch predictor

 - ▣ Correlating branch predictor

 - ▣ Global vs. local branch predictors

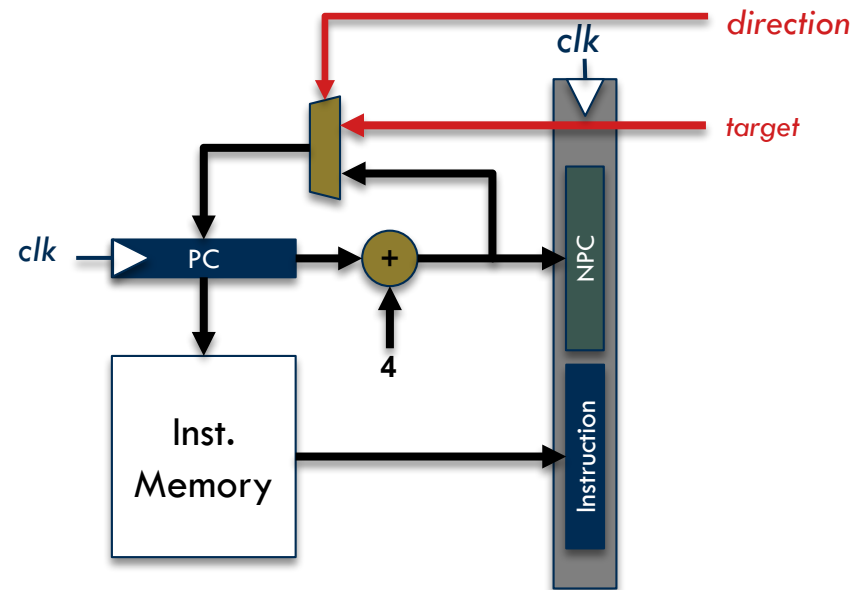
Big Picture: Why Branch Prediction?

- **Problem:** performance is mainly limited by the number of instructions fetched per second
- **Solution:** deeper and wider frontend
- **Challenge:** handling branch instructions



Big Picture: How to Predict Branch?

- Static prediction (based on direction or profile)
 - Always not-taken
 - Target = next PC
 - Always taken
 - Target = unknown
- Dynamic prediction
 - Special hardware using PC



Recall: Dynamic Branch Prediction

- Hardware unit capable of learning at runtime
 - ▣ 1. Prediction logic
 - Direction (taken or not-taken)
 - Target address (where to fetch next)
 - ▣ 2. Outcome validation and training
 - Outcome is computed regardless of prediction
 - ▣ 3. Recovery from misprediction
 - Nullify the effect of instructions on the wrong path

Branch Prediction

- **Goal:** avoiding stall cycles caused by branches
- **Solution:** static or dynamic branch predictor
 - ▣ 1. prediction
 - ▣ 2. validation and training
 - ▣ 3. recovery from misprediction
- **Performance** is influenced by the frequency of branches **(b)**, prediction accuracy **(a)**, and misprediction cost **(c)**

Branch Prediction

- **Goal:** avoiding stall cycles caused by branches
- **Solution:** static or dynamic branch predictor
 - ▣ 1. prediction
 - ▣ 2. validation and training
 - ▣ 3. recovery from misprediction
- **Performance** is influenced by the frequency of branches (**b**), prediction accuracy (**a**), and misprediction cost (**c**)

$$Speedup = \frac{Old\ Time}{New\ Time} = \frac{CPI_{old}}{CPI_{new}} = \frac{1 + bc}{1 + (1 - a)bc}$$

Problem

- A pipelined processor requires 3 stall cycles to compute the outcome of every branch before fetching next instruction; due to perfect forwarding/bypassing, no stall cycles are required for data/structural hazards; every 5th instruction is a branch.
- Compute speedup gained by a branch predictor with 90% accuracy

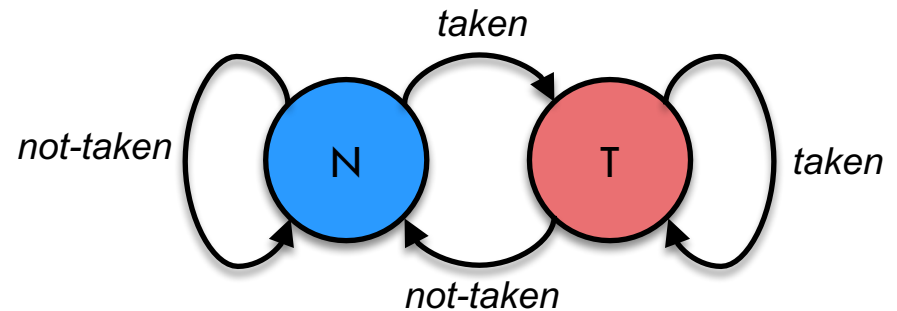
Problem

- A pipelined processor requires 3 stall cycles to compute the outcome of every branch before fetching next instruction; due to perfect forwarding/bypassing, no stall cycles are required for data/structural hazards; every 5th instruction is a branch.
- Compute speedup gained by a branch predictor with 90% accuracy

$$\text{Speedup} = (1 + 0.2 \times 3) / (1 + 0.1 \times 0.2 \times 3) = 1.5$$

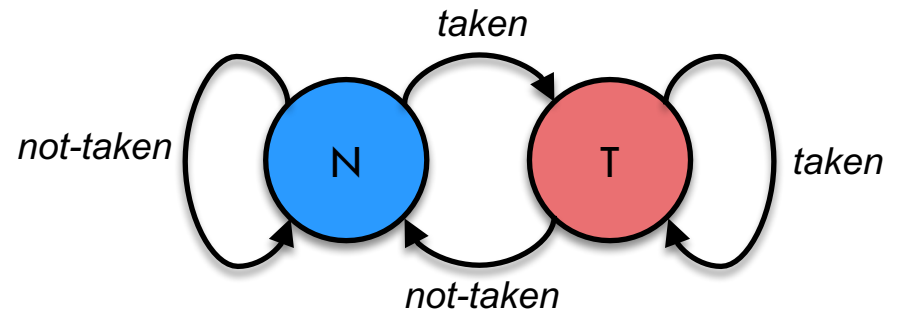
Bimodal Branch Predictors

- One-bit branch predictor
 - ▣ Keep track of and use the outcome of last branch



Bimodal Branch Predictors

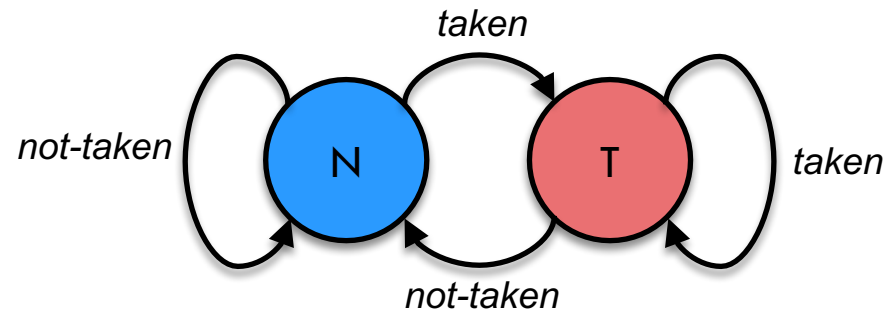
- One-bit branch predictor
 - ▣ Keep track of and use the outcome of last branch



```
while(1) {  
    for(i=0; i<10; i++) {  
    }  
    for(j=0; j<20; j++) {  
    }  
}
```

Bimodal Branch Predictors

- One-bit branch predictor
 - ▣ Keep track of and use the outcome of last branch



```
while(1) {  
    for(i=0; i<10; i++) {  
    }  
    for(j=0; j<20; j++) {  
    }  
}
```

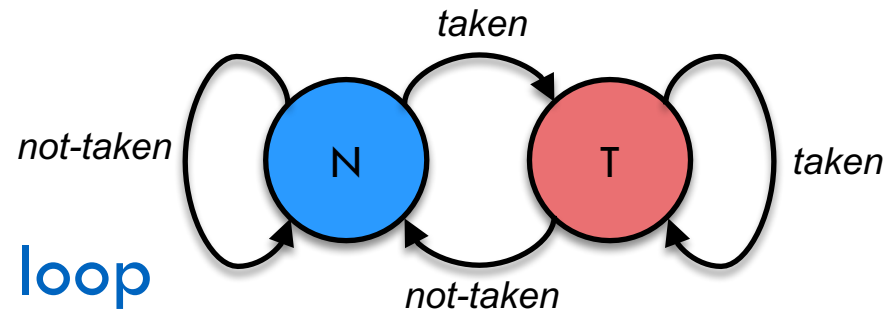
branch-1

branch-2

Bimodal Branch Predictors

- One-bit branch predictor
 - ▣ Keep track of and use the outcome of last branch

- Shared predictor
- Two mispredictions per loop



```
while(1) {  
    for(i=0; i<10; i++) {  
    }  
    for(j=0; j<20; j++) {  
    }  
}
```

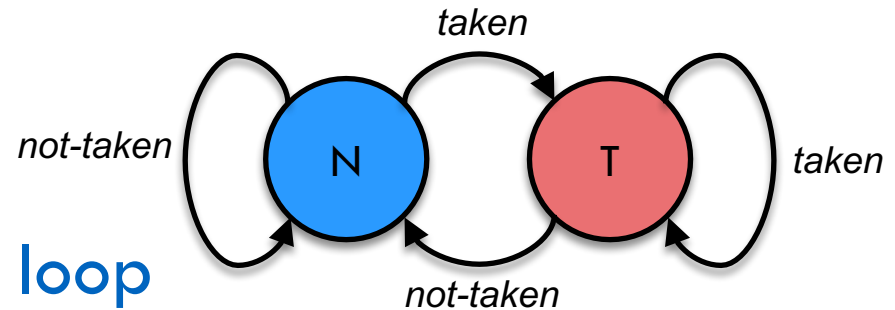
branch-1

branch-2

Bimodal Branch Predictors

- One-bit branch predictor
 - ▣ Keep track of and use the outcome of last branch

- Shared predictor
- Two mispredictions per loop



Accuracy = $26/30 = 0.86$

How to improve?

```
while(1) {  
    for(i=0; i<10; i++) {  
    }  
    for(j=0; j<20; j++) {  
    }  
}
```

branch-1
branch-2

Bimodal Branch Predictors

- Two-bit branch predictor
 - ▣ Increment if taken
 - ▣ Decrement if untaken

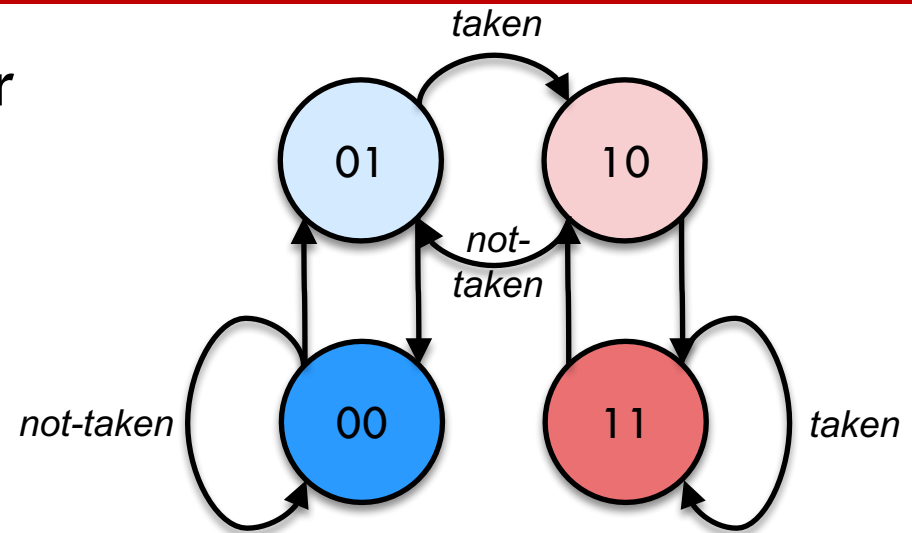
```
while(1) {  
    for(i=0; i<10; i++) {  
    }  
    for(j=0; j<20; j++) {  
    }  
}
```

branch-1

branch-2

Bimodal Branch Predictors

- Two-bit branch predictor
 - ▣ Increment if taken
 - ▣ Decrement if untaken



```
while(1) {  
    for(i=0; i<10; i++) {  
    }  
    for(j=0; j<20; j++) {  
    }  
}
```

branch-1

branch-2

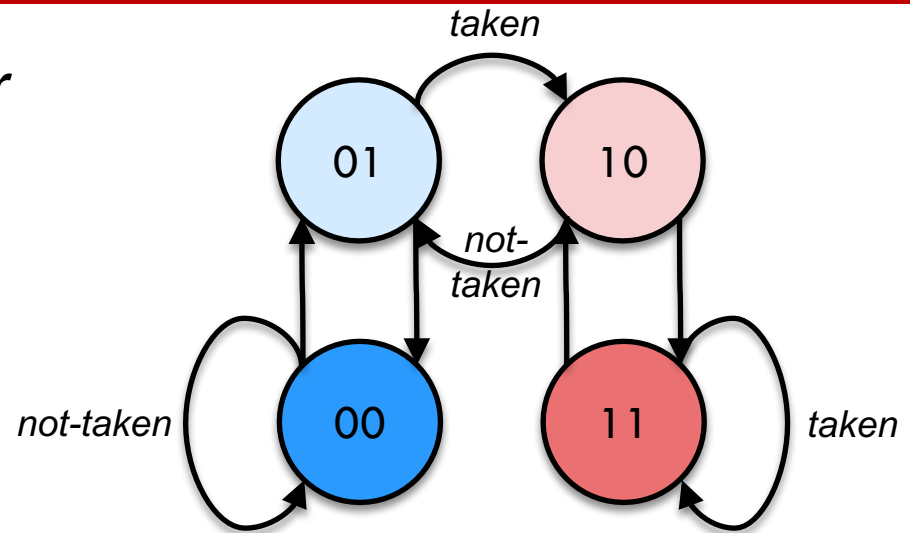
Bimodal Branch Predictors

□ Two-bit branch predictor

▣ Increment if taken

▣ Decrement if untaken

- One misprediction on loop exit
- Accuracy = $28/30 = 0.93$



```
while(1) {  
    for(i=0; i<10; i++) {  
    }  
    for(j=0; j<20; j++) {  
    }  
}
```

branch-1

branch-2

Bimodal Branch Predictors

□ Two-bit branch predictor

▣ Increment if taken

▣ Decrement if untaken

- One misprediction on loop exit

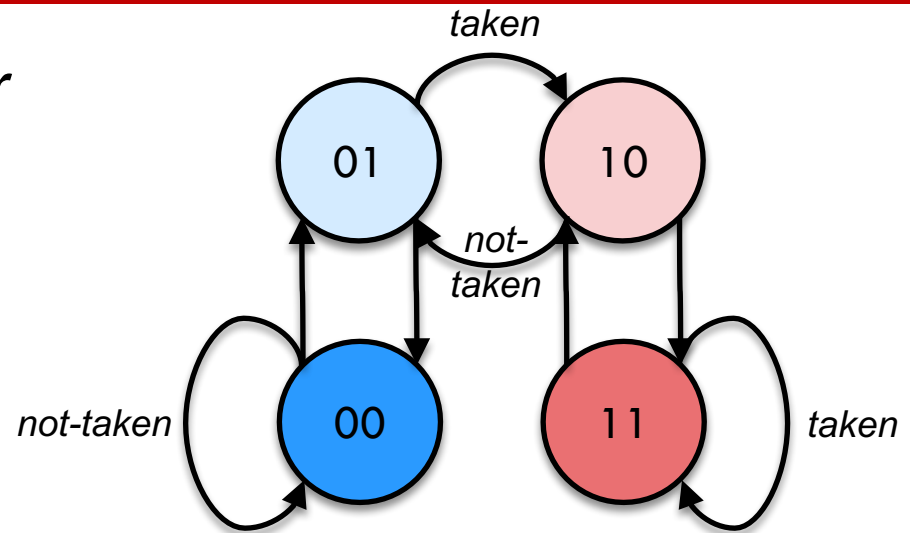
• Accuracy = $28/30 = 0.93$

• How to improve?

- 3-bit predictor?

• Problem?

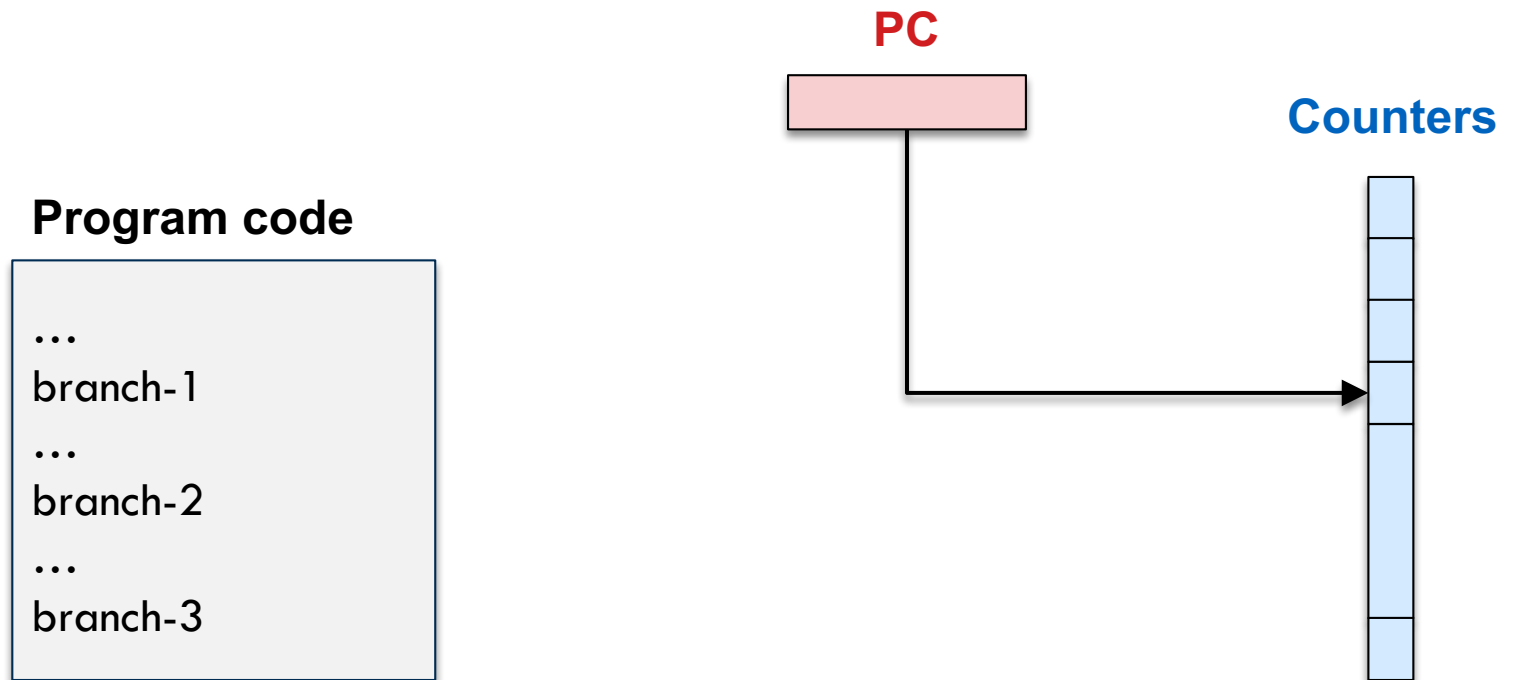
- A single predictor shared among many branches



```
while(1) {
    for(i=0; i<10; i++) {      branch-1
    }
    for(j=0; j<20; j++) {      branch-2
    }
}
```


Using Multiple Counters

- How to assign a branch to each counter?



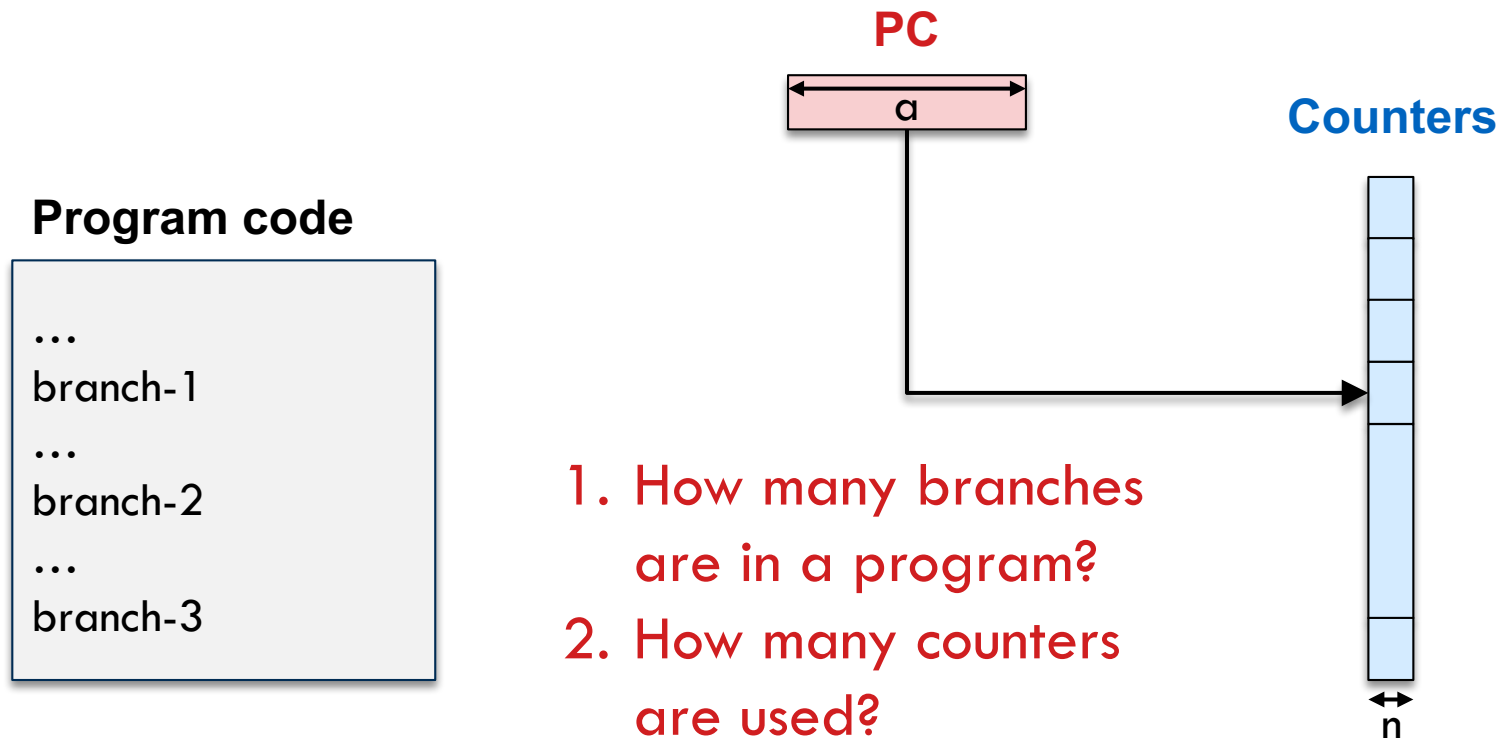
Using Multiple Counters

- How to assign a branch to each counter?



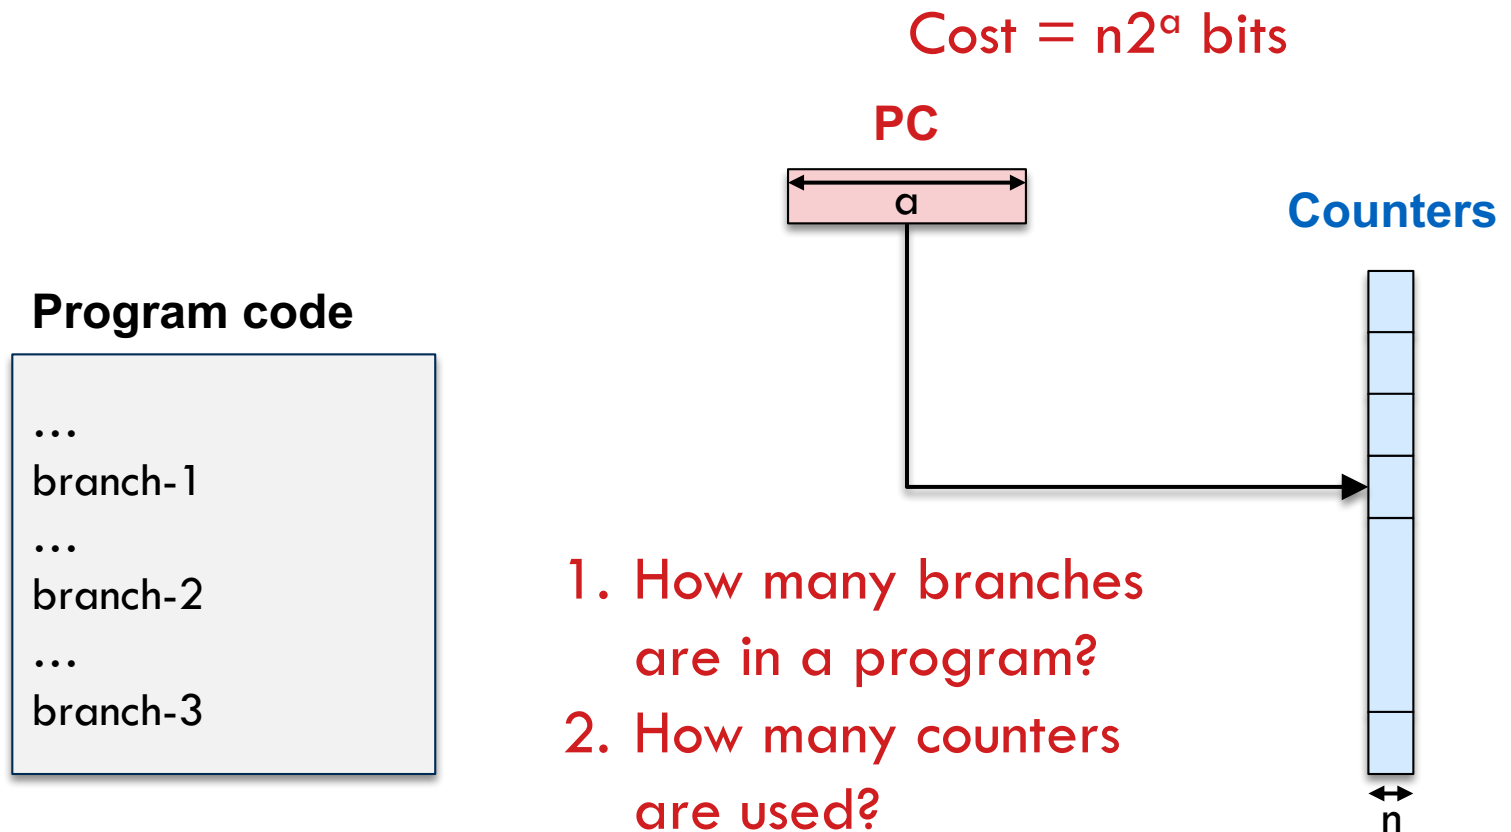
Using Multiple Counters

- How to assign a branch to each counter?



Using Multiple Counters

- How to assign a branch to each counter?

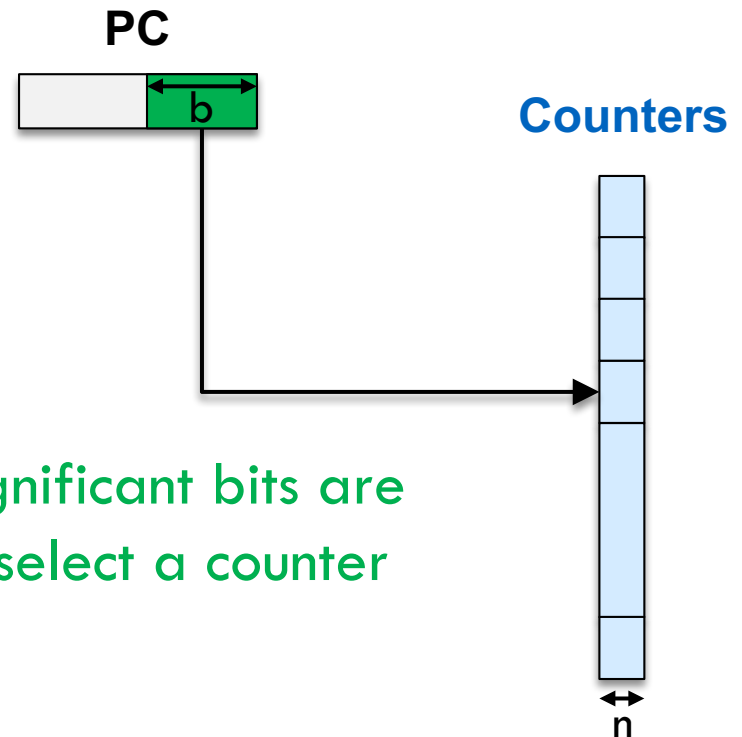


Using Multiple Counters

- How to assign a branch to each counter?
 - ▣ Decode History Table (DHT)
 - Reduced HW with aliasing

Program code

```
...  
branch-1  
...  
branch-2  
...  
branch-3
```

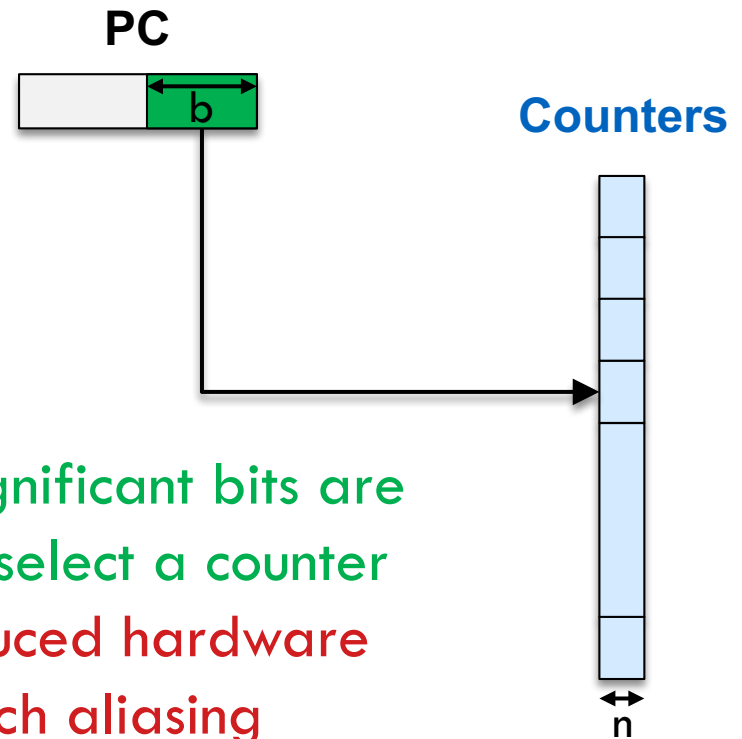


Using Multiple Counters

- How to assign a branch to each counter?
 - ▣ Decode History Table (DHT)
 - Reduced HW with aliasing

Program code

```
...  
branch-1  
...  
branch-2  
...  
branch-3
```



Using Multiple Counters

□ How to assign a branch to each counter?

▣ Decode History Table (DHT)

■ Reduced HW with aliasing

Program code

```
...  
branch-1  
...  
branch-2  
...  
branch-3
```

Cost = $n2^b$ bits

PC



Counters



Least significant bits are
used to select a counter
(+) Reduced hardware
(-) Branch aliasing

Using Multiple Counters

□ How to assign a branch to each counter?

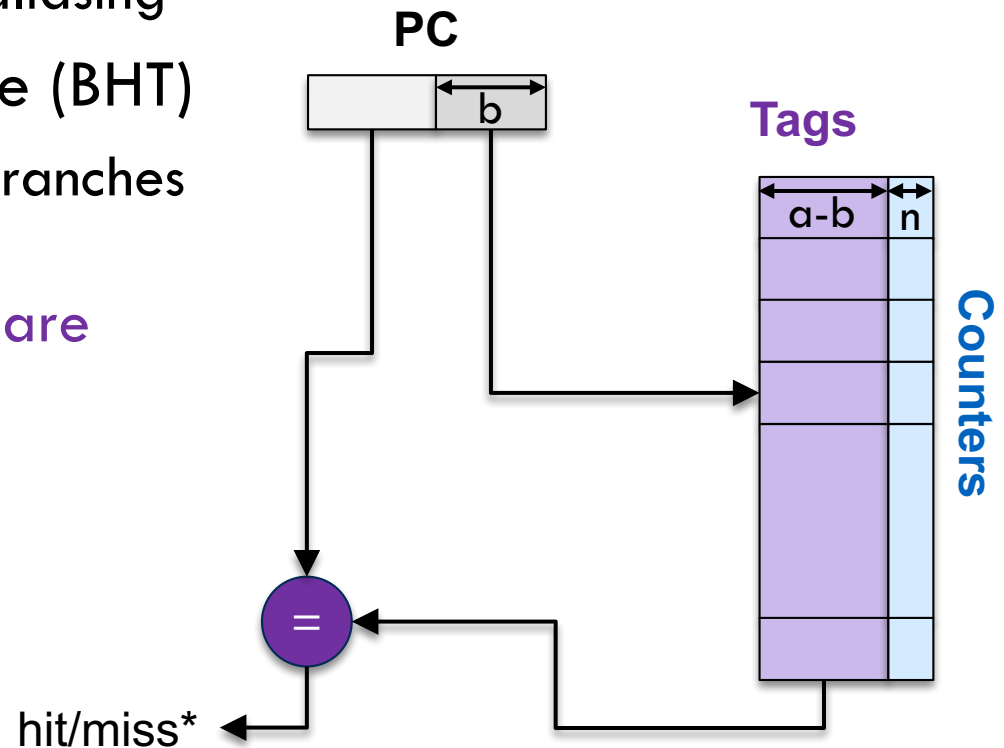
▣ Decode History Table (DHT)

■ Reduced HW with aliasing

▣ Branch History Table (BHT)

■ Precisely tracking branches

Most significant bits are
used as tags



Using Multiple Counters

□ How to assign a branch to each counter?

▣ Decode History Table (DHT)

■ Reduced HW with aliasing

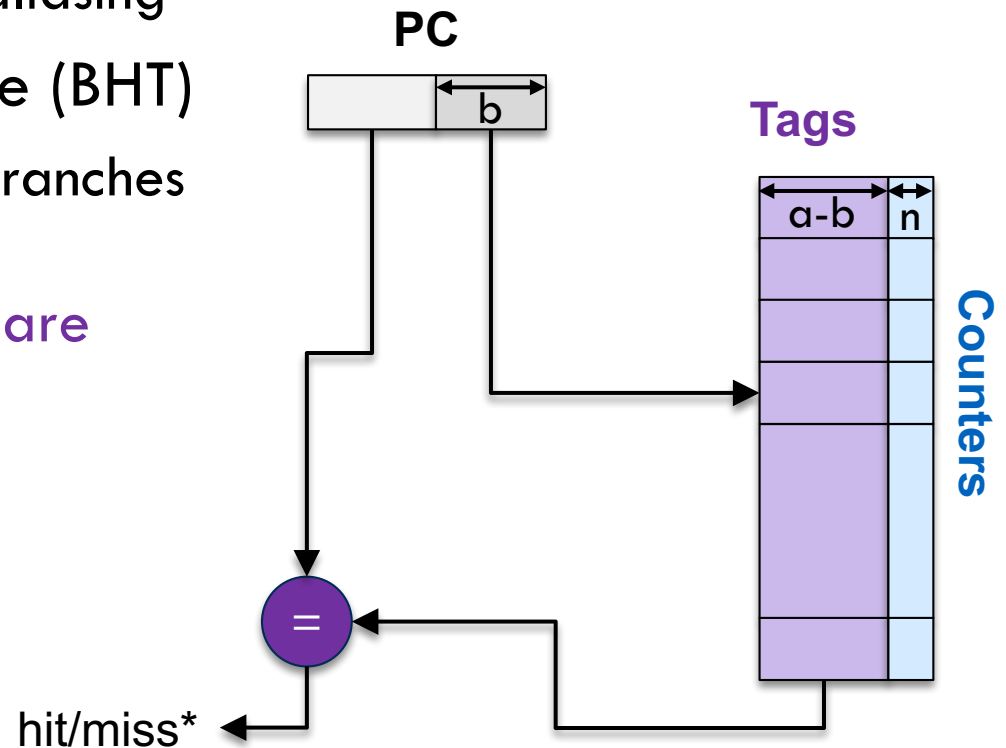
▣ Branch History Table (BHT)

■ Precisely tracking branches

Most significant bits are
used as tags

(+) No aliasing

(-) Missing entries



Using Multiple Counters

□ How to assign a branch to each counter?

▣ Decode History Table (DHT)

■ Reduced HW with aliasing

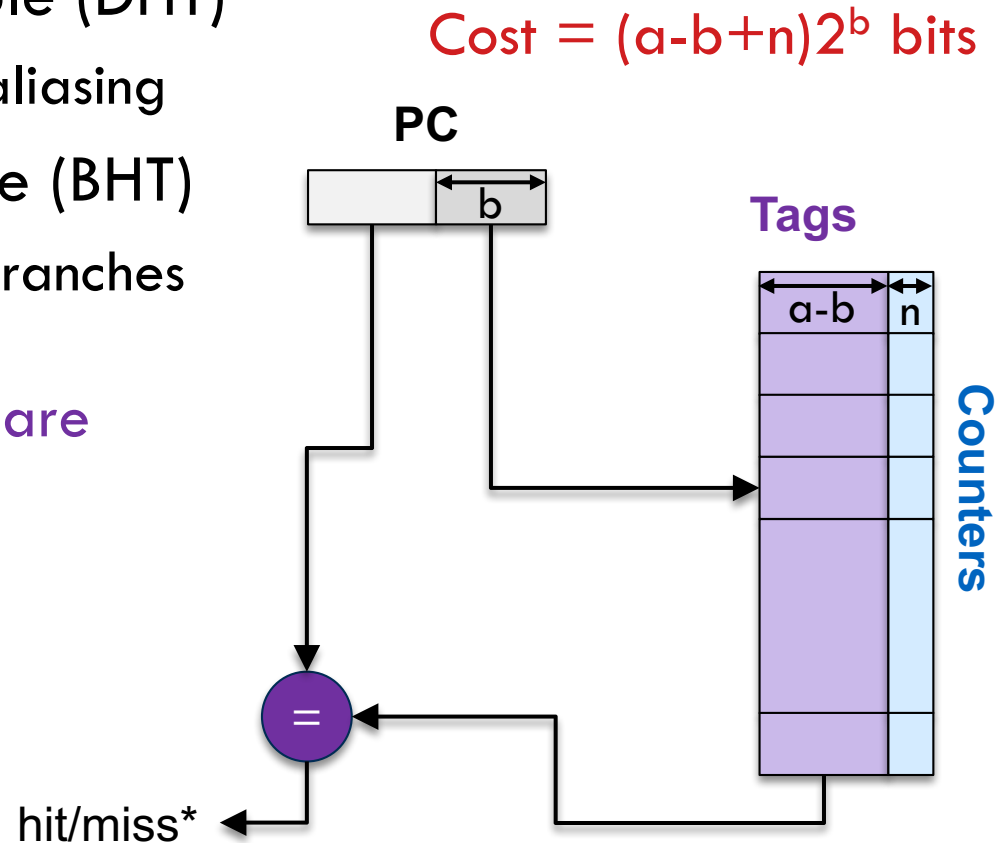
▣ Branch History Table (BHT)

■ Precisely tracking branches

Most significant bits are
used as tags

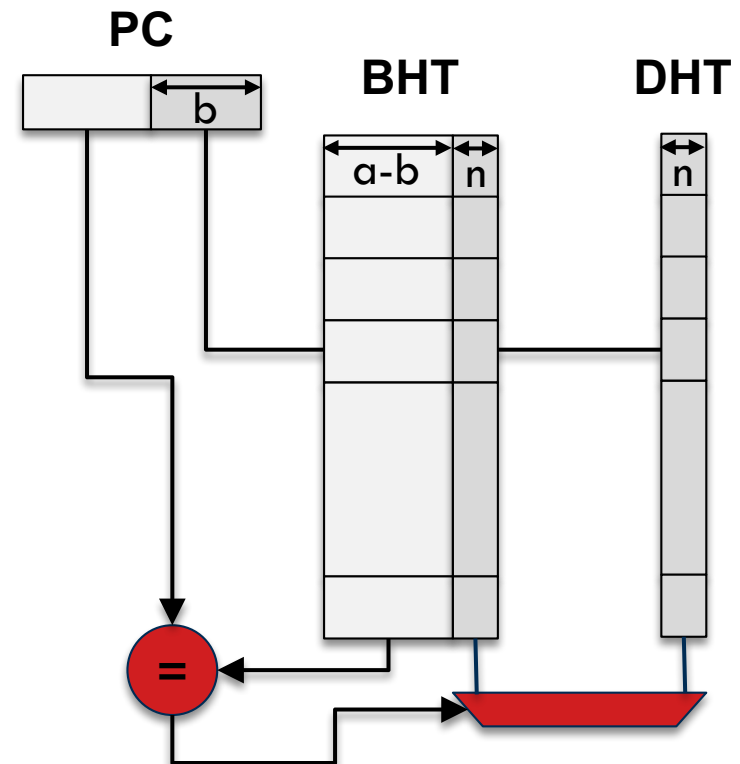
(+) No aliasing

(-) Missing entries



Using Multiple Counters

- How to assign a branch to each counter?
 - ▣ Decode History Table (DHT)
 - Reduced HW with aliasing
 - ▣ Branch History Table (BHT)
 - Precisely tracking branches
 - ▣ Combined BHT and DHT
 - BHT is used on a hit
 - DHT is used/updated on a miss



Using Multiple Counters

□ How to assign a branch to each counter?

▣ Decode History Table (DHT)

- Reduced HW with aliasing

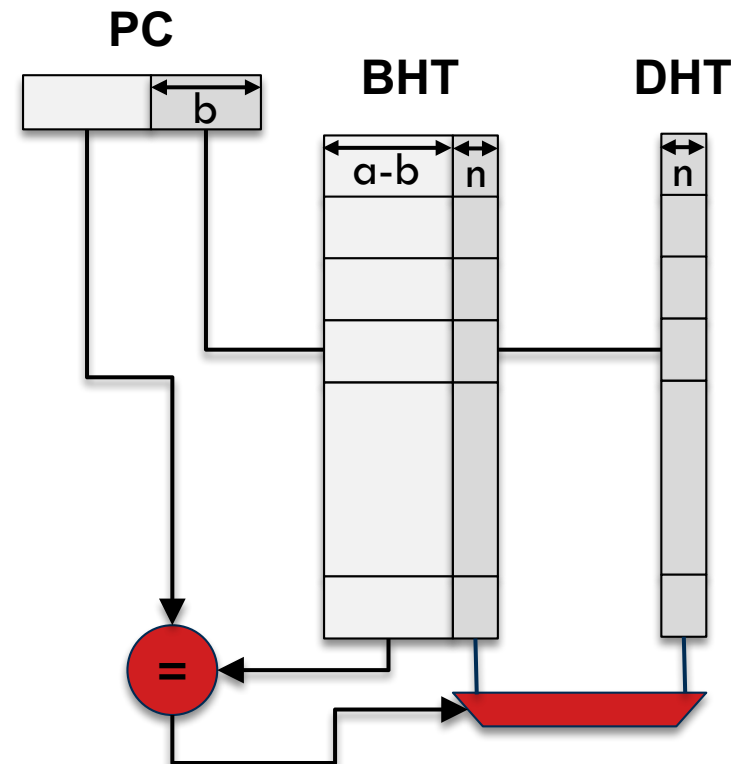
▣ Branch History Table (BHT)

- Precisely tracking branches

▣ Combined BHT and DHT

- BHT is used on a hit
- DHT is used/updated on a miss

$$\text{Cost} = (a-b+2n)2^b \text{ bits}$$



Using Multiple Counters

□ How to assign a branch to each counter?

▣ Decode History Table (DHT)

- Reduced HW with aliasing

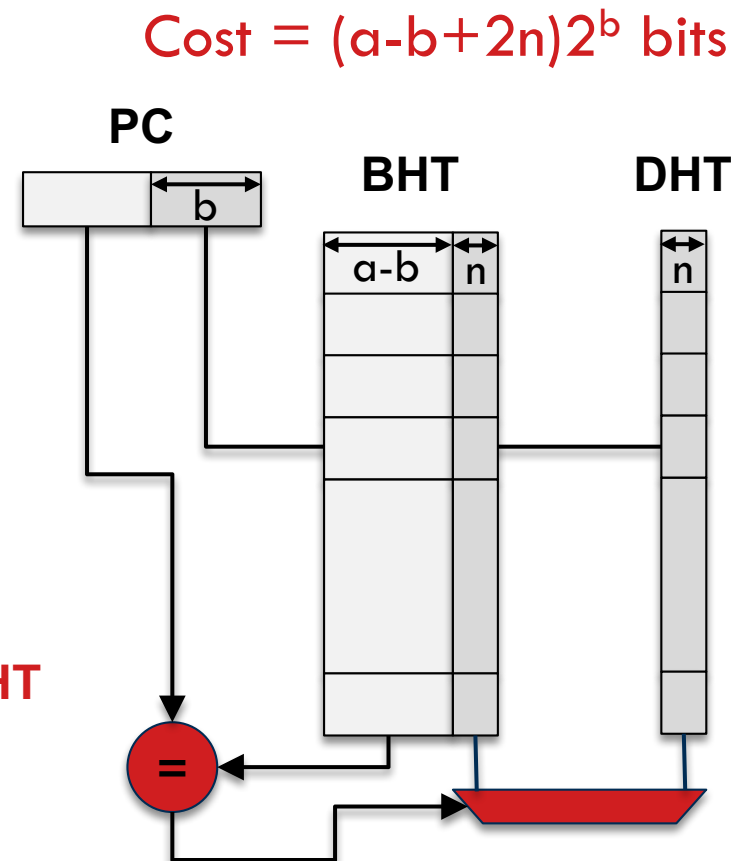
▣ Branch History Table (BHT)

- Precisely tracking branches

▣ Combined BHT and DHT

- BHT is used on a hit
- DHT is used/updated on a miss

DHT typically has more entries than BHT



Correlating Branch Predictor

- Executed branches of a program stream may be correlated

```
while (1) {  
    if(x == 0)  
        y = 0;  
    ...  
    if(y == 0)  
        x = 1;  
}
```


Correlating Branch Predictor

- Executed branches of a program stream may be correlated

```
while (1) {  
    if(x == 0)  
        y = 0;  
    ...  
    if(y == 0)  
        x = 1;  
}
```

branch-1

branch-2

Correlating Branch Predictor

- Executed branches of a program stream may be correlated

```
while (1) {  
    if(x == 0)  
        y = 0;  
    ...  
    if(y == 0)  
        x = 1;  
}
```

branch-1

branch-2

```
while:  
    BNEQ R1, R0, skp1  
    ADDI R2, R0, #0  
skp1: ...  
    BNEQ R2, R0, skp2  
    ADDI R1, R0, #1  
skp2: J while
```

Correlating Branch Predictor

- Executed branches of a program stream may be correlated

```
while (1) {  
    if(x == 0)  
        y = 0;  
    ...  
    if(y == 0)  
        x = 1;  
}
```

branch-1

branch-2

Global History Register: an r-bit shift register that maintains outcome history



Correlating Branch Predictor

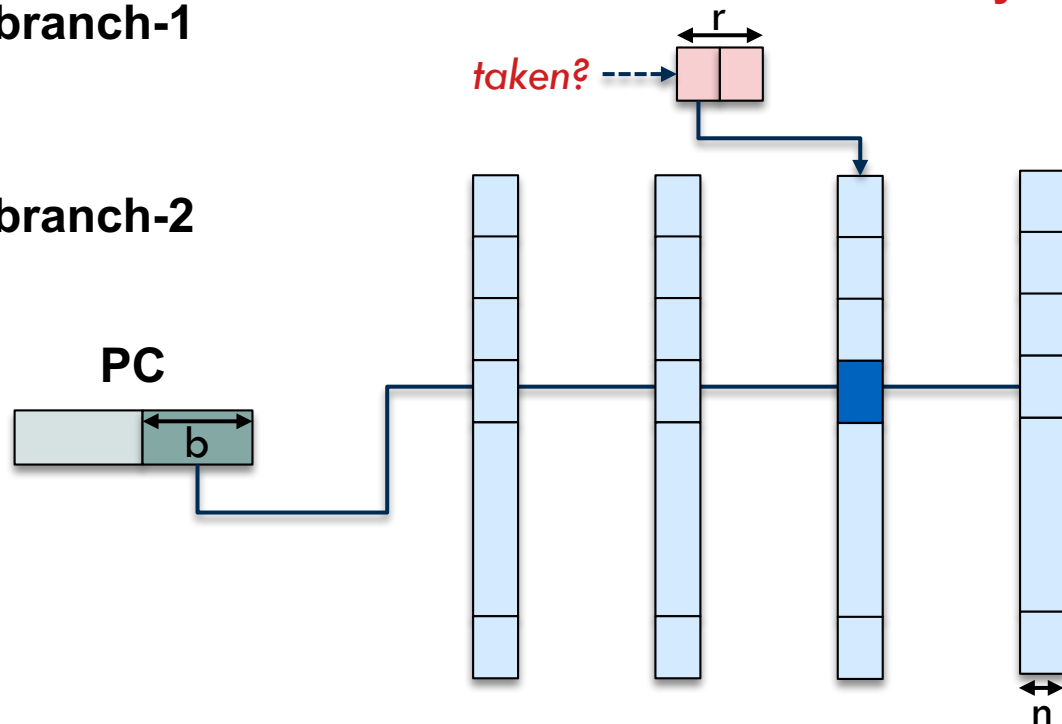
- Executed branches of a program stream may be correlated

```
while (1) {  
    if(x == 0)  
        y = 0;  
    ...  
    if(y == 0)  
        x = 1;  
}
```

branch-1

branch-2

PC



Global History Register: an r -bit shift register that maintains outcome history

Correlating Branch Predictor

- Executed branches of a program stream may be correlated

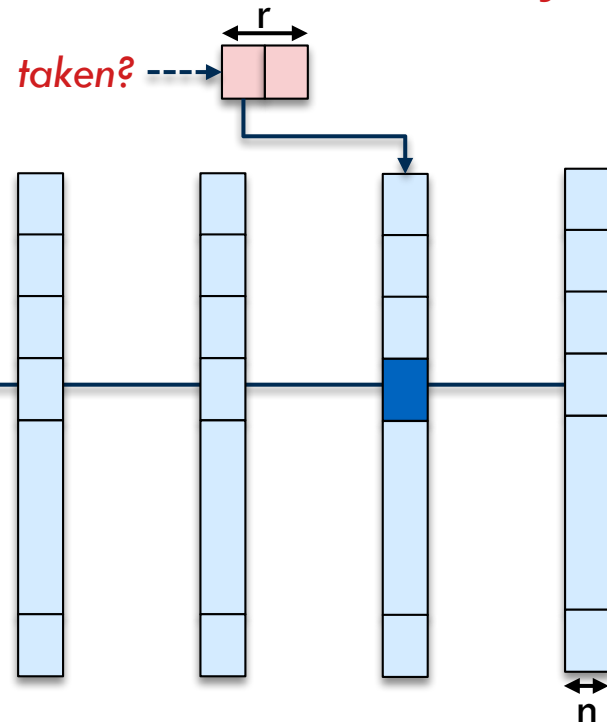
```
while (1) {  
    if(x == 0)  
        y = 0;  
    ...  
    if(y == 0)  
        x = 1;  
}
```

branch-1

branch-2

PC

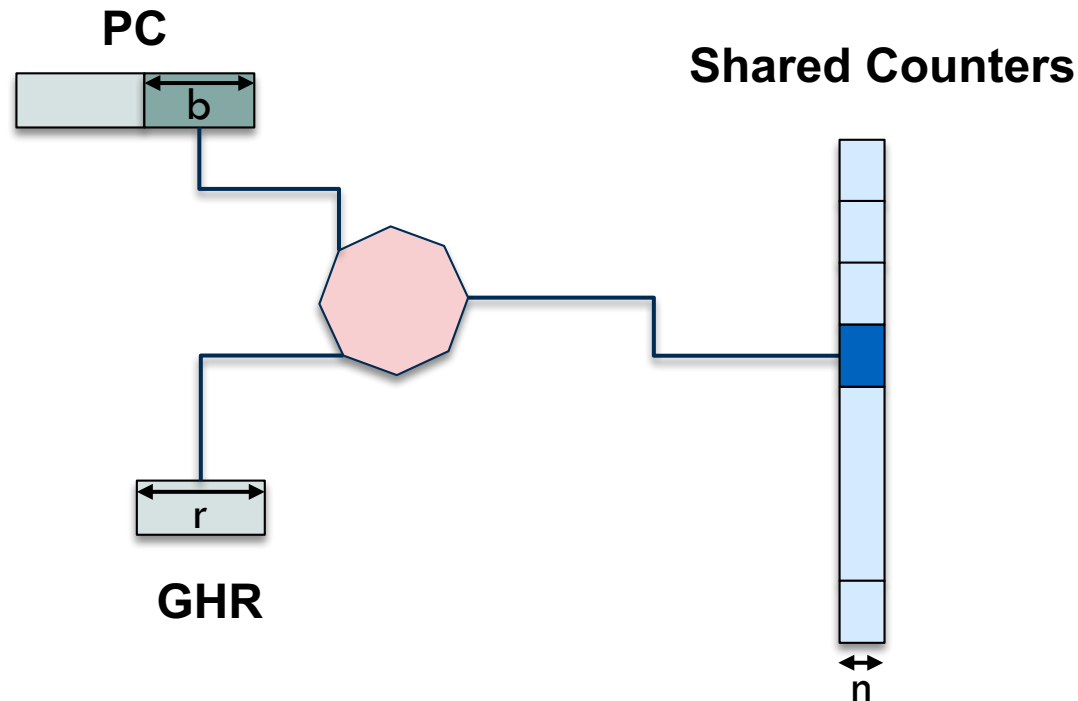
Global History Register: an r-bit shift register that maintains outcome history



Cost = $r + n2^{b+r}$ bits

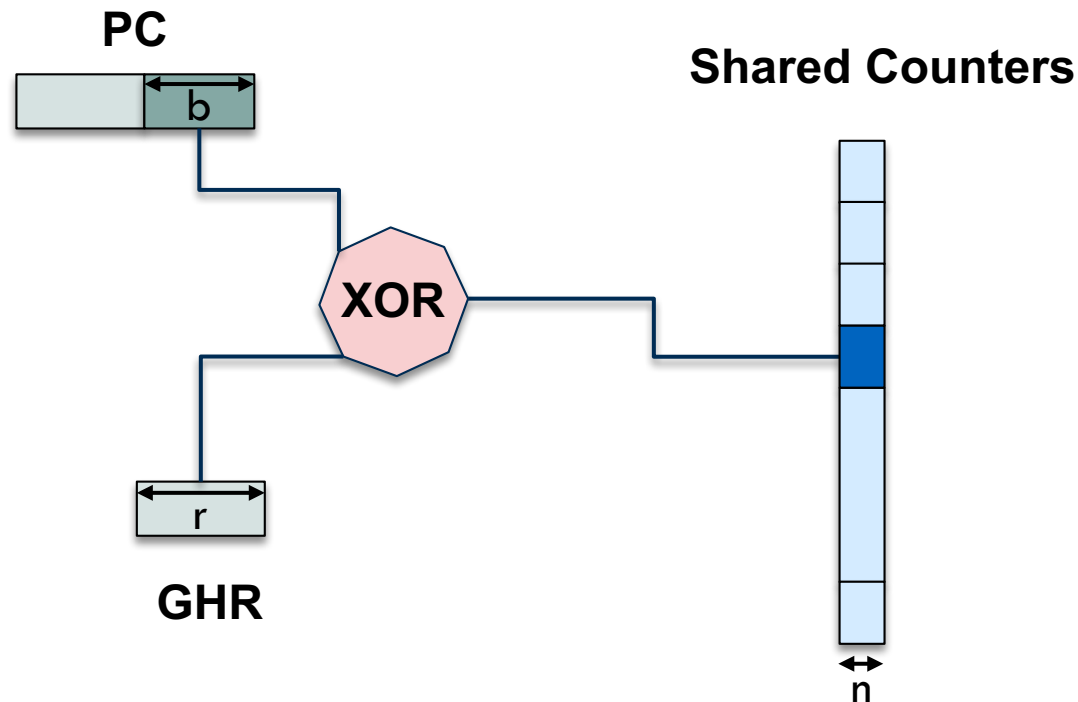
Global Branch Predictor

- GHR is merged with PC bits to choose a counter



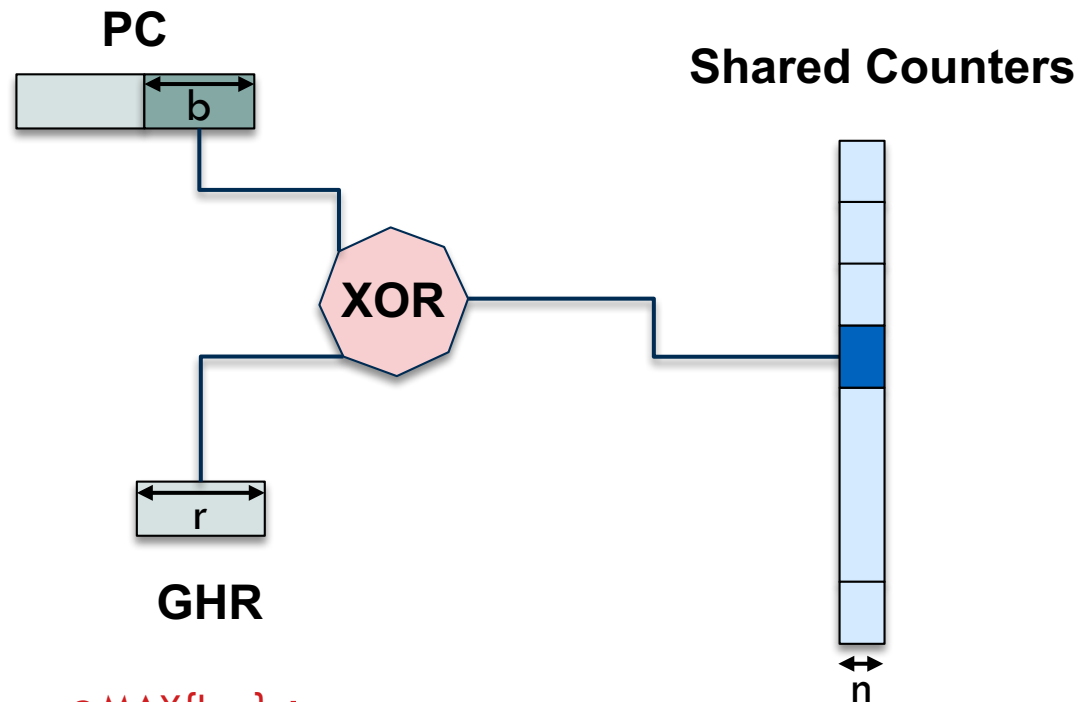
Global Branch Predictor

- GHR is merged with PC bits to choose a counter



Global Branch Predictor

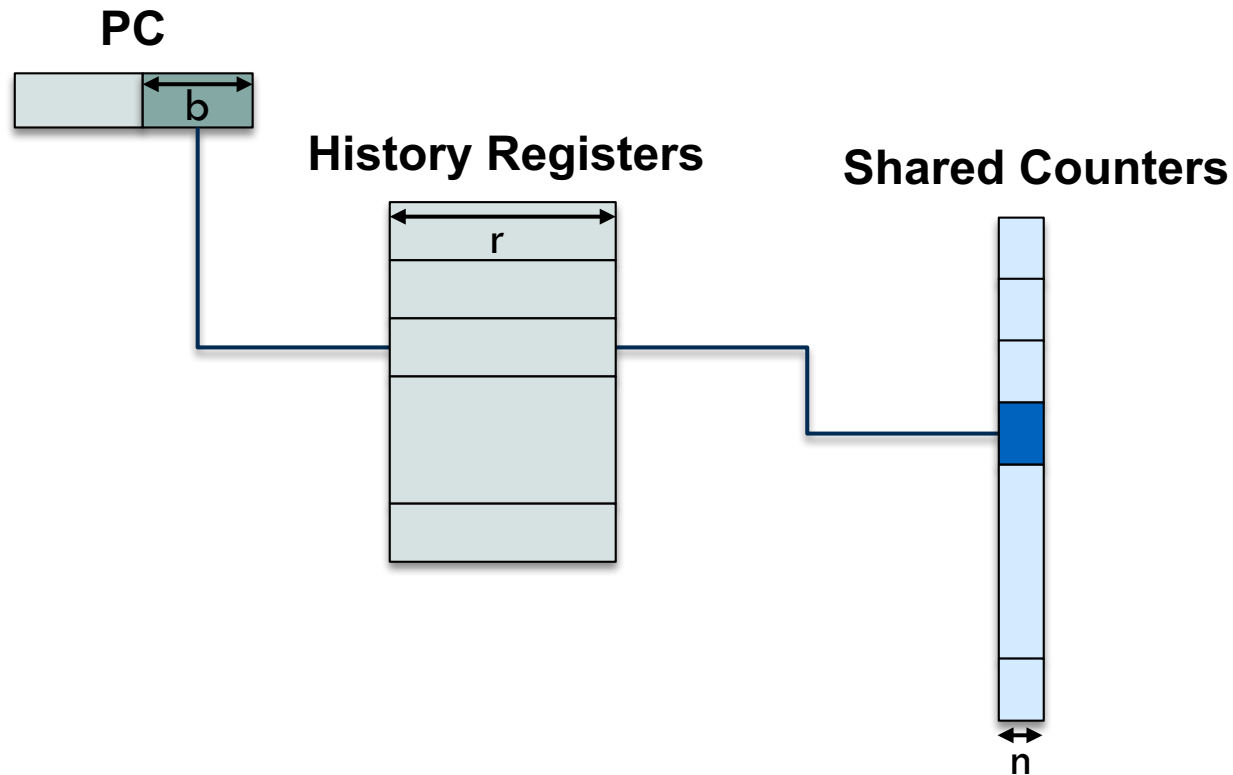
- GHR is merged with PC bits to choose a counter



$$\text{Cost} = r + n2^{\text{MAX}\{b, r\}} \text{ bits}$$

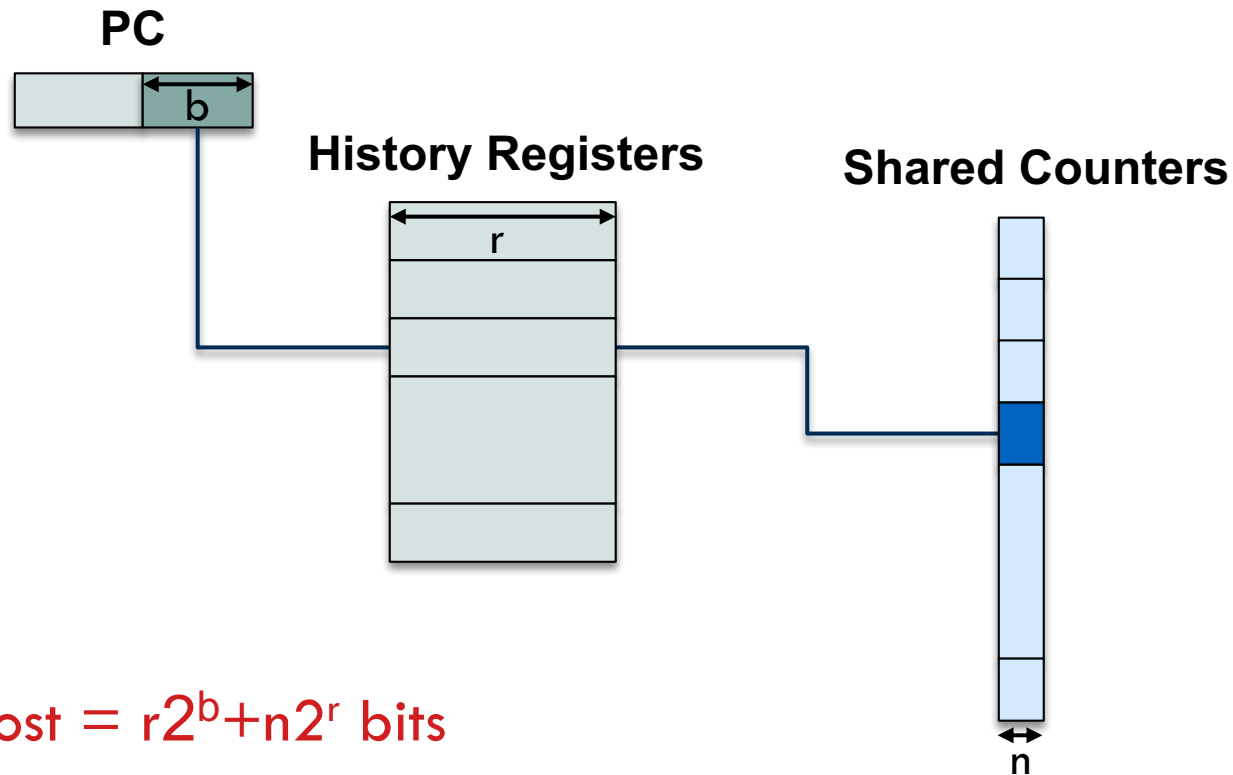
Local Branch Predictor

- One GHR per branch



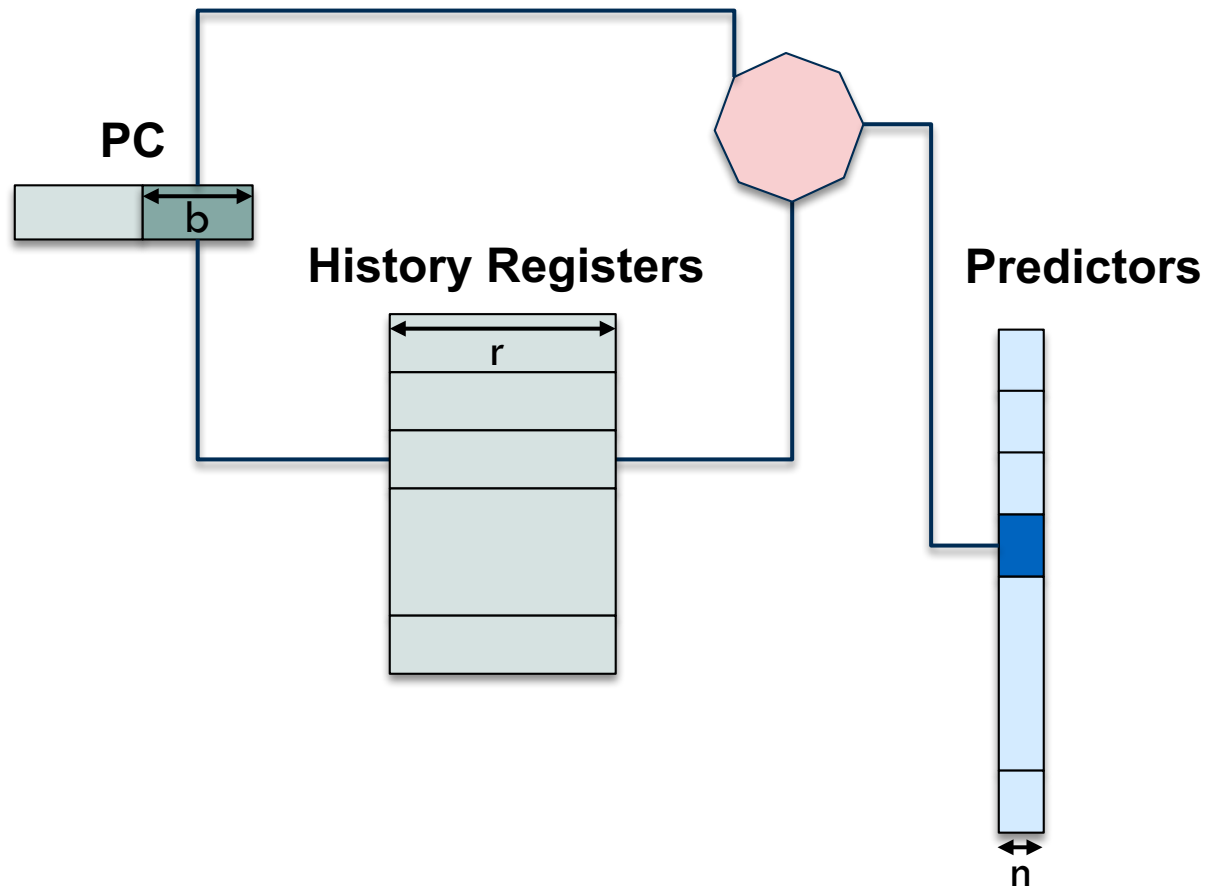
Local Branch Predictor

- One GHR per branch



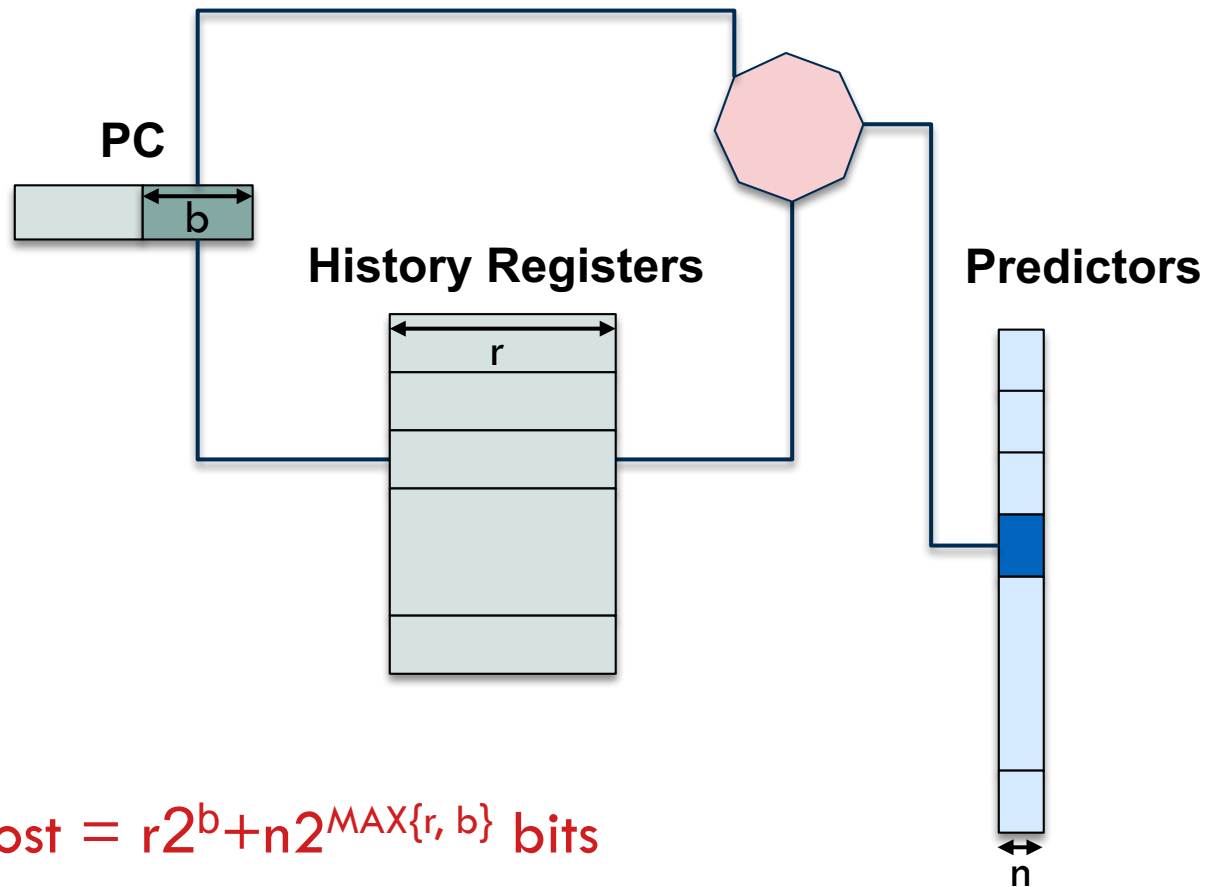
Local Branch Predictor

- One GHR per branch



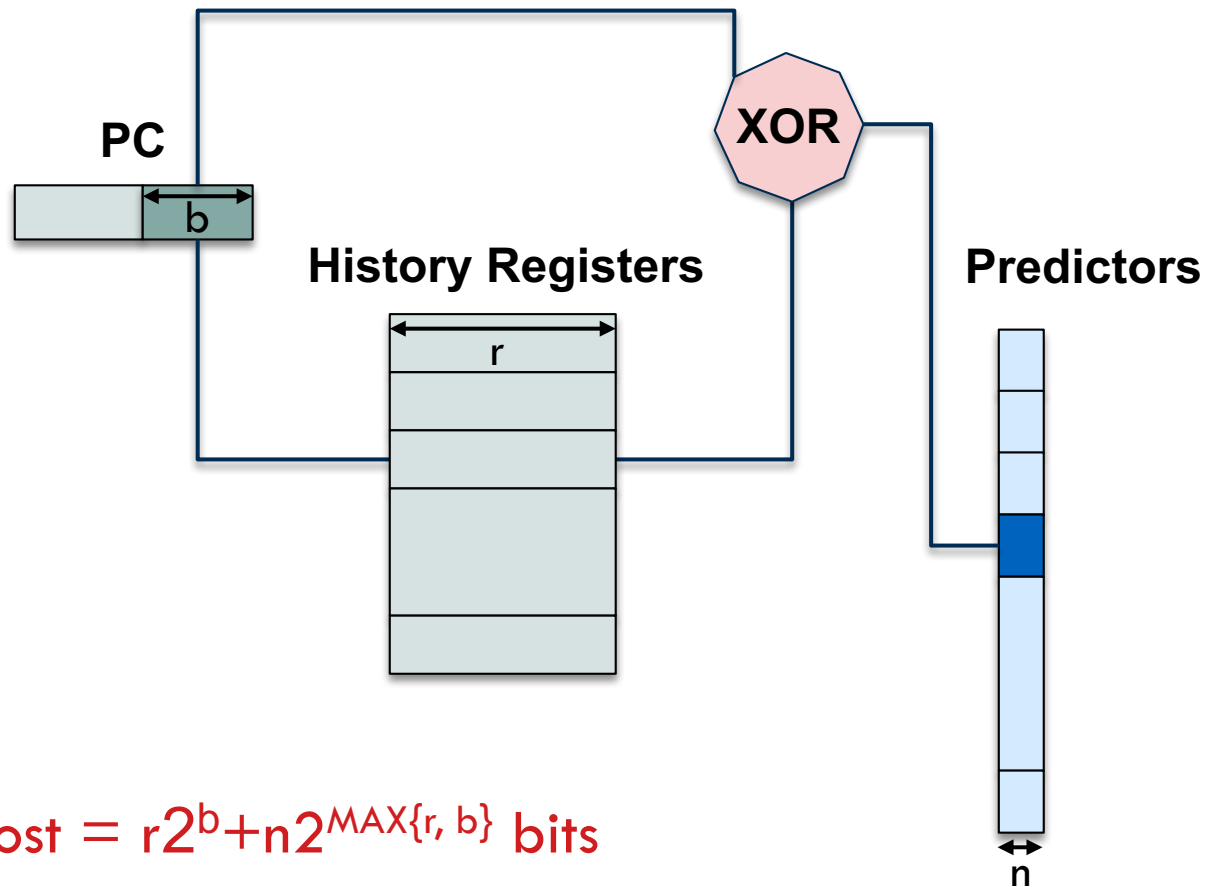
Local Branch Predictor

- One GHR per branch



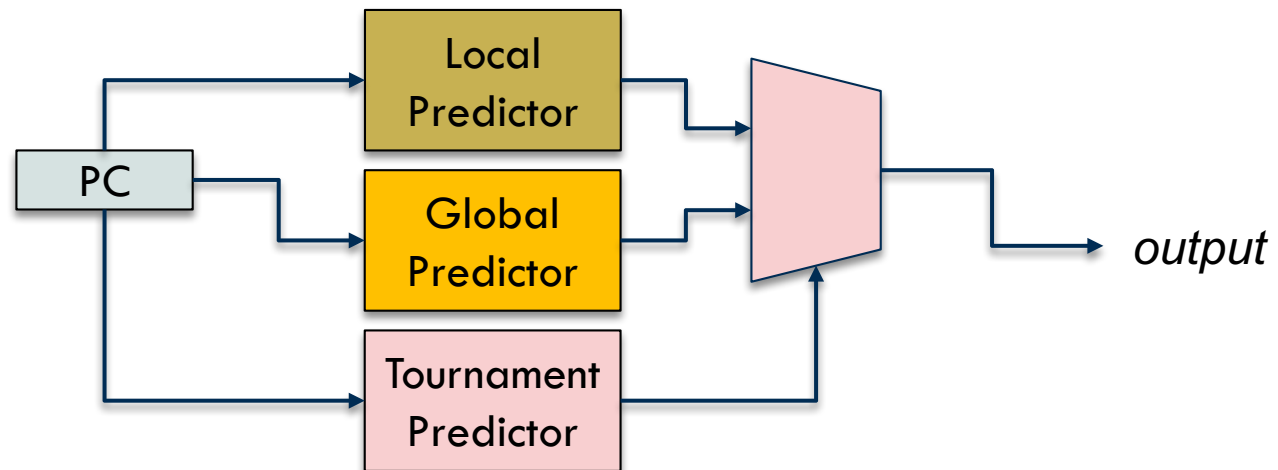
Local Branch Predictor

- One GHR per branch



Tournament Branch Predictor

- Local predictor may work well for some applications, while global predictor works well for some other programs
- ▣ Include both and identify/use the best one for each branch



Two bit saturating counters

Branch Prediction Summary

- Dedicated predictor per branch
 - ▣ Program counter is used for assigning predictors to branches
- Capturing correlation among branches
 - ▣ Shift register is used to track history
- Predicting branch direction is not enough
 - ▣ Which instruction to be fetched if taken?
- Storing the target instruction can eliminate fetching
 - ▣ Extra hardware is required

Branch Target Buffer

- Store tags and target addresses for each branch

