

# CACHE INTERCONNECTS

Mahdi Nazm Bojnordi

Assistant Professor

School of Computing

University of Utah


# Overview

- Upcoming deadline
  - ▣ Feb. 1<sup>st</sup>: project group formation
  - ▣ Note: email me once you form a group
- This lecture
  - ▣ Content aware optimizations
  - ▣ Cache interconnect optimizations
  - ▣ Encoding based optimizations

# Recall: Cache Power Optimization

- ❑ Caches are power and performance critical components
- ❑ Performance
  - ▣ Bridging the CPU-Mem gap
- ❑ Static power
  - ▣ Large number of leaky cells
- ❑ Dynamic power
  - ▣ Access through long interconnects

## Example: FX Processors




[source: AMD]

# Content Aware Optimizations

# Dynamic Zero Compression

- More than 70% of the bits in data cache accesses are 0s




**Example of a small cache**

|                             | Read |       | Write |       |
|-----------------------------|------|-------|-------|-------|
|                             | (pJ) | (%)   | (pJ)  | (%)   |
| Total                       | 44.4 | 100.0 | 99.1  | 100.0 |
| Decoder                     | 5.5  | 12.4  | 5.5   | 5.5   |
| word lines                  | 1.1  | 2.5   | 1.1   | 1.1   |
| Tag bitlines and sense-amp  | 3.0  | 6.2   | 3.0   | 3.0   |
| Data bitlines and sense-amp | 14.5 | 32.7  | 69.2  | 69.9  |
| I/O buses                   | 12.1 | 27.3  | 12.1  | 12.2  |
| Other                       | 8.4  | 18.9  | 8.4   | 8.5   |

[Villa'00]


# Dynamic Zero Compression

- Zero Indicator Bit; one bit per grouping of bits; set if bits are zeros; controls wordline gating




# Dynamic Zero Compression

## □ Data cache bitline swing reduction



# Dynamic Zero Compression

## □ Data cache energy savings




[Villa'00]



# Cache Interconnect Optimizations


# Large Cache Organization

- Fewer subarrays gives increased area efficiency, but larger delay due to longer wordlines/bitlines




# Large Cache Energy Consumption

- H-tree is clearly the dominant component of energy consumption




# Heterogeneous Interconnects


- A global wire management at the microarchitecture level
- A heterogeneous interconnect that is comprised of wires with varying **latency**, **bandwidth**, and **energy** characteristics




Delay Optimized



Bandwidth Optimized



Power Optimized




Power and Bandwidth Optimized

# Heterogeneous Interconnects


- Better energy-efficiency for a dynamically scheduled partitioned architecture
  - ▣  $ED^2$  is reduced by 11%
- A low-latency low-bandwidth network can be effectively used to hide wire latencies and improve performance
- A high-bandwidth low-energy network and an instruction assignment heuristic are effective at reducing contention cycles and total processor energy.

# Non-Uniform Cache Architecture

- NUCA optimizes energy and time based on the proximity of the cache blocks to the cache controller.



Bank Access time = 3 cycles  
Interconnect delay = 8 cycles




Bank Access time = 3 cycles  
Interconnect delay = 44 cycles

# Non-Uniform Cache Architecture

## □ S-NUCA-1

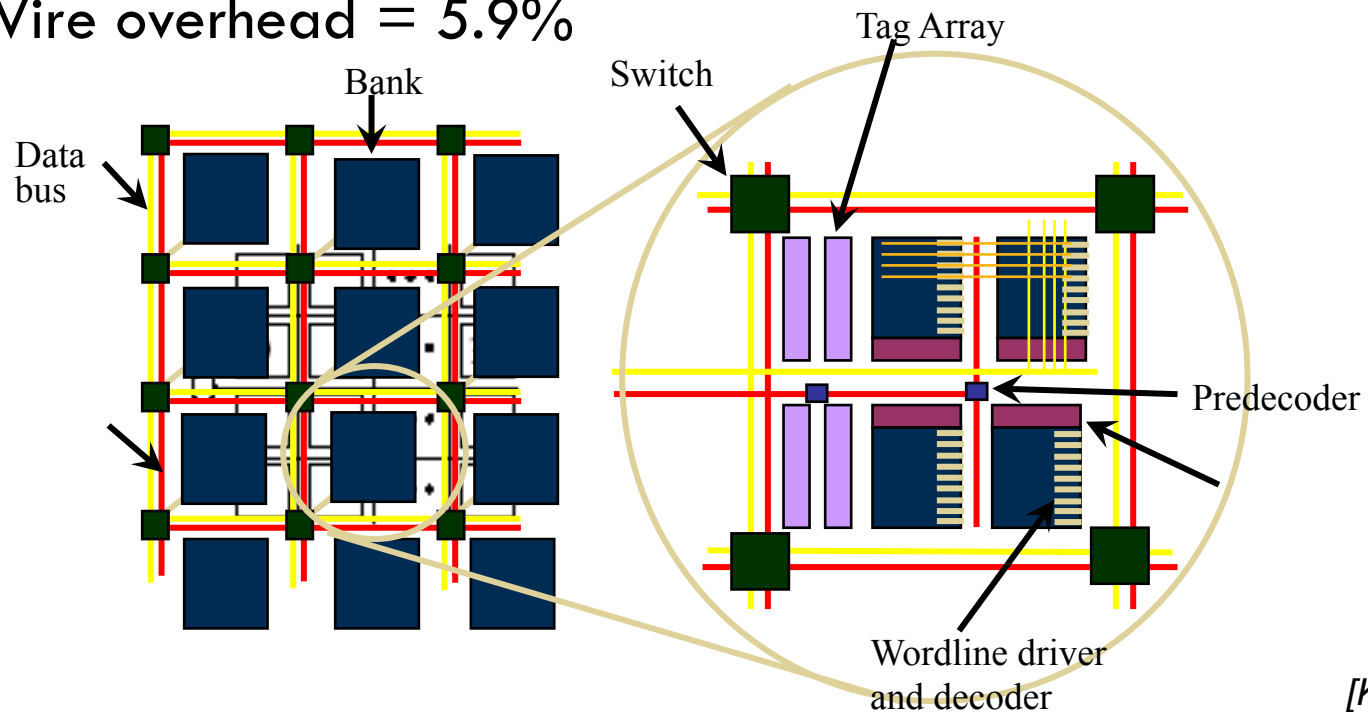
- Use private per-bank channel
- Each bank has its distinct access latency
- Statically decide data location for its given address
- Average access latency = 34.2 cycles
- Wire overhead = 20.9% → an issue



# Non-Uniform Cache Architecture

## □ S-NUCA-2

- Use a 2D switched network to alleviate wire area overhead
- Average access latency = 24.2 cycles
- Wire overhead = 5.9%




[Kim'04]




# Non-Uniform Cache Architecture

- Dynamic NUCA
  - ▣ Data can dynamically migrate
  - ▣ Move frequently used cache lines closer to CPU




# Non-Uniform Cache Architecture

- Fair mapping
  - ▣ Average access time across all bank sets are equal



# Non-Uniform Cache Architecture

- Shared mapping
  - ▣ Sharing the closet banks for farther banks




# Encoding Based Optimizations

# Cache Interconnect Optimizations


- Bus invert coding transfers either the data or its complement to minimize the number of bit flips on the bus.

$$P_{\text{switching}} = \alpha C V_{DD}^2 f$$



# Time-Based Data Transfer


- The percentage of processor energy expended on an 8MB cache when running a set of parallel applications on a Sun Niagara-like multicore processor



[Bojnordi'13]

# Time-Based Data Transfer

- Communication over the long, capacitive H-tree interconnect is the dominant source of energy consumption (80% on average) in the L2 cache




[Bojnordi'13]

# Time-Based Data Transfer

**Key idea:** represent information by the number of clock cycles between two consecutive pulses to reduce interconnect activity factor.

**Example: transmitting the value 5**



[Bojnordi'13]




# Time-Based Data Transfer

- Cache blocks are partitioned into small, contiguous **chunks**.

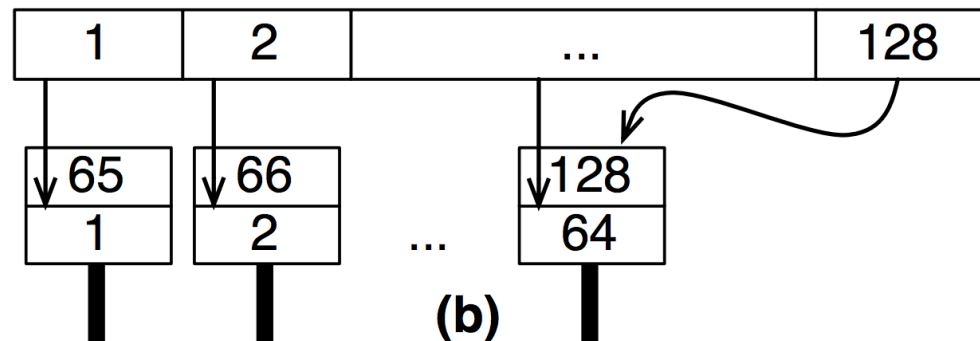
Cache Block  
Partitioned into Chunks

Communication Wires  
and FIFO Queues




Cache Block  
Partitioned into Chunks

Communication Wires  
and FIFO Queues




# Time-Based Data Transfer

## a) Transmitter



Wire Delay


## b) Receiver



[Bojnordi'13]

# Time-Based Data Transfer

- L2 cache energy is reduced by 1.8x at the cost of less than 2% increase in the execution time.



[Bojnordi'13]