

DIRECTORY COHERENCE

Mahdi Nazm Bojnordi

Assistant Professor

School of Computing

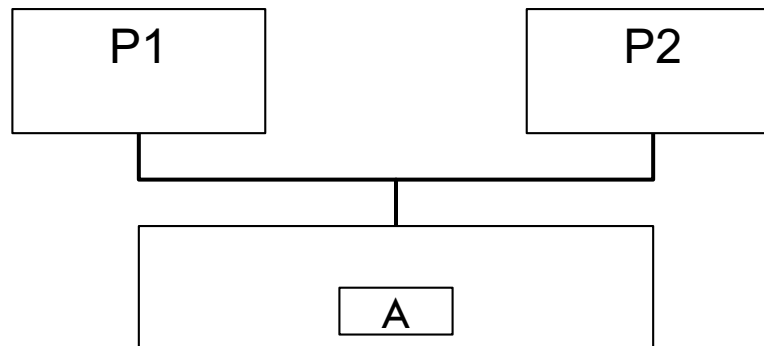
University of Utah

Overview

- Upcoming deadline
 - ▣ Feb. 8th: project proposal
- This lecture
 - ▣ Snooping wrap-up
 - ▣ Directory coherence
 - ▣ Implementation challenges
 - ▣ Token-based coherence protocol

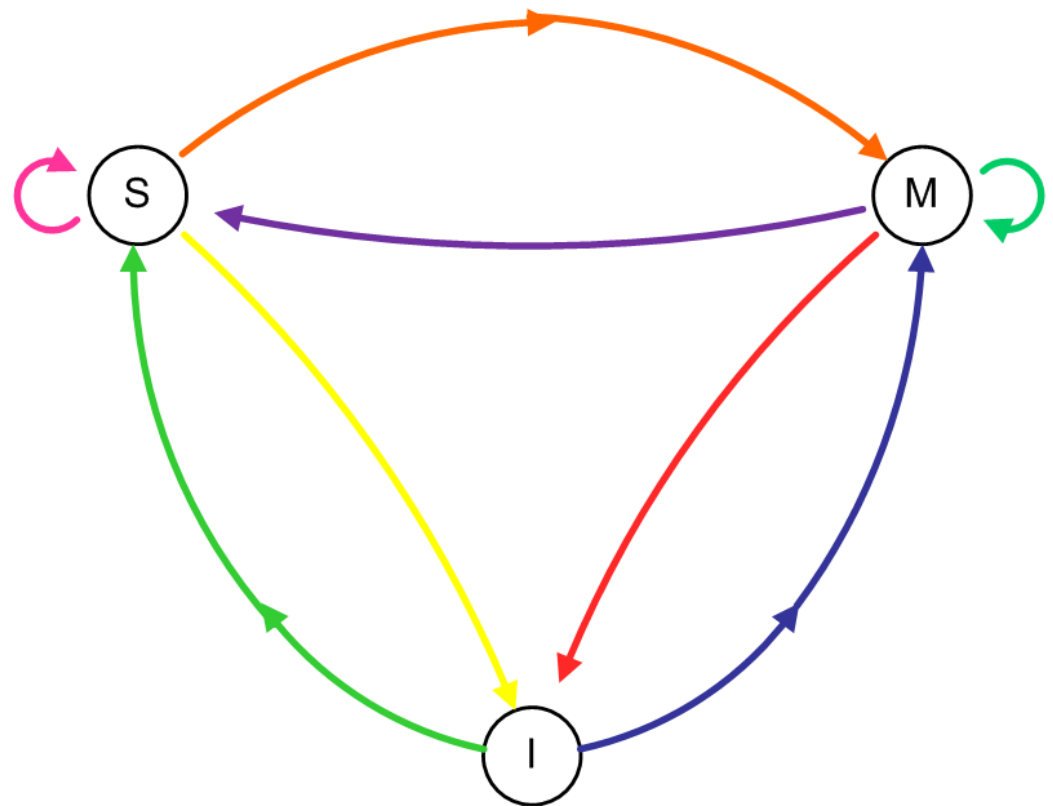
Recall: Cache Coherence

- Definition of coherence
 - ▣ Write propagation
 - Write is visible to other processors
 - ▣ Write serialization
 - All writes to the same location are seen in the same order by all processes



Implementation Challenges

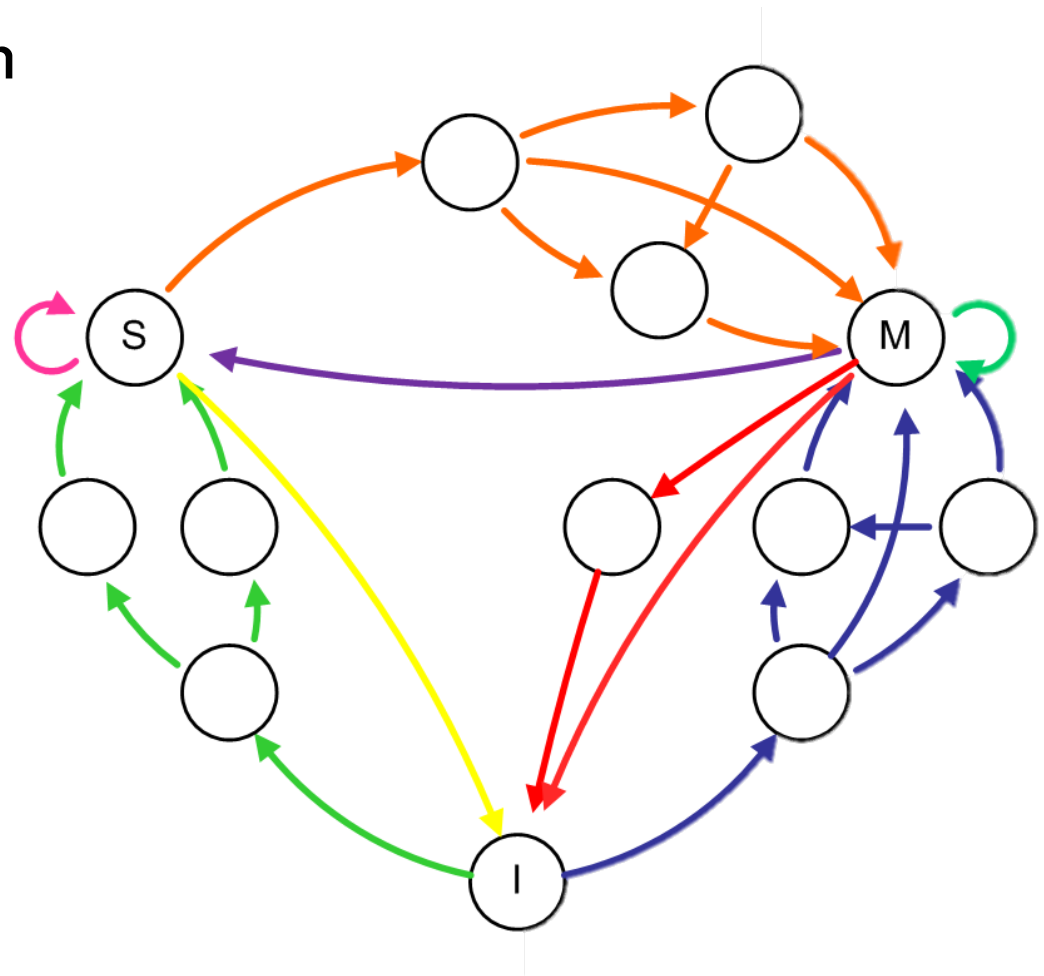
- MSI implementation
 - ▣ Stable States



[Vantrease'11]

Implementation Challenges

- MSI implementation
 - ▣ Stable States
 - ▣ Busy states



[Vantrease'11]

Implementation Challenges

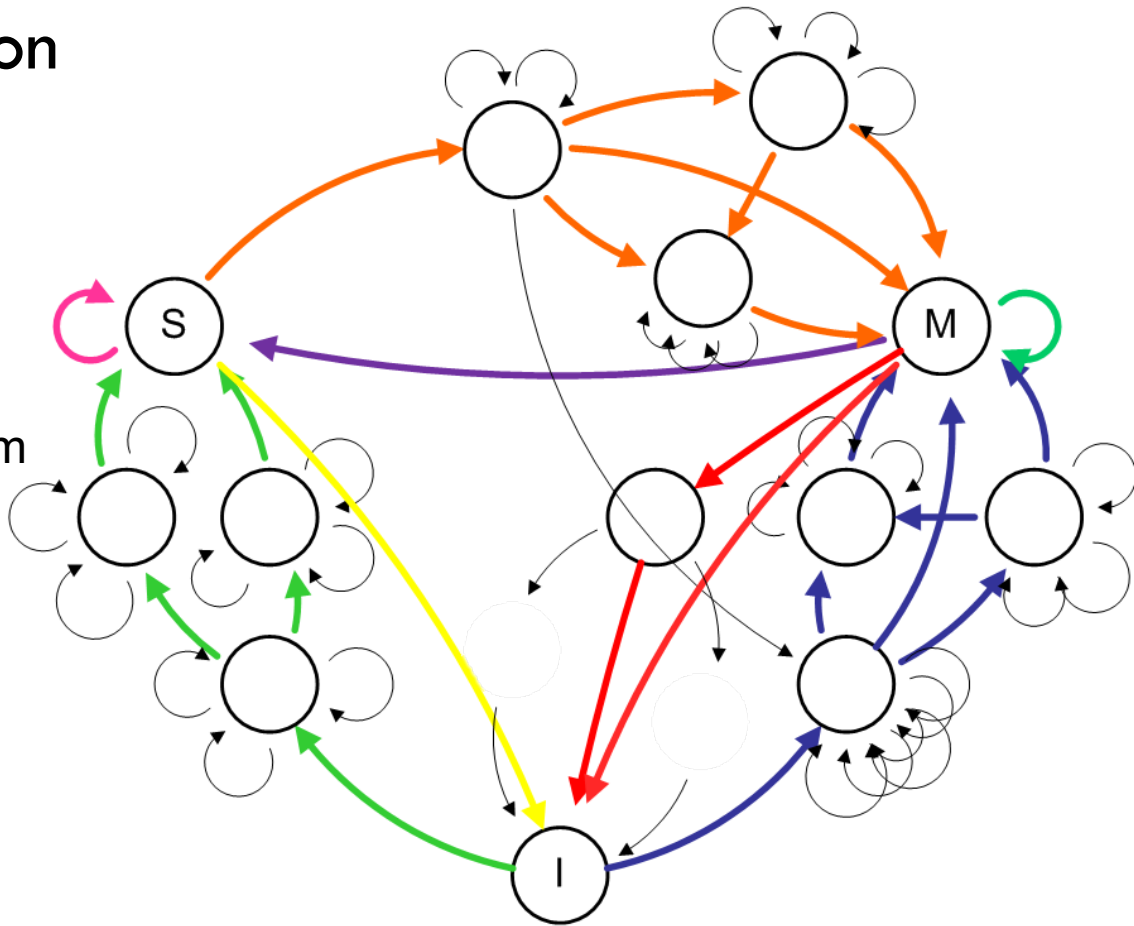
□ MSI implementation

▣ Stable States

▣ Busy states

▣ Races

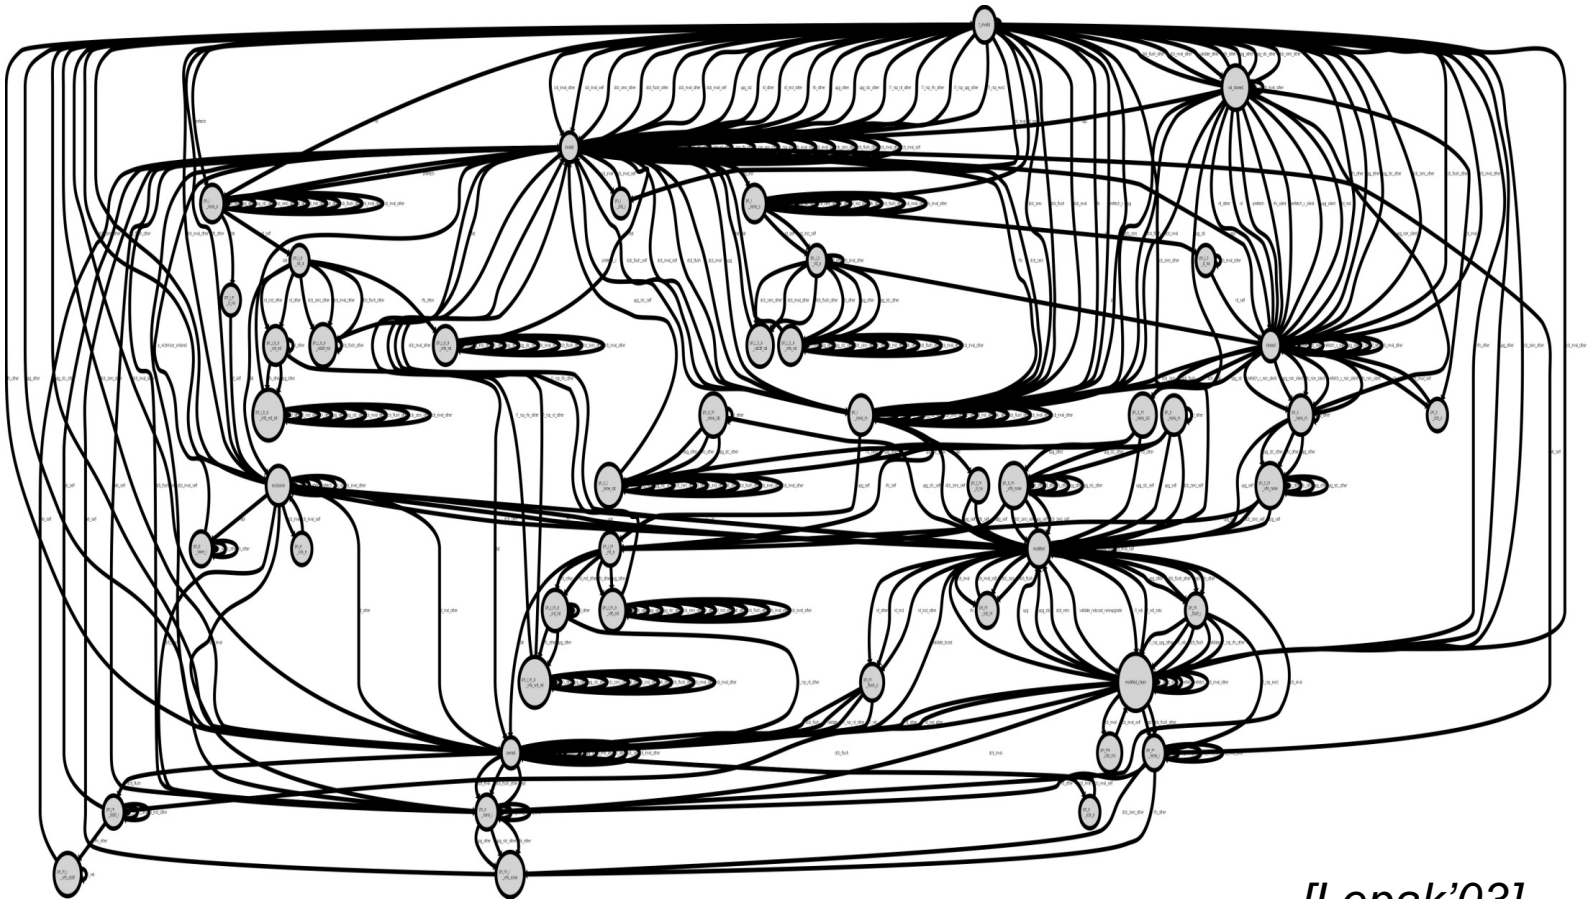
Unexpected events from
concurrent requests to
same block



[Vantrease'11]

Cache Coherence Complexity

- A broadcast snooping bus (L2 MOETSI)



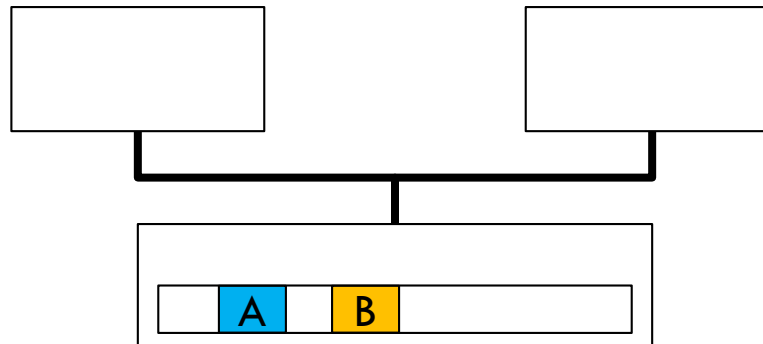
[Lepak'03]

Implementation Tradeoffs

- Reduce unnecessary invalidates and transfers of blocks
 - Optimize the protocol with more states and prediction mechanisms
- Adding more states and optimizations
 - Difficult to design and verify
 - lead to more cases to take care of
 - race conditions
 - Gained benefit may be less than costs (diminishing returns)

Coherence Cache Miss

- **Recall:** cache miss classification
 - ▣ Cold (compulsory): first access to block
 - ▣ Capacity: due to limited capacity
 - ▣ Conflict: many blocks are mapped to the same set
- **New class:** misses due to sharing
 - ▣ True vs. false sharing



Summary of Snooping Protocols

□ Advantages

- Short miss latency
- Shared bus provides global point of serialization
- Simple implementation based on buses in uniprocessors

□ Disadvantages

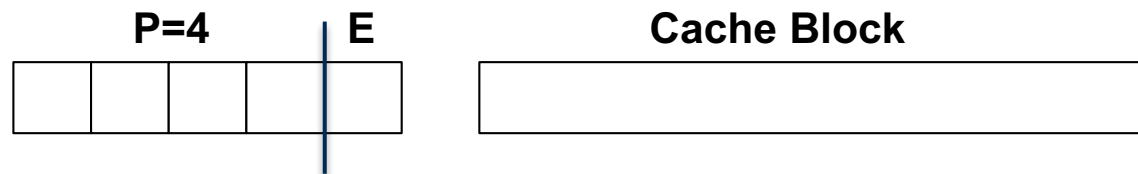
- Must broadcast messages to preserve the order
- The global point of serialization is not scalable
 - It needs a virtual bus (or a totally-ordered interconnect)

Scalable Coherence Protocols

- **Problem:** shared interconnect is not scalable
- **Solution:** make explicit requests for blocks
- Directory-based coherence: every cache block has additional information
 - ▣ To track of copies of cached blocks and their states
 - ▣ To track ownership for each block
 - ▣ To coordinate invalidation appropriately

Directory Information

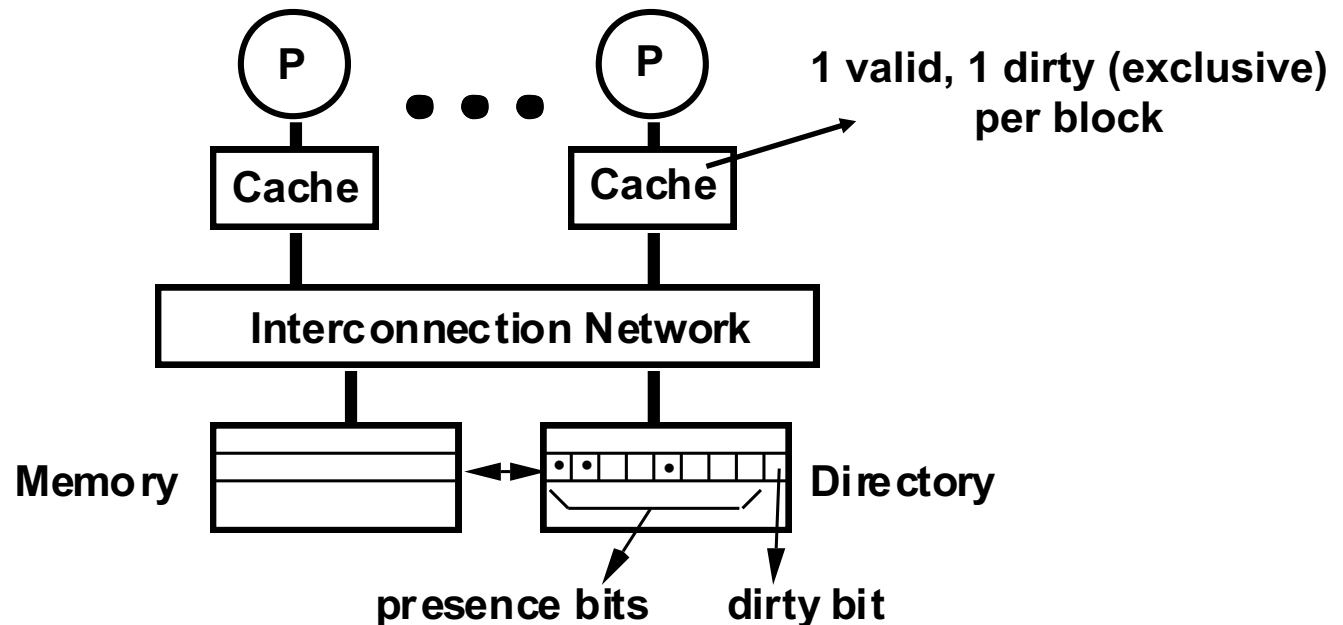
- $P+1$ additional bits for every cache block
 - ▣ One bit used to indicate the block is in each cache
 - ▣ One exclusive bit to indicate the cache has the only copy (can update without notifying others)
- On a read, set the cache's bit and arrange the supply of data
- On a write, invalidate all caches that have the block and reset their bits



How to organize directory information?

Directory Organization

- Example: central directory for P processors
 - For each cache block in memory
 - p presence bits, 1 dirty bit
 - For each cache block in cache
 - 1 valid bit, 1 dirty (owner) bit



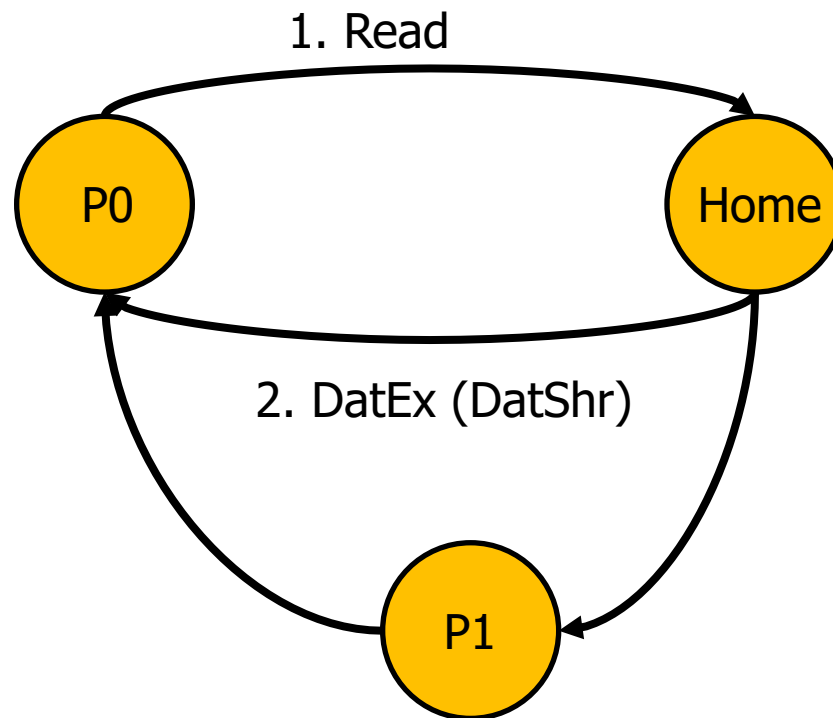
Directory Protocol

- Three states (similar to snoopy protocol)
 - ▣ **Shared:** more than one processors have data, memory up-to-date
 - ▣ **Uncached:** no processor has it; not valid in any cache
 - ▣ **Exclusive:** one processor has data; memory out-of-date

- Basic terminology
 - ▣ **Local node**, where a request originates
 - ▣ **Home node**, where the memory location of an address resides
 - ▣ **Remote node**, has copy of a cache block, whether exclusive or shared

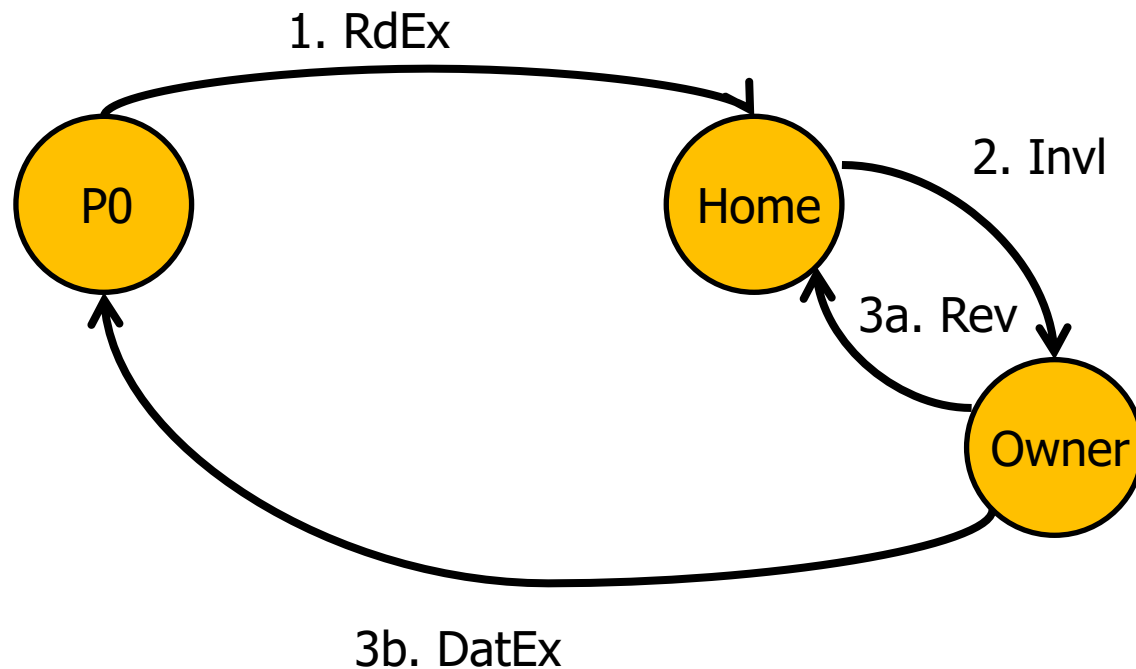
Read Request

- P0 reads a cache location



ReadEx Request

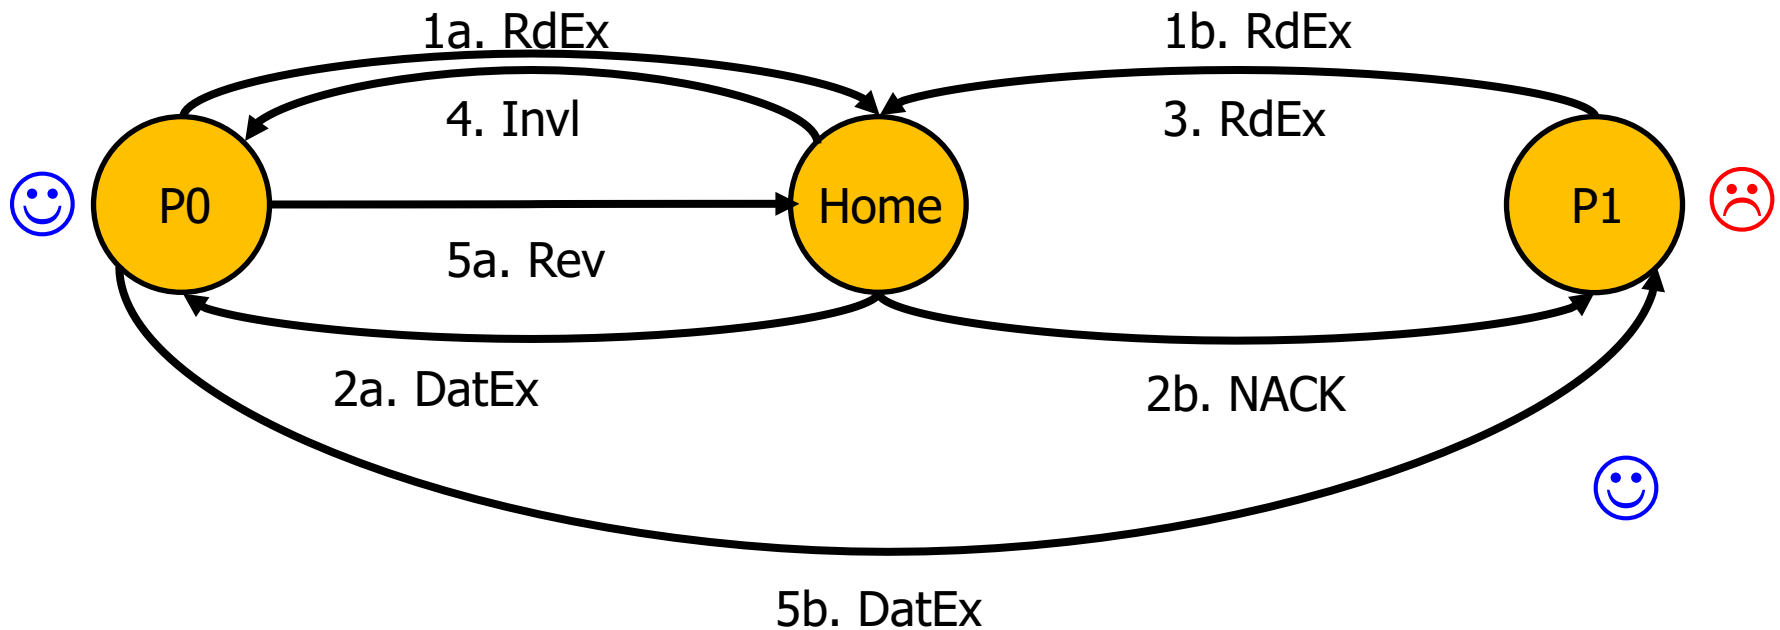
- Avoid roundtrip to home by sending data directly from owner



[Culler/Singh]

Write Contention

□ NACKing mechanism



What are the challenges?

[Culler/Singh]

Design Challenges

- Fairness: which requester is preferred on a conflict?
 - ▣ Consider distance and delivery order of interconnect

- Race condition: how to keep the proper sequence
 - ▣ NACK requests to busy blocks (pending invalidate)
 - Original requestor retries
 - ▣ Queuing requests and granting in sequence

Summary of Directory Protocols

□ Advantages

- Does not require broadcast to all caches
- Exactly as scalable as interconnect and directory storage (much more scalable than bus)

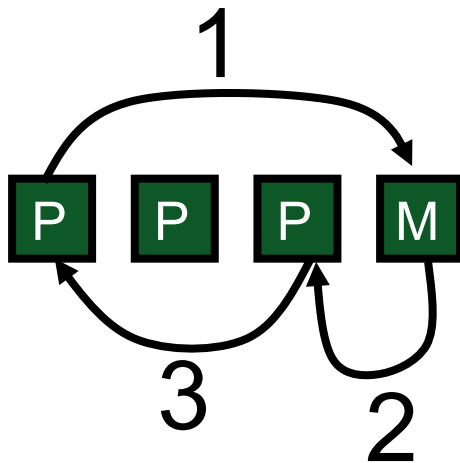
□ Disadvantages

- Adds **indirection** to miss latency (critical path)
 - request → directory → memory
- Requires extra storage space to track directory states
- Protocols and race conditions are more complex

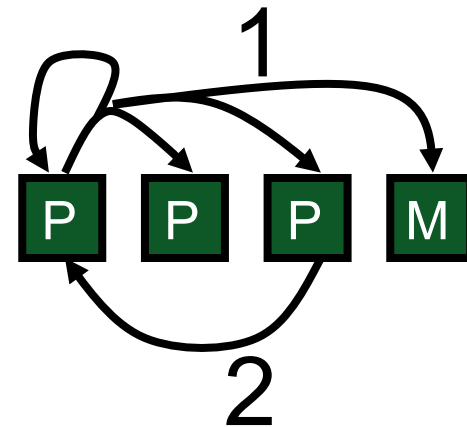
Avoid Indirection

- Can we get the best of both snooping and directory protocols?
 - ▣ Direct cache-to-cache misses (broadcast is ok)
 - ▣ What if unordered interconnect (e.g., mesh) was used?

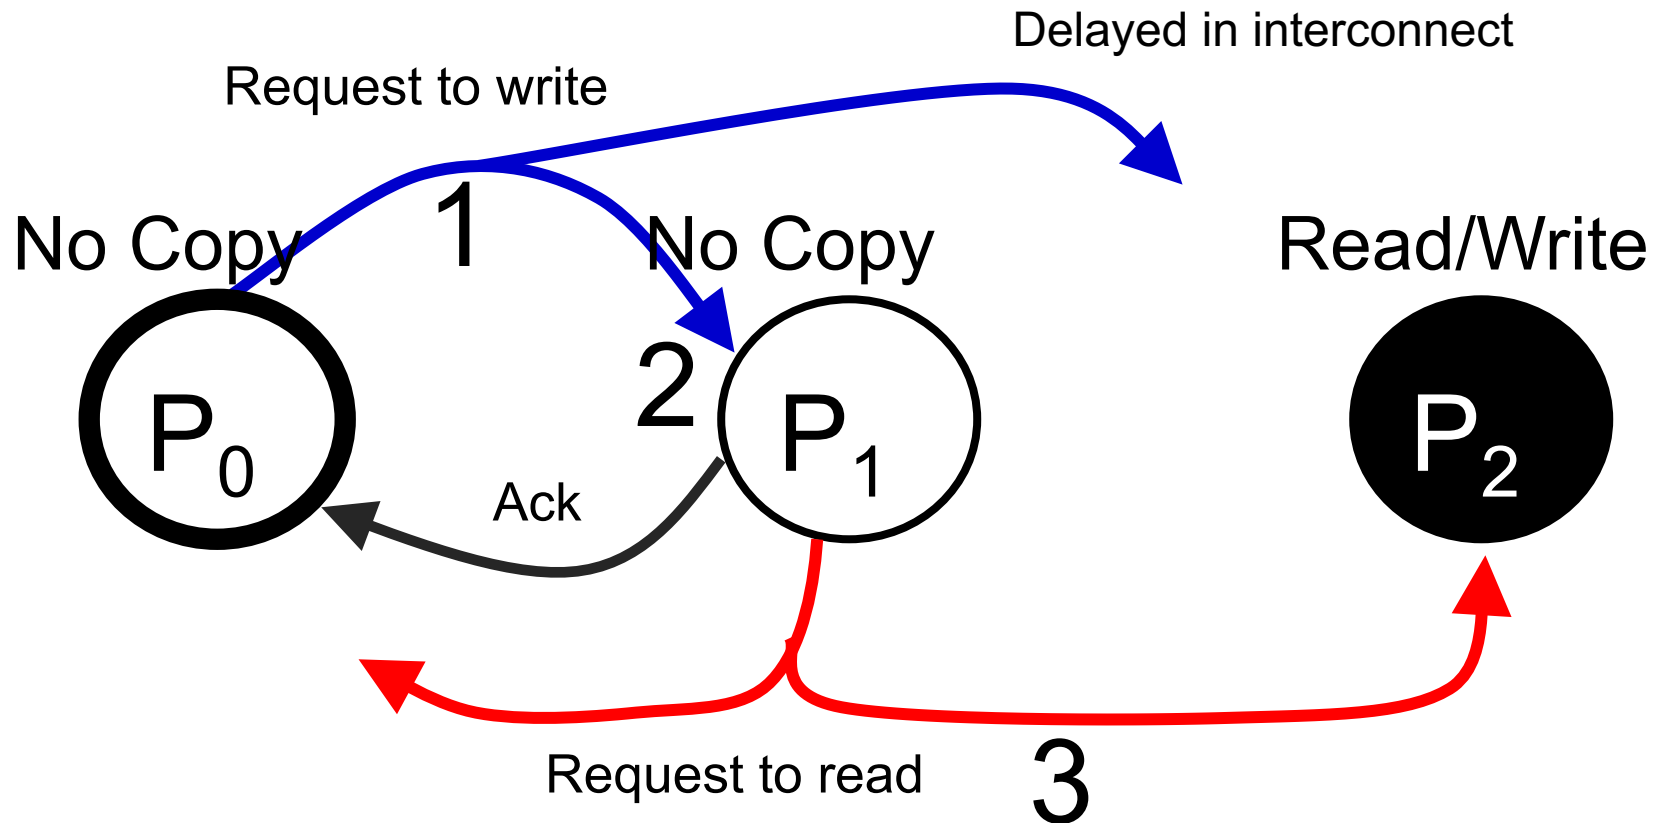
Directory Protocol



Hybrid Protocol

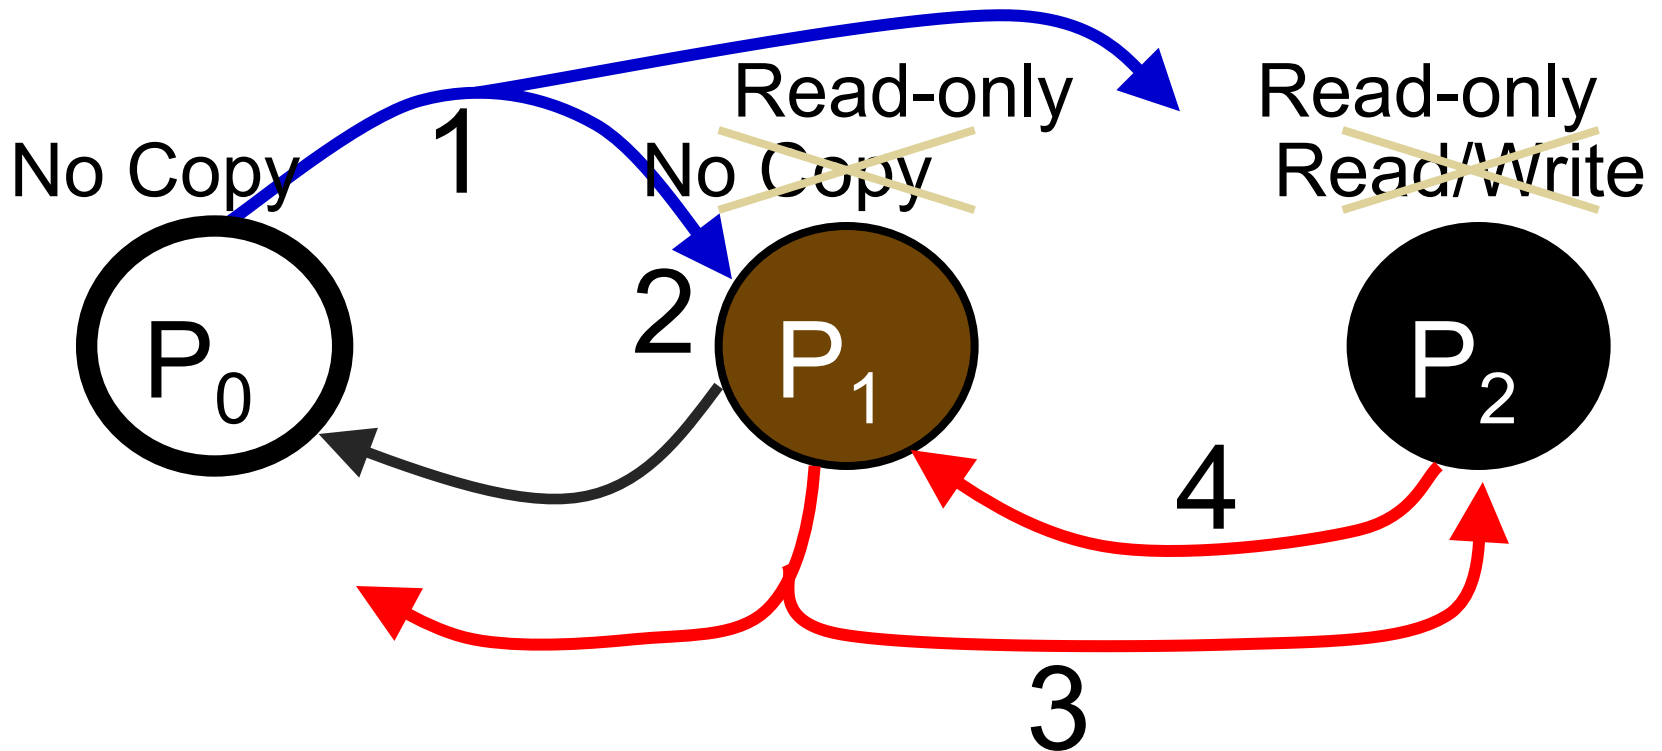


An Example Problem



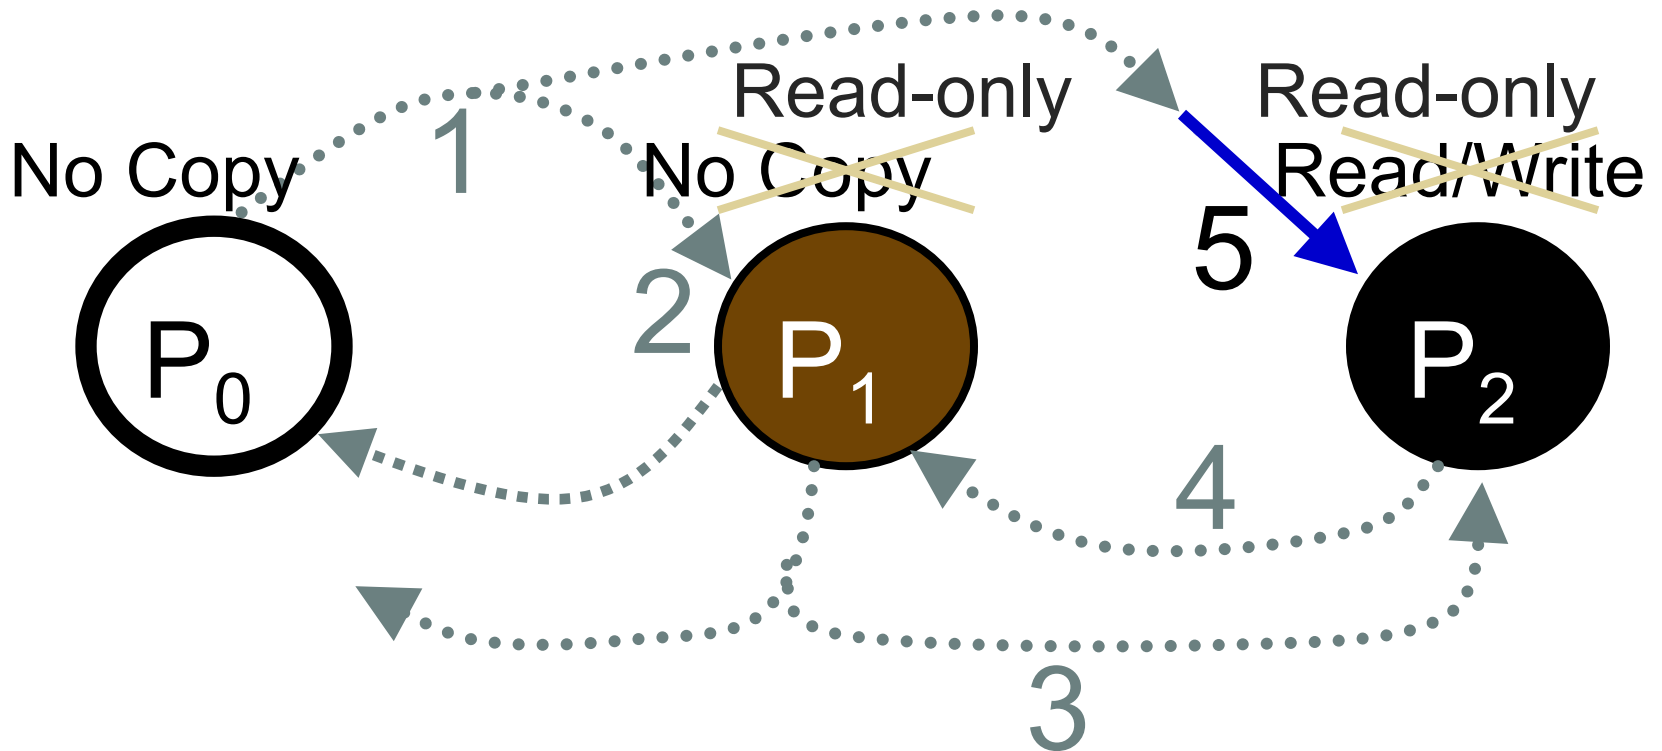
- P_0 issues a request to write (delayed to P_2)
- P_1 issues a request to read

An Example Problem



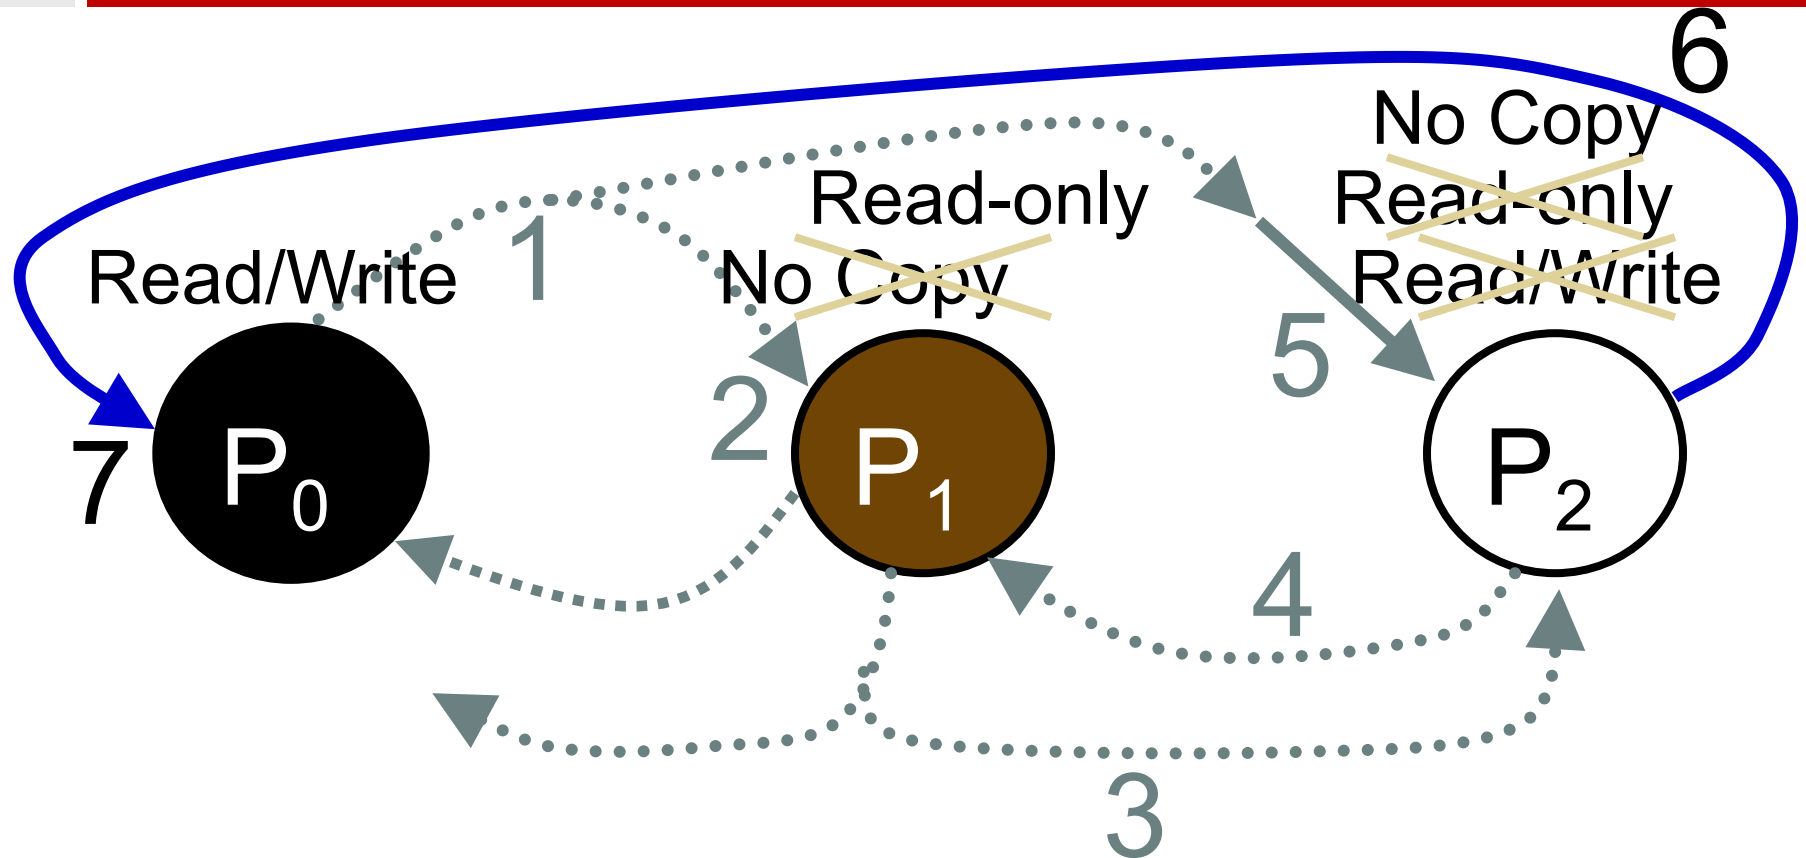
- P_2 responds with data to P_1

An Example Problem



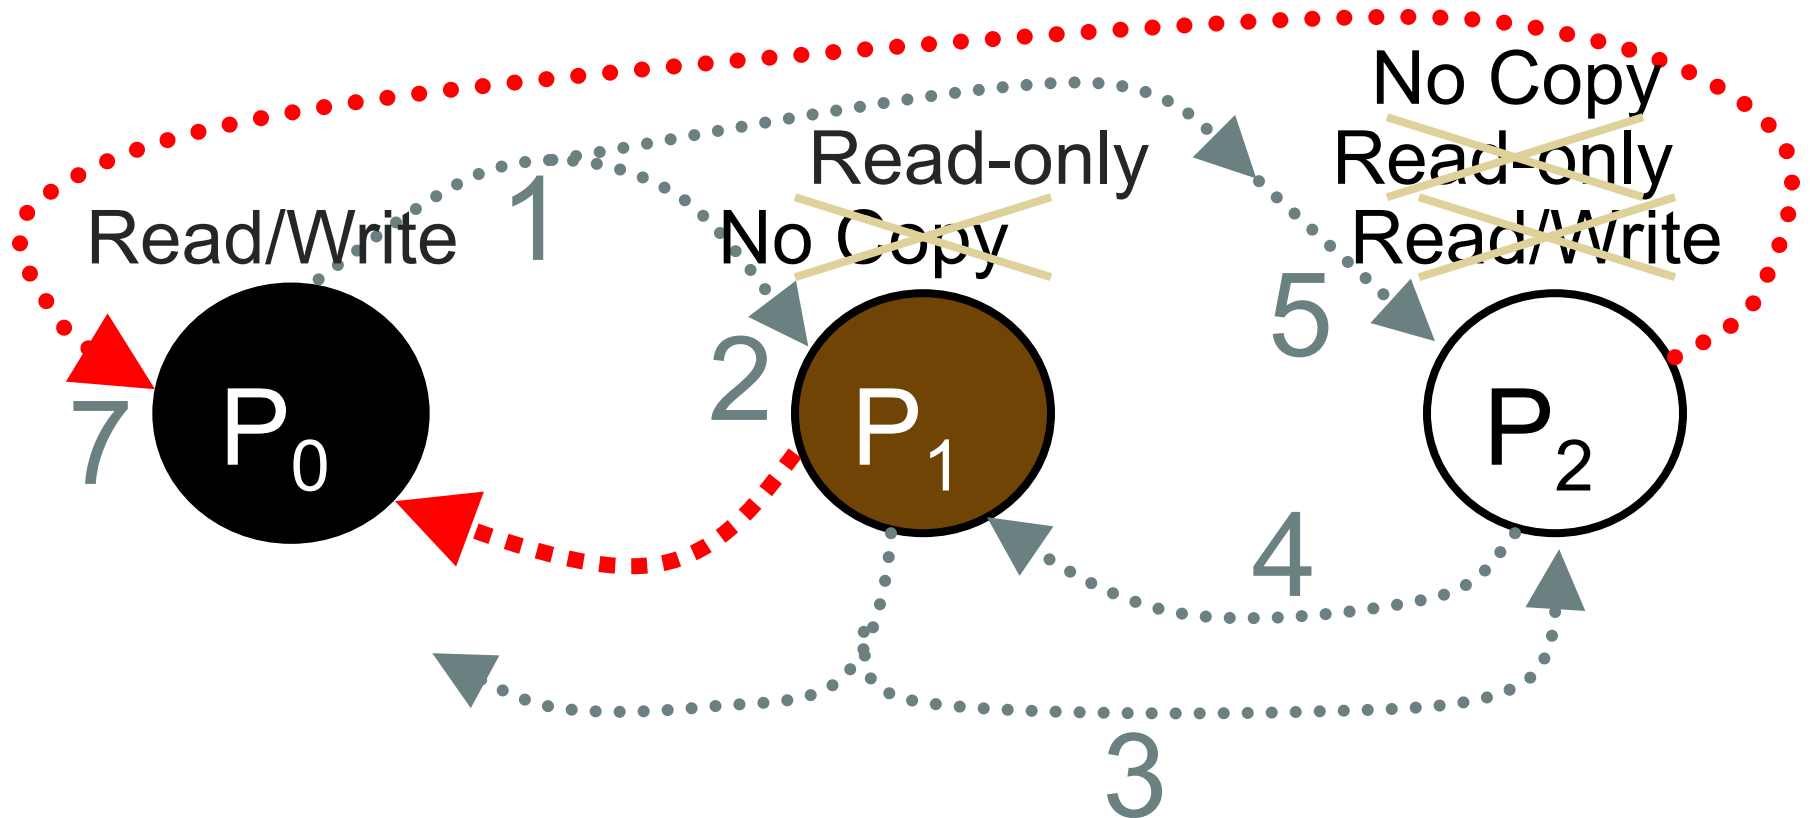
- P_0 's delayed request arrives at P_2

An Example Problem



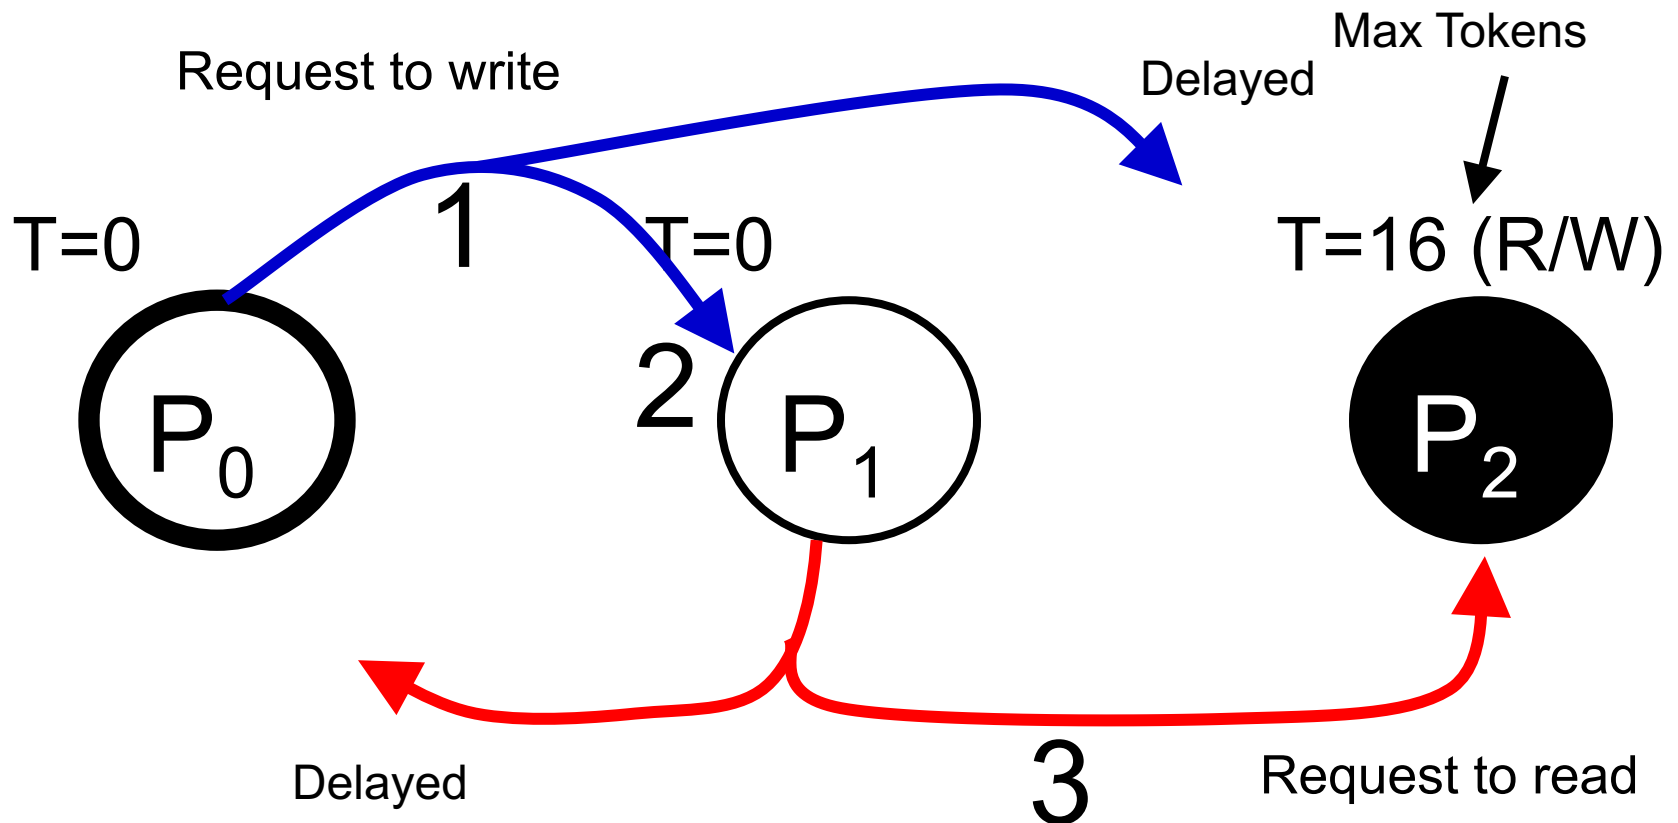
- P_2 responds to P_0

An Example Problem



Problem: P_0 and P_1 are in inconsistent states
Locally “correct” operation, globally inconsistent

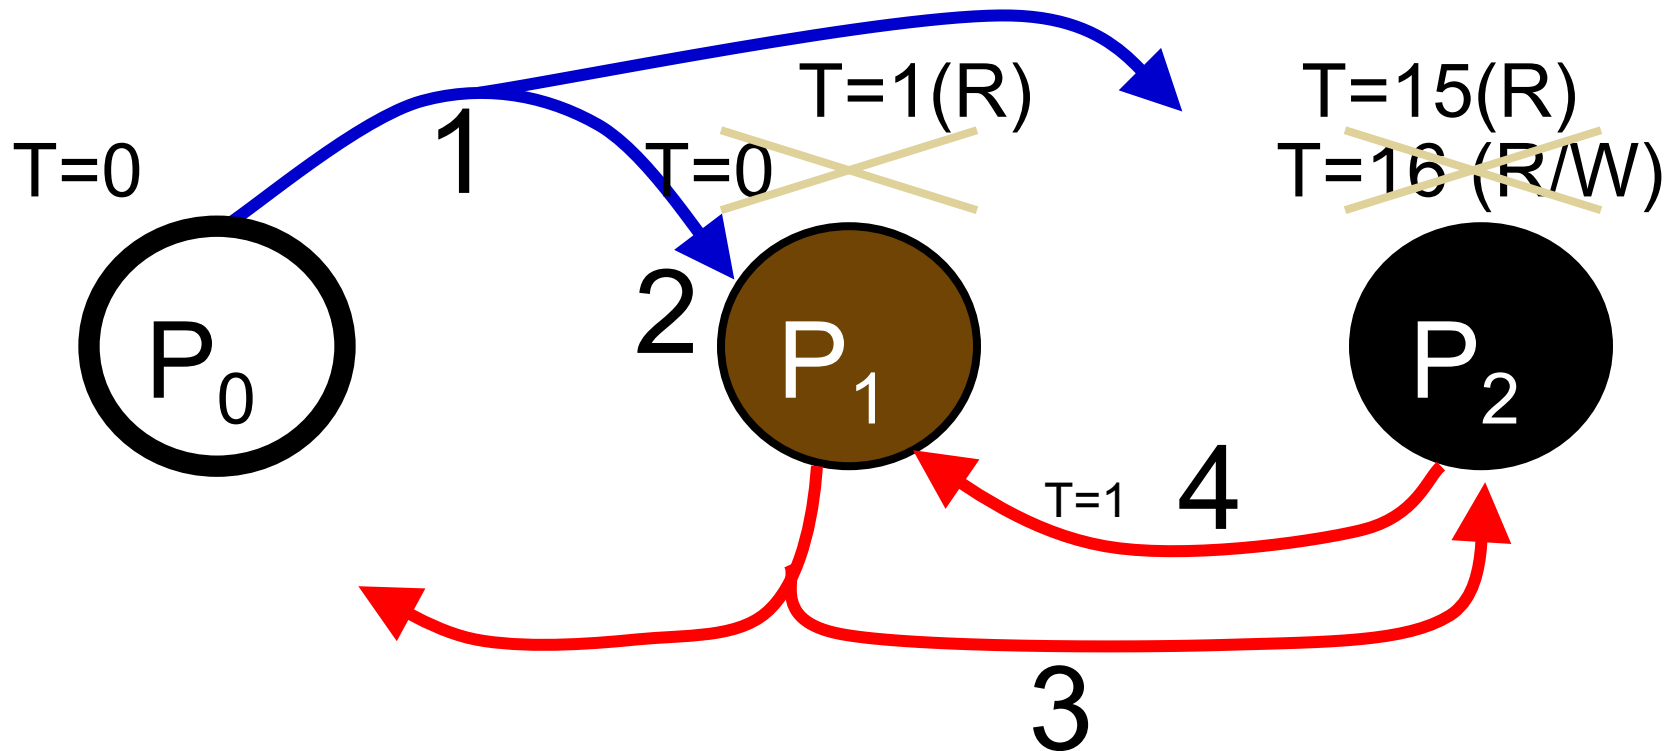
Token Coherence



- P_0 issues a request to write (delayed to P_2)
- P_1 issues a request to read

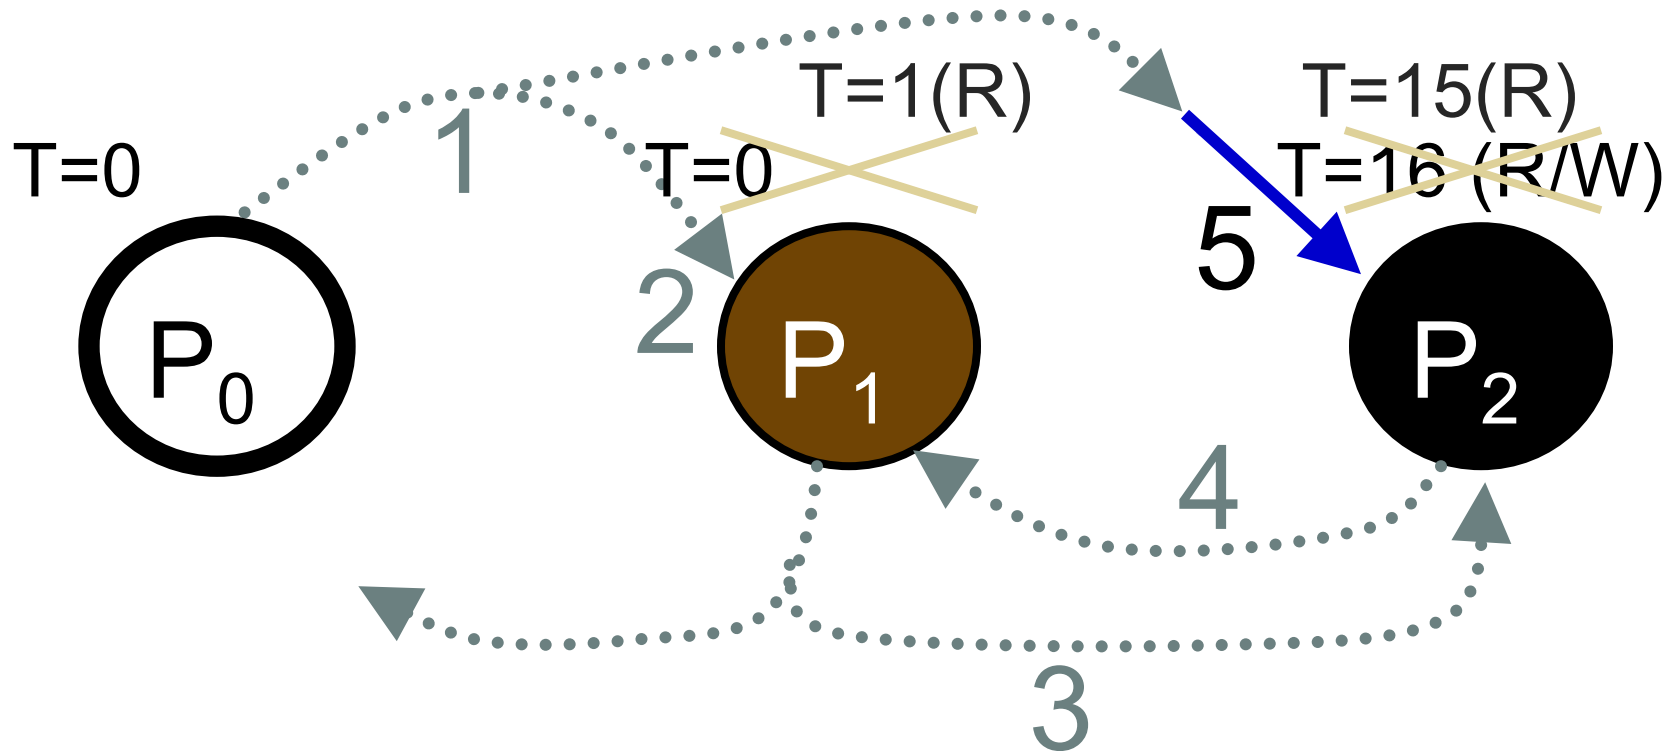
[Martin'03]

Token Coherence



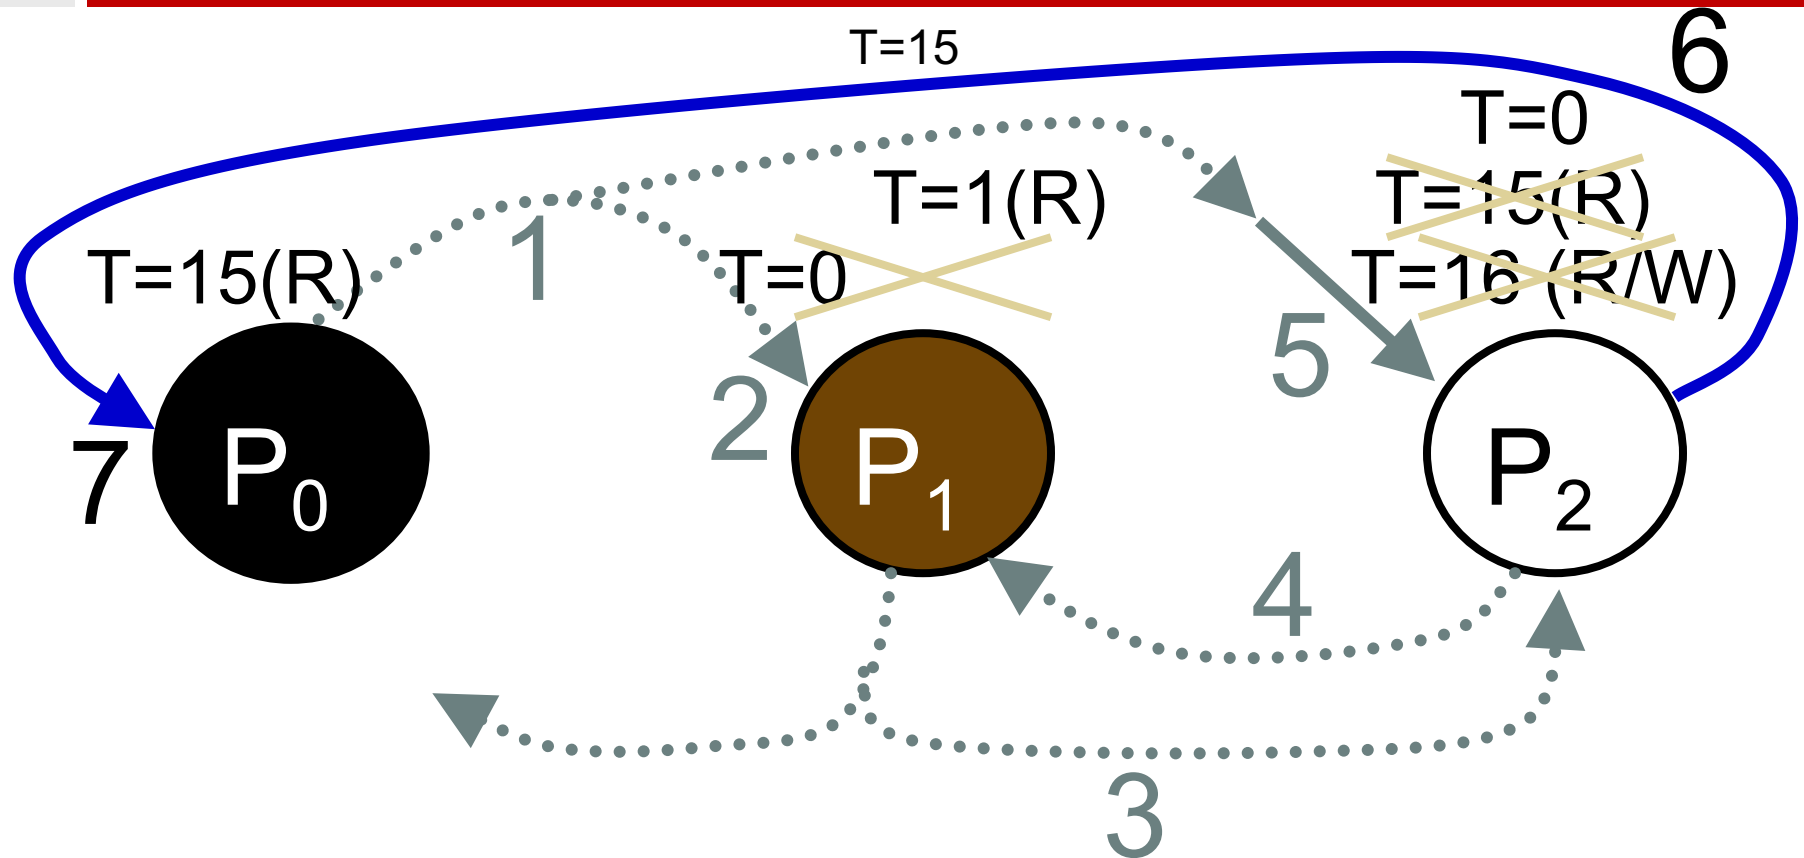
- P_2 responds with data to P_1

Token Coherence



- P_0 's delayed request arrives at P_2

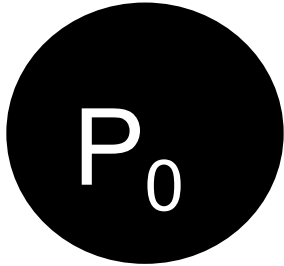
Token Coherence



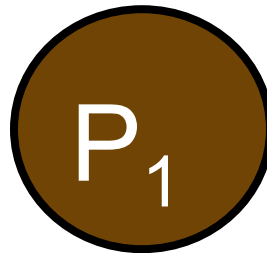
- P_2 responds to P_0

Token Coherence

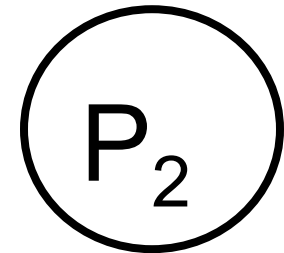
$T=15(R)$



$T=1(R)$

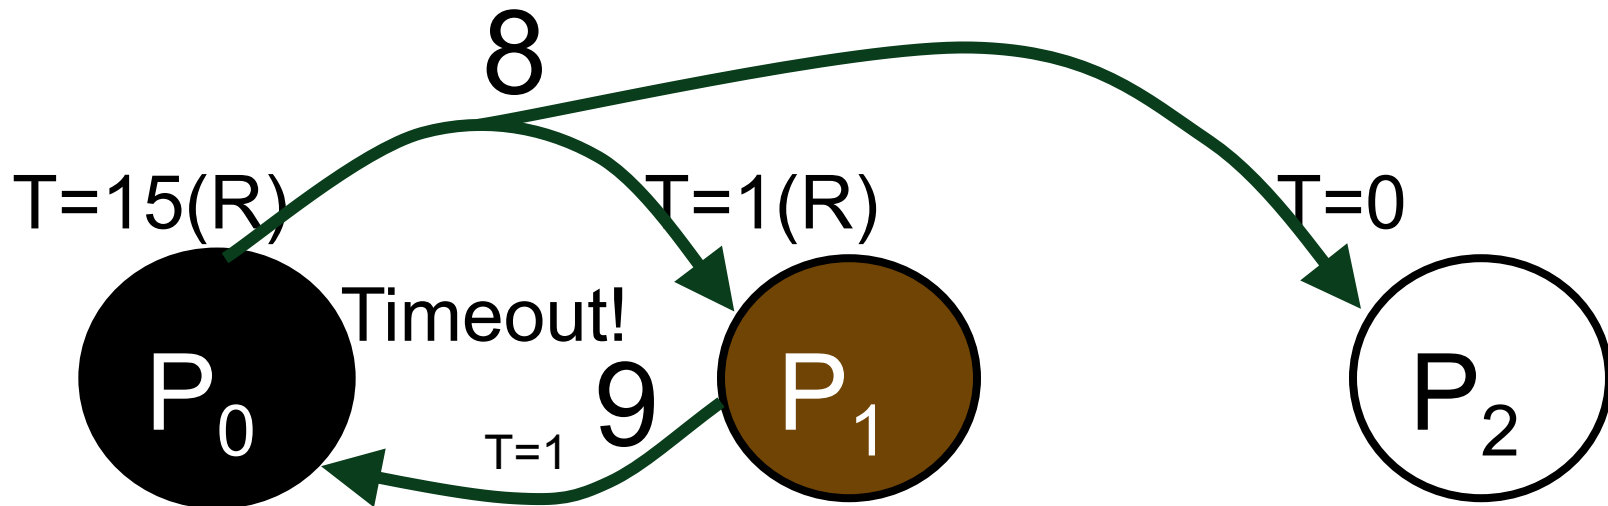


$T=0$



Now what? (P_0 wants **all** tokens)
[Martin'03]

Token Coherence

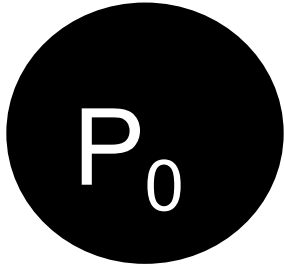


- P_0 reissues request
- P_1 responds with a token

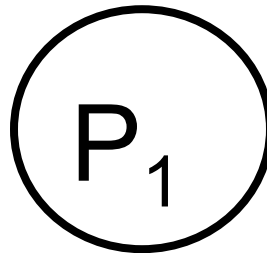
[Martin'03]

Token Coherence

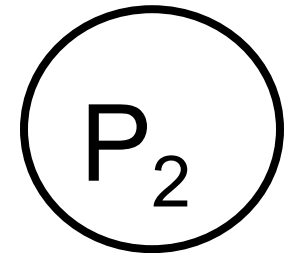
T=16 (R/W)



T=0



T=0



One final issue: What about starvation?

- P_0 's request completed

[Martin'03]