MEMORY SYSTEMS

Mahdi Nazm Bojnordi

Assistant Professor

School of Computing

University of Utah

Overview

- Upcoming deadline
  - Feb. 27th: homework assignment will be posted
- Group projects

	Subject	Presentation
Group 1	Reducing cache energy via skipping common values	April 17
Group 2	Integrating high bandwidth memory with high capacity DIMMs for data-centric workloads	April 17
Group 3	A near-threshold spiking neural network accelerator	April 17
Group 4	Improving energy efficiency of low power micro-controllers used in various IoT nodes	April 19
Group 5	In-situ hardware accelerator for clustering applications	April 19
Group 6	Improving performance of time-based data encoding	April 19

Overview

- Student paper presentation
  - Prepare for exactly 20m talk followed by 5m Q&A

Presenter Names	Date
Kohl, Meher, Shirley	March 29
Karl, Anirban, Chandrasekhar	April 3
Suryanarayanan, Tim, Arjun	April 5
Pranav, Goverdhan, Yomi	April 10
Munzer, Manikanth, Amandeep	April 12

- This lecture
  - Main memory systems
  - DRAM architecture basics

Main Memory System

- A critical component of all computing systems
  - server, mobile, embedded, desktop, sensor

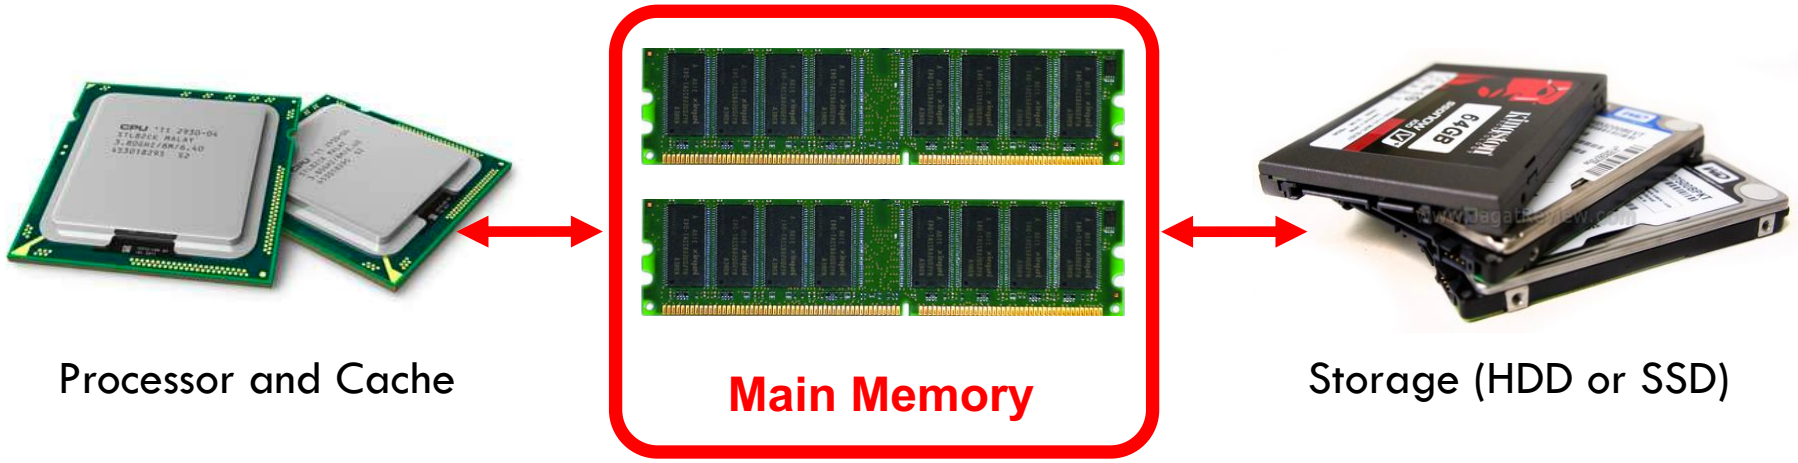

- Must scale to maintain performance growth
  - size, technology, efficiency, cost, and control algorithms

Why Main Memory is Important?

- Shared resource
  - Multiple applications running on different processor cores
  - Different objectives and requirements
  - Highly contented resource

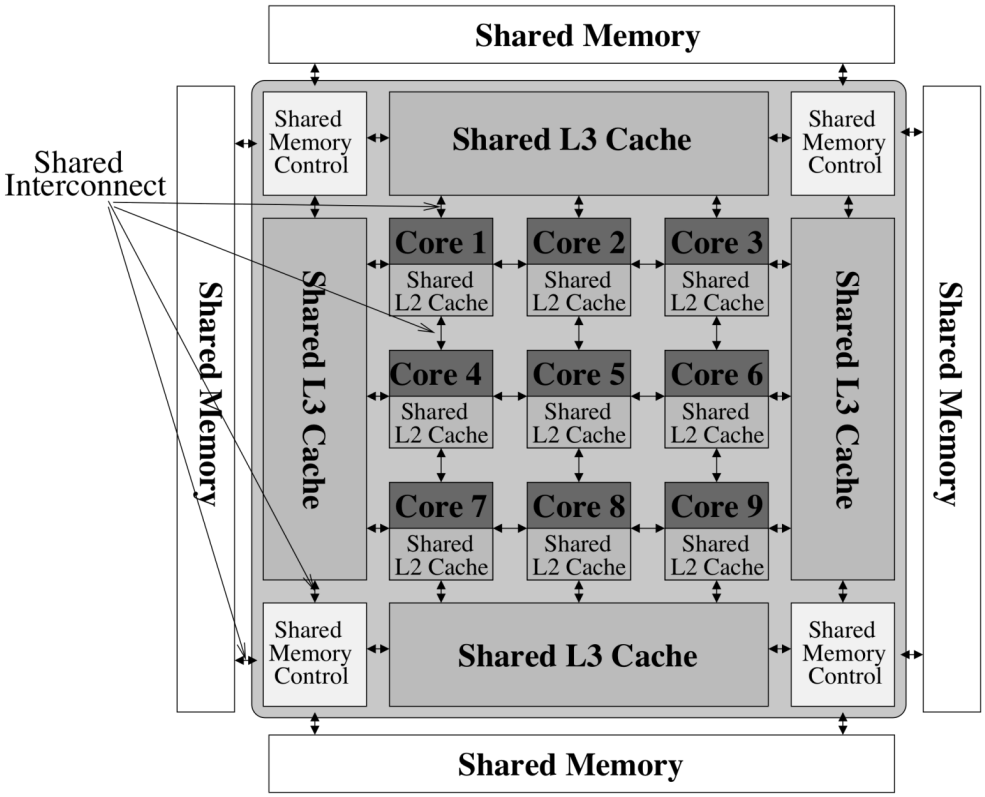

Complex control policies and microarchitectures are required.

Scalability Challenges

- Increasing need for memory capacity, bandwidth, and quality-of-service maintenance
  - increasing number of cores (multicores)
  - increasing demand for data (big data processing)
  - cloud computing, GPUs, mobile (consolidation)

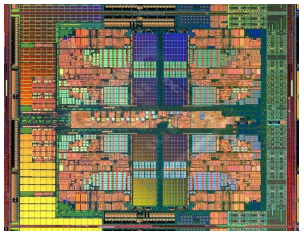

AMD Barcelona

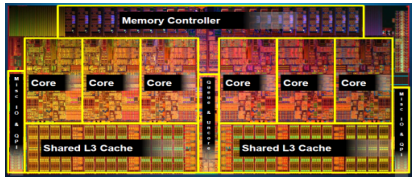

Intel Core i7
8 cores

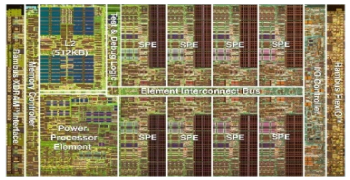

IBM Cell BE
8+1 cores

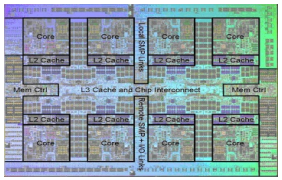

IBM POWER7
8 cores

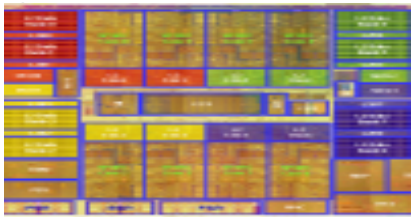

Sun Niagara II
8 cores

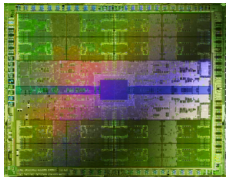

Nvidia Fermi
448 “cores”

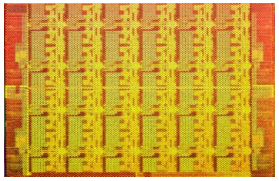

Intel SCC
48 cores, networked

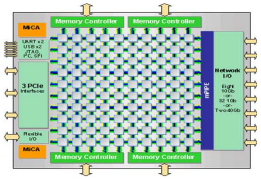

Tilera TILE Gx
100 cores, networked

CPU-DRAM Gap

- Core count doubling ~ every 2 years
- DRAM DIMM capacity doubling ~ every 3 years

Memory capacity per core expected to drop by 30% every two years

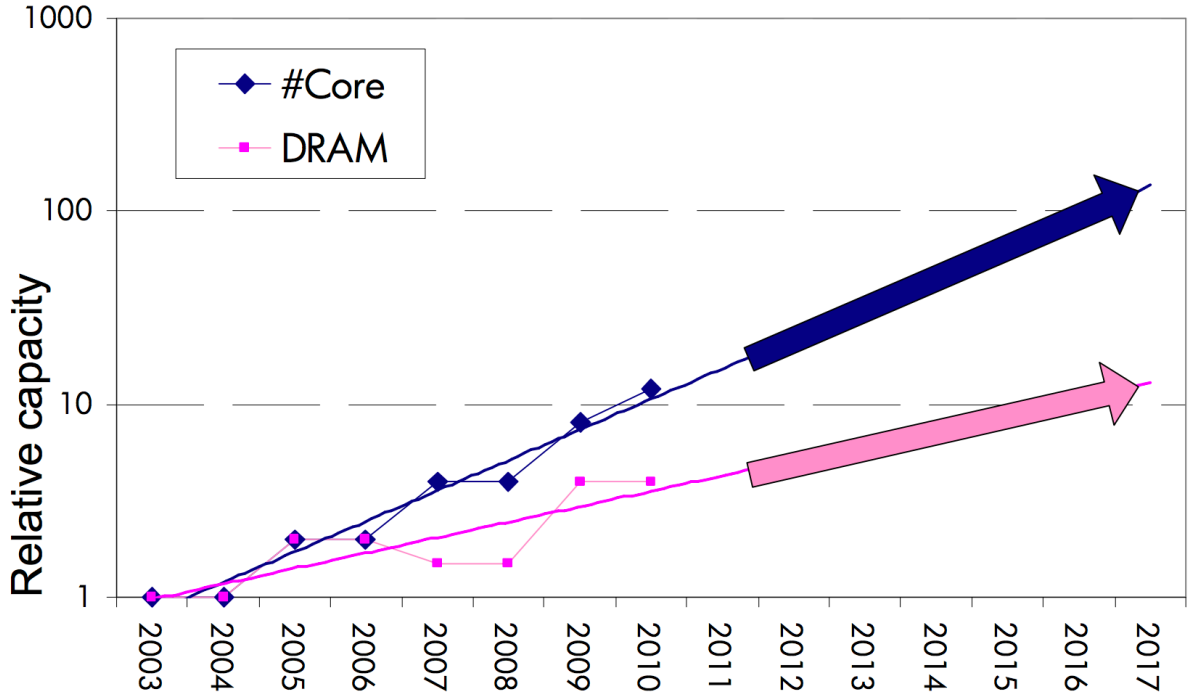

[Lim'09]

DRAM: Design Challenges

- Main memory energy/power is a key system design concern
  - Energy spent in off-chip memory hierarchy is about 40-50% [Lefurgy'03]
  - DRAM consumes power even when not used
    - periodic refresh
- DRAM technology scaling is ending
  - stops gaining higher capacity, lower cost, lower energy

DRAM: Logical Organization

- Five DRAM coordinates
  - Channel, rank, bank, row, column

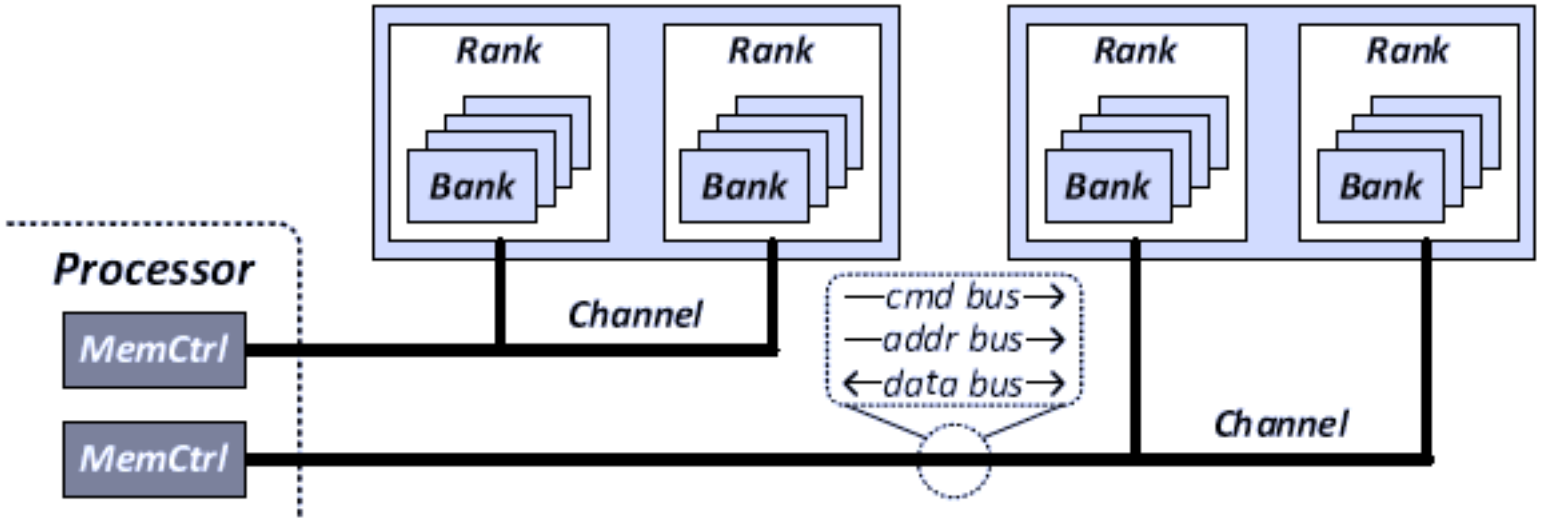

[Kim'12]

DRAM: Physical Organization

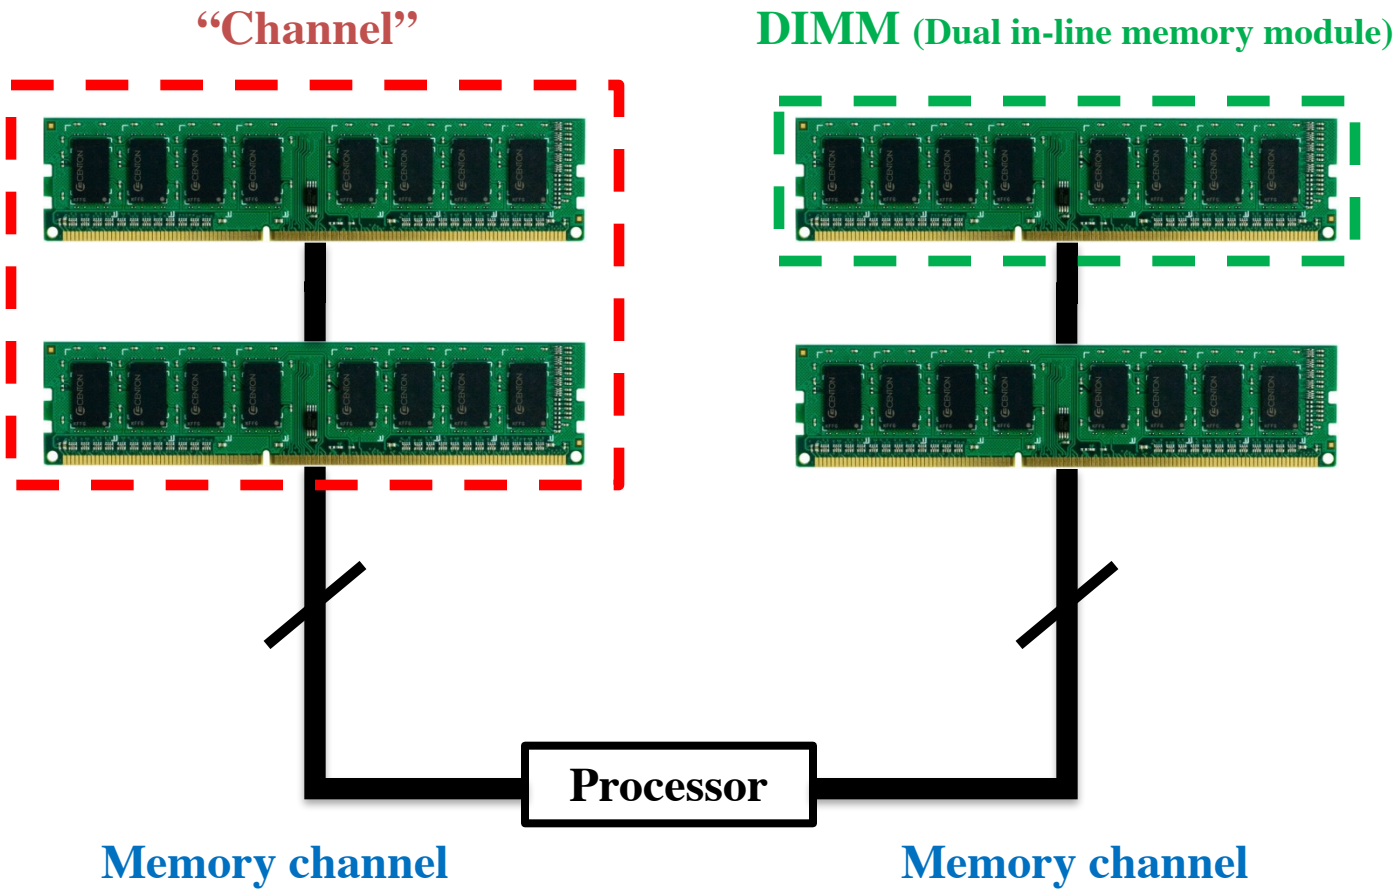

DIMM Structure

DIMM (Dual in-line memory module)

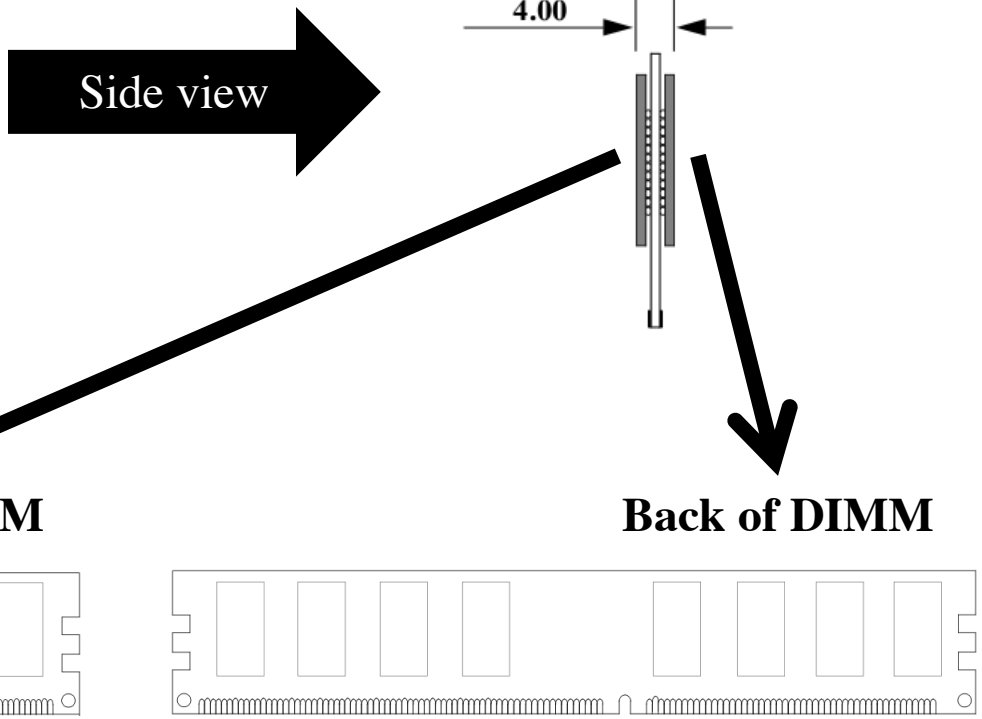

DIMM Structure

DIMM (Dual in-line memory module)

Side view

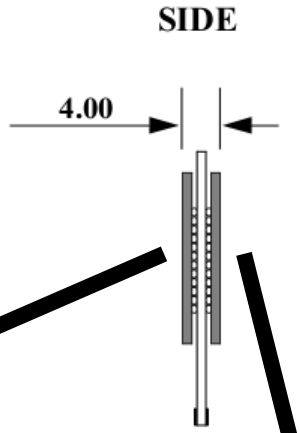

Front of DIMM

Back of DIMM

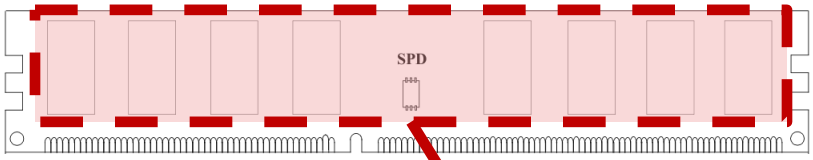

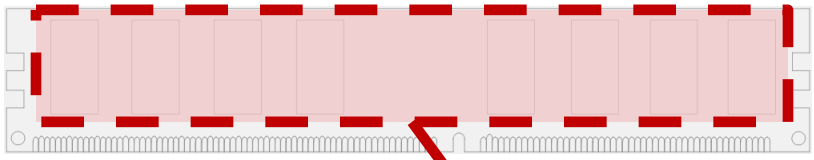

Rank 0: collection of 8 chips

Rank 1

Rank Organization

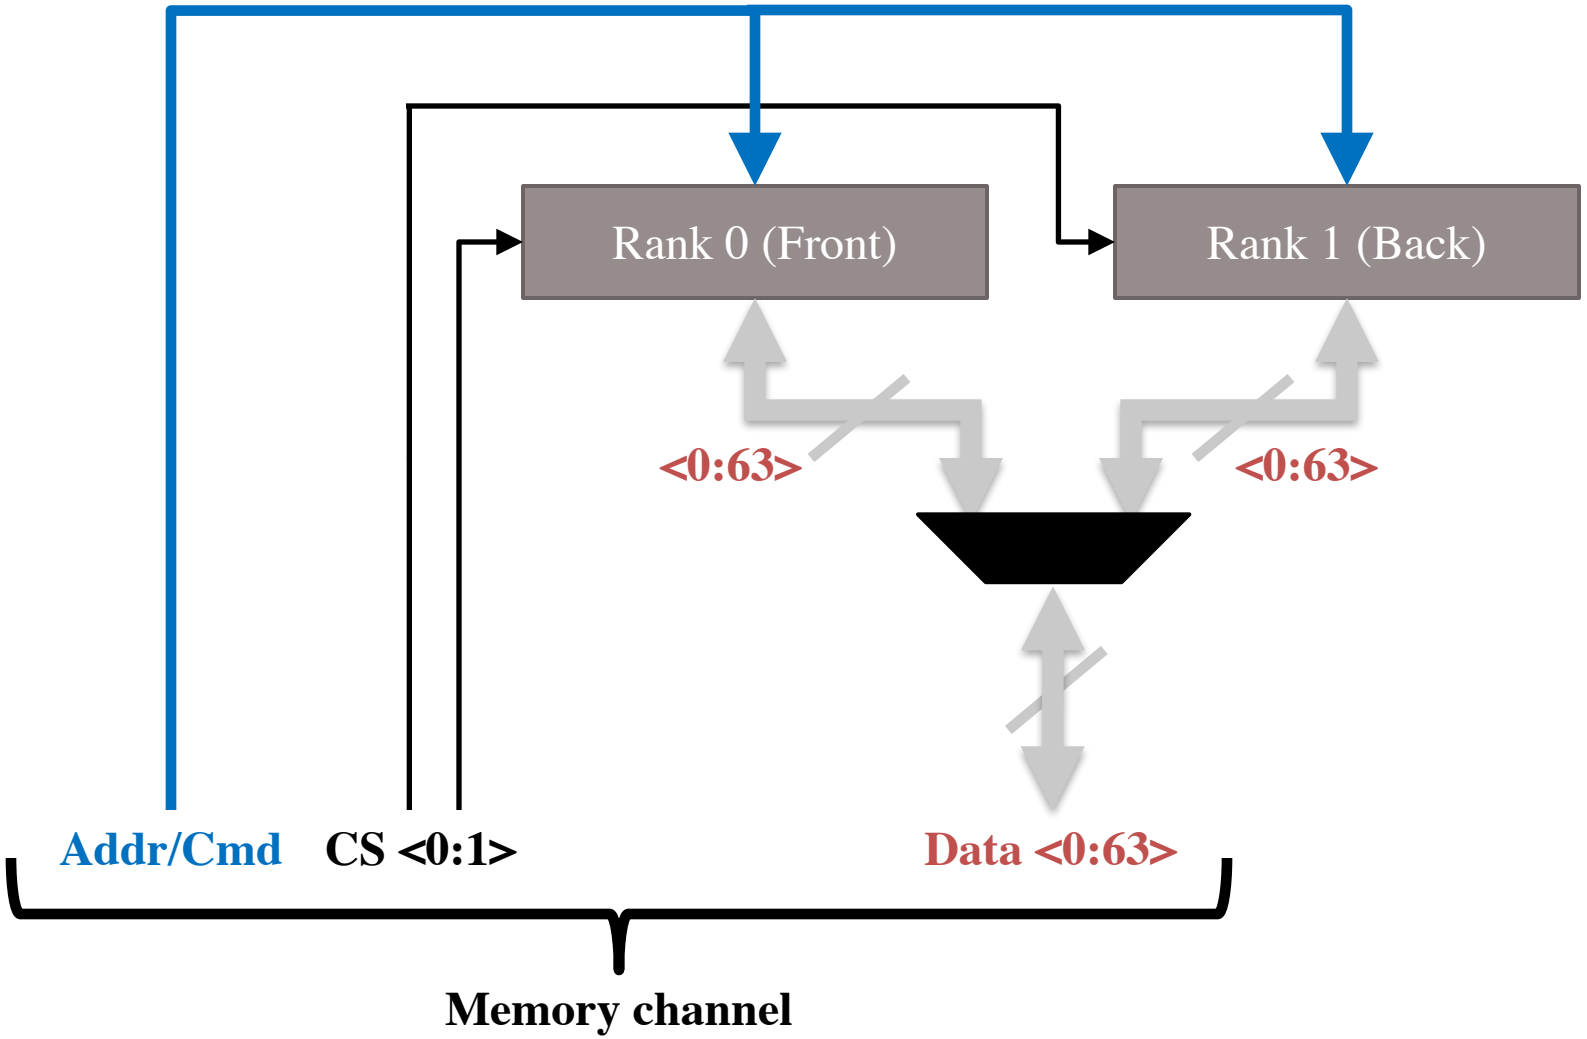

Rank Breakdown

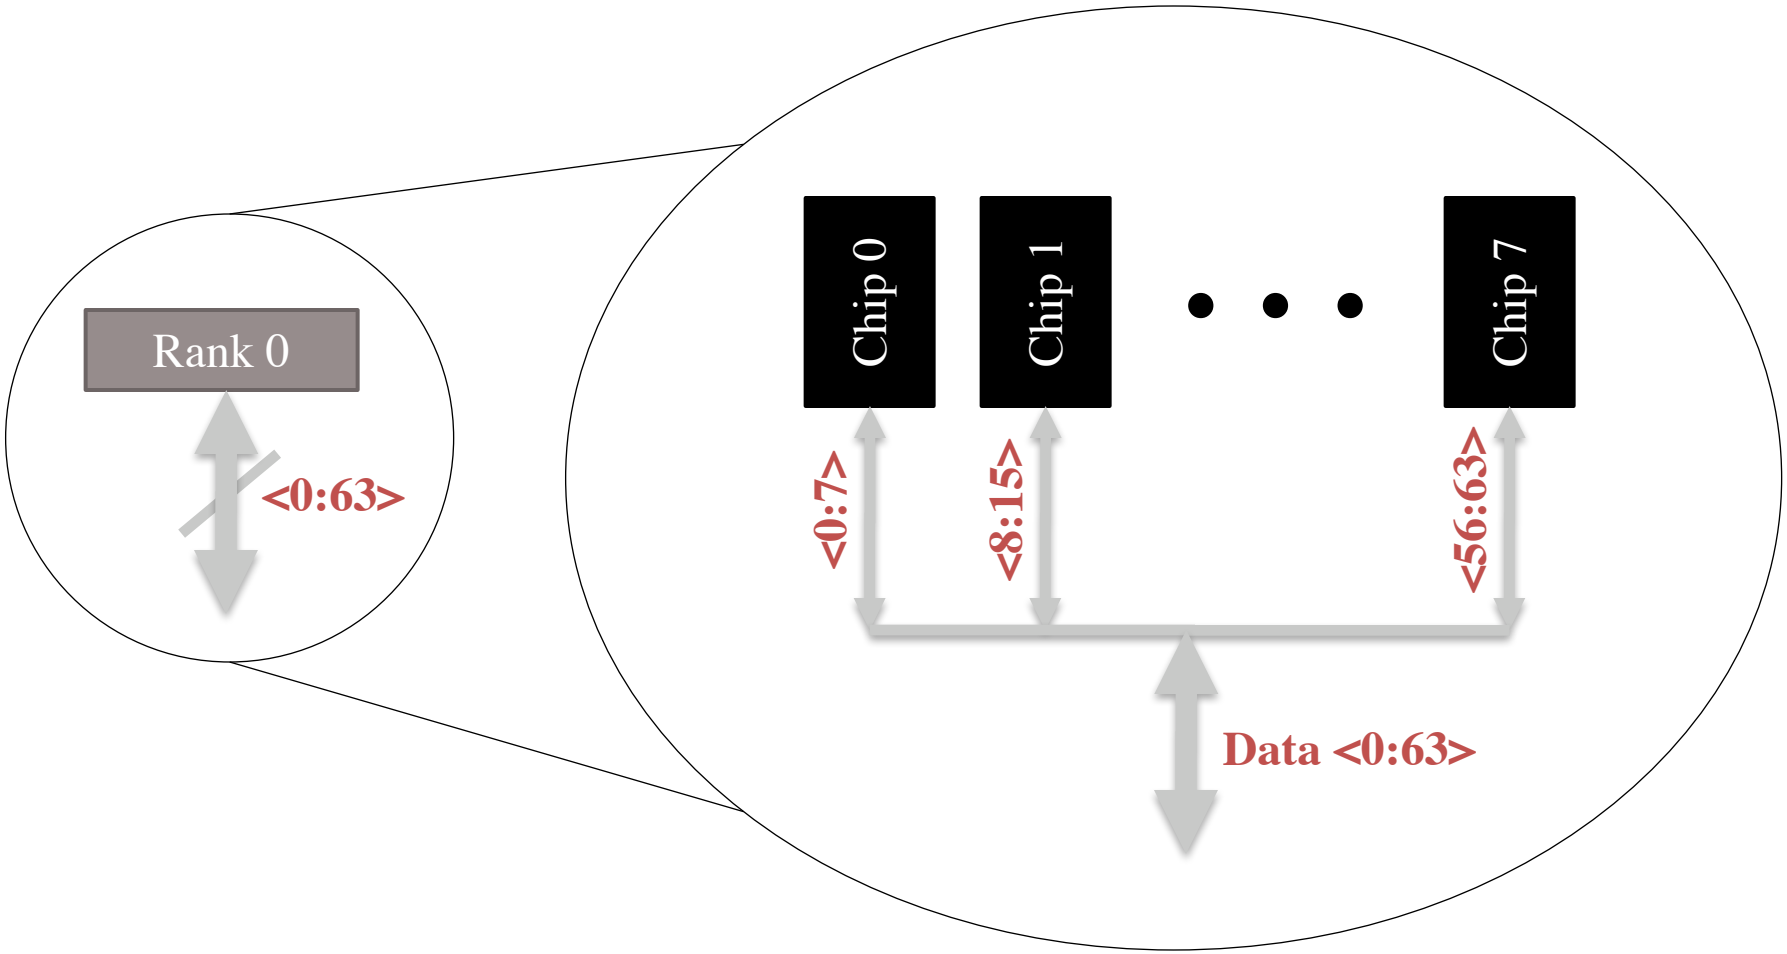

DRAM DIMM and Rank

- Multiple chips operated together to form a wide interface
  - All chips within a rank are controlled at the same time
  - Respond to a single command
  - Share address and command; different data bits

- A DRAM module consists of one or more ranks
  - e.g., DIMM (dual inline memory module)
  - If we have chips with 8-bit interface, to read 8 bytes in a single access, use 8 chips in a DIMM

Chip Structure

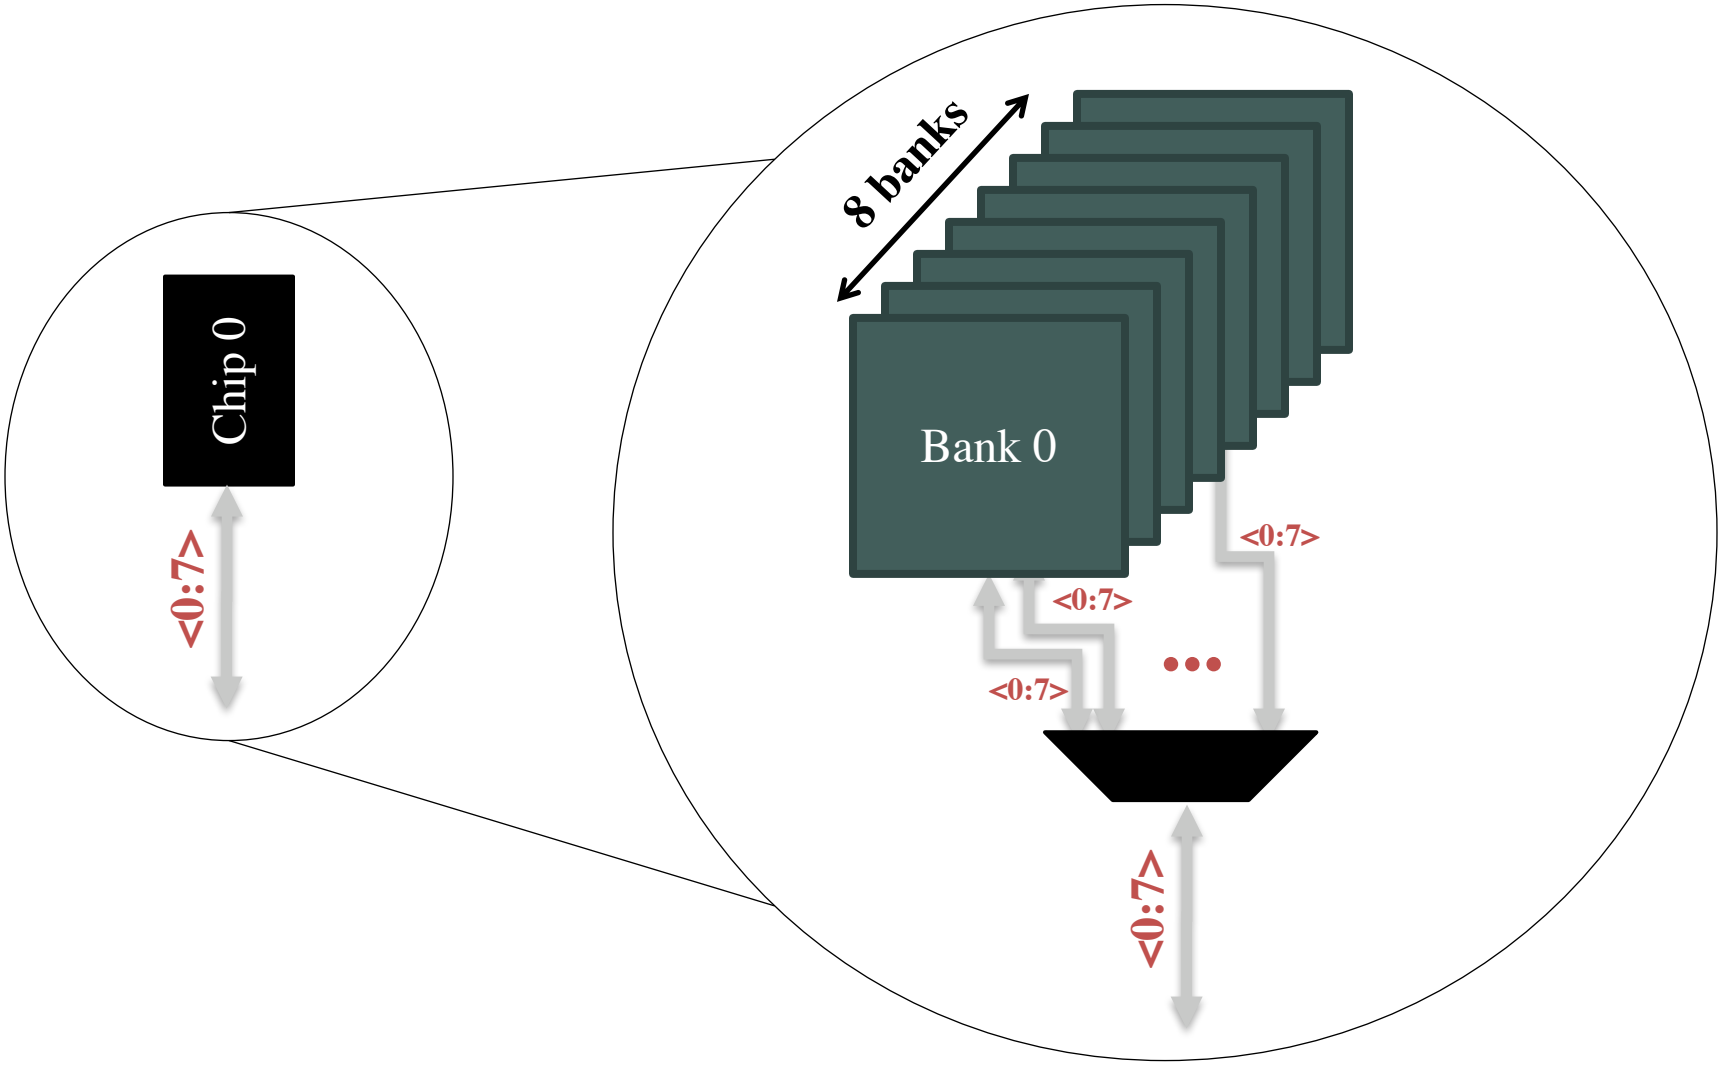

Bank Organization

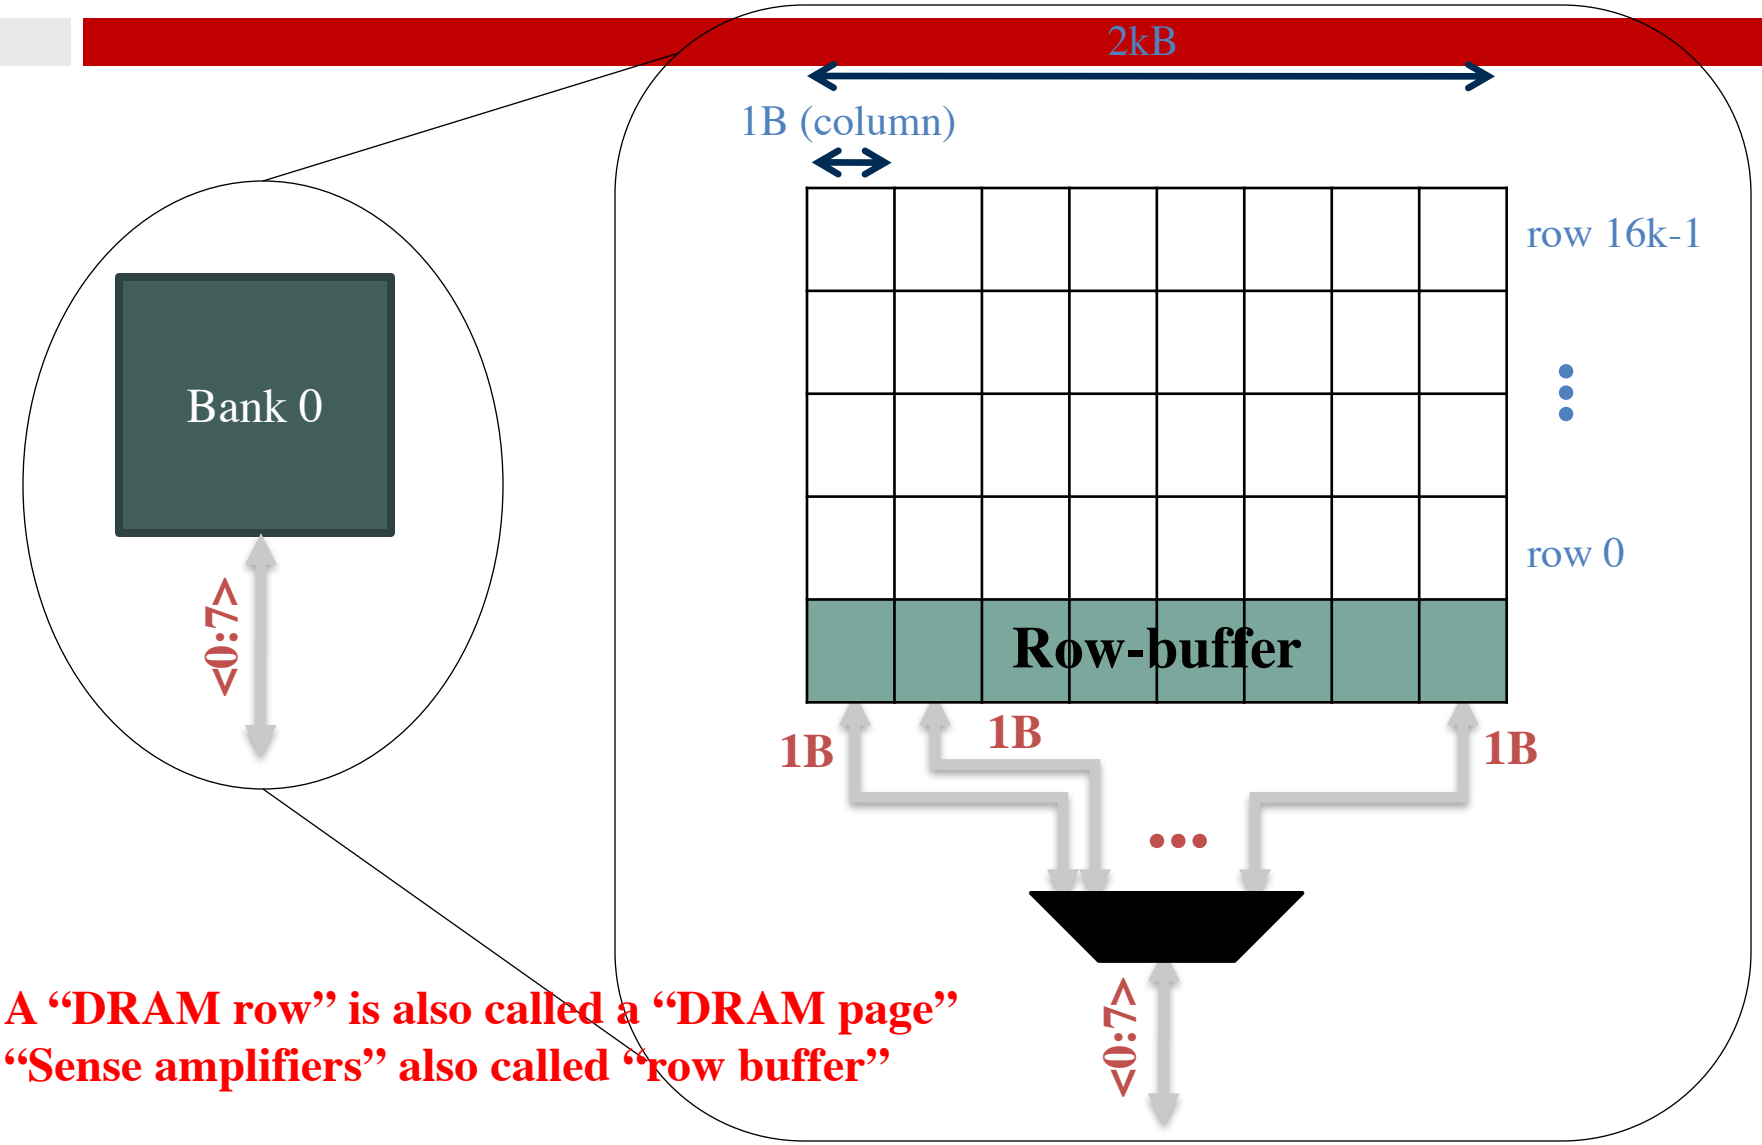

A "DRAM row" is also called a "DRAM page"
"Sense amplifiers" also called "row buffer"

DRAM Page Access

- Access to a closed row
  - Activate command opens row (placed into row buffer)
  - Read/write command reads/writes column in the row buffer
- Precharge command closes the row and prepares the bank for next access
- Access to an “open row”
  - No need for activate command

DRAM Page Access

Access Address:
(Row 0, Column 0)
(Row 0, Column 1)
(Row 0, Column 85)
(Row 1, Column 0)

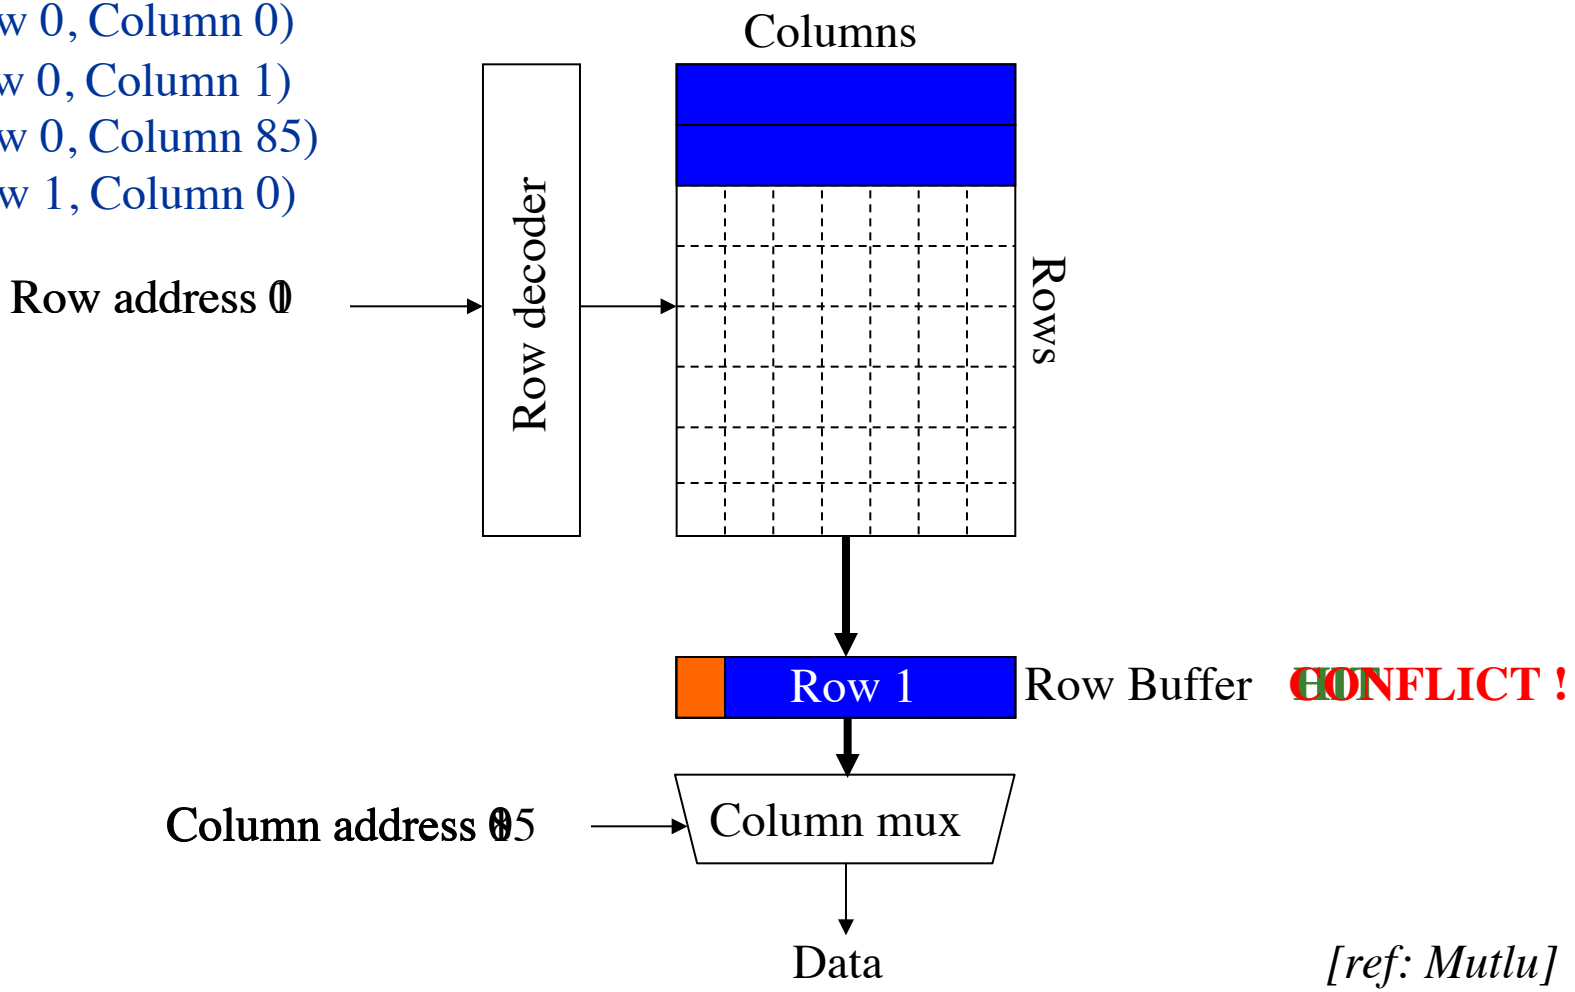

[ref: Mutlu]

DRAM: Parallel Access

- DRAM subsystem comprises multiple banks
  - Organized under independent channels and ranks

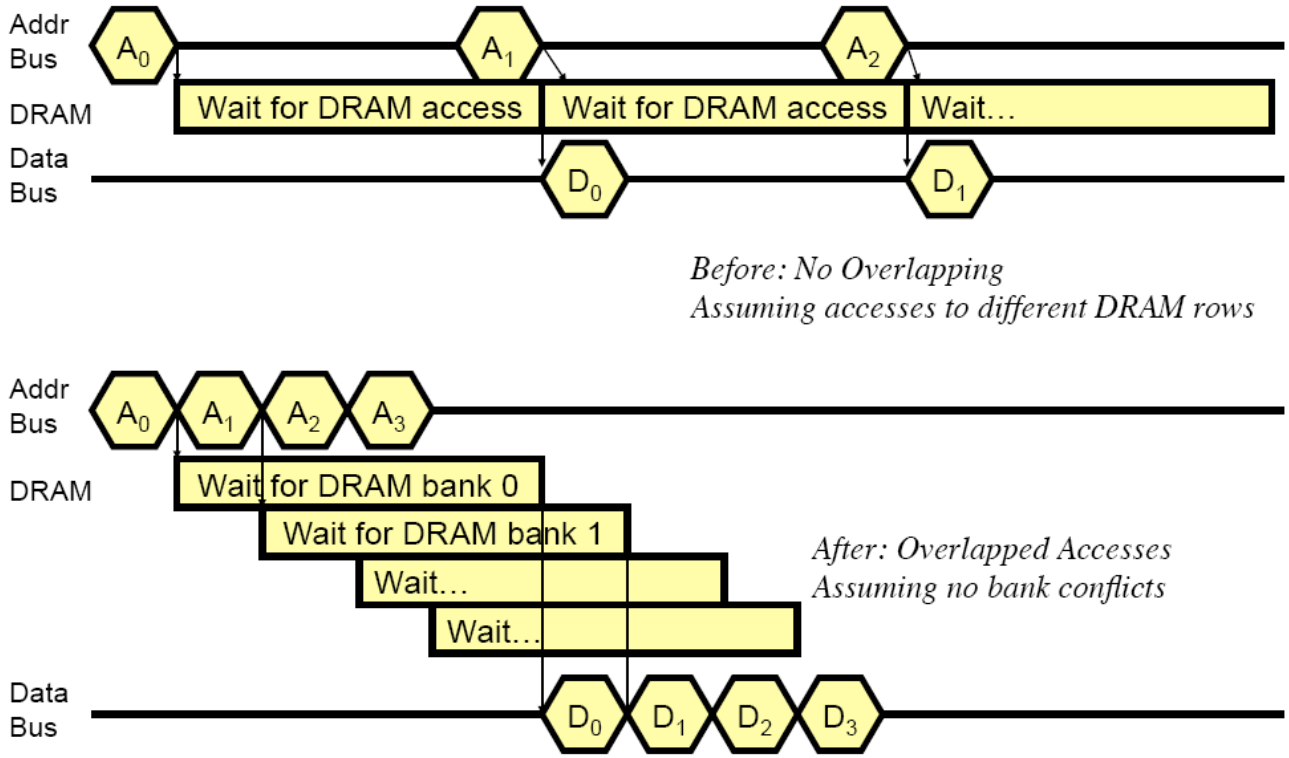

DRAM Controller

- Ensure correct operation of DRAM (refresh and timing)
- Service DRAM requests while obeying timing constraints of DRAM chips
  - Constraints: resource conflicts (bank, bus, channel), minimum write-to-read delays
  - Translate requests to DRAM command sequences
- Buffer and schedule requests to improve performance
  - Reordering, row-buffer, bank, rank, bus management
- Manage power consumption and thermals in DRAM
  - Turn on/off DRAM chips, manage power modes

DRAM Controller

- Ensuring DDRx timing constraints

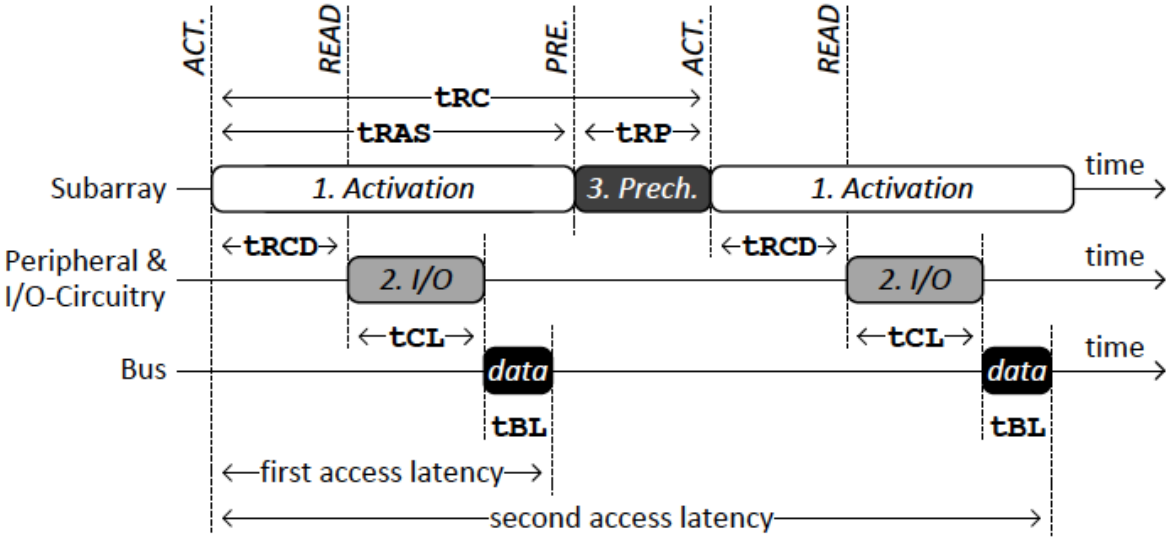

Figure 5. Three Phases of DRAM Access

Table 2. Timing Constraints (DDR3-1066) [43]

Phase	Commands	Name	Value
1	ACT → READ ACT → WRITE	tRCD	15ns
	ACT → PRE	tRAS	37.5ns
2	READ → *data*	tCL	15ns
	WRITE → *data*	tCWL	11.25ns
	data burst	tBL	7.5ns
3	PRE → ACT	tRP	15ns
1 & 3	ACT → ACT	tRC (tRAS+tRP)	52.5ns

[Kim'12]