

DRAM REFRESH MANAGEMENT

Mahdi Nazm Bojnordi

Assistant Professor

School of Computing

University of Utah

Overview

- Upcoming deadline
 - ▣ Feb. 27th: homework assignment will be posted
- This lecture
 - ▣ DRAM address mapping
 - ▣ DRAM refresh basics
 - ▣ Smart refresh
 - ▣ Elastic refresh
 - ▣ Avoiding or pausing refreshes

DRAM Address Mapping


- Where to store cache lines in main memory?

Typical Mapping



Application A:

Good distribution of memory requests
among DRAM banks.



DRAM Address Mapping


- Where to store cache lines in main memory?

Typical Mapping



Application B:

Unbalanced distribution of memory requests among DRAM banks.



DRAM Address Mapping

- How to compute bank ID?


Custom Mapping




Application B:

Good distribution of memory requests
among DRAM banks.


DRAM Banks



Cache Line Interleaving

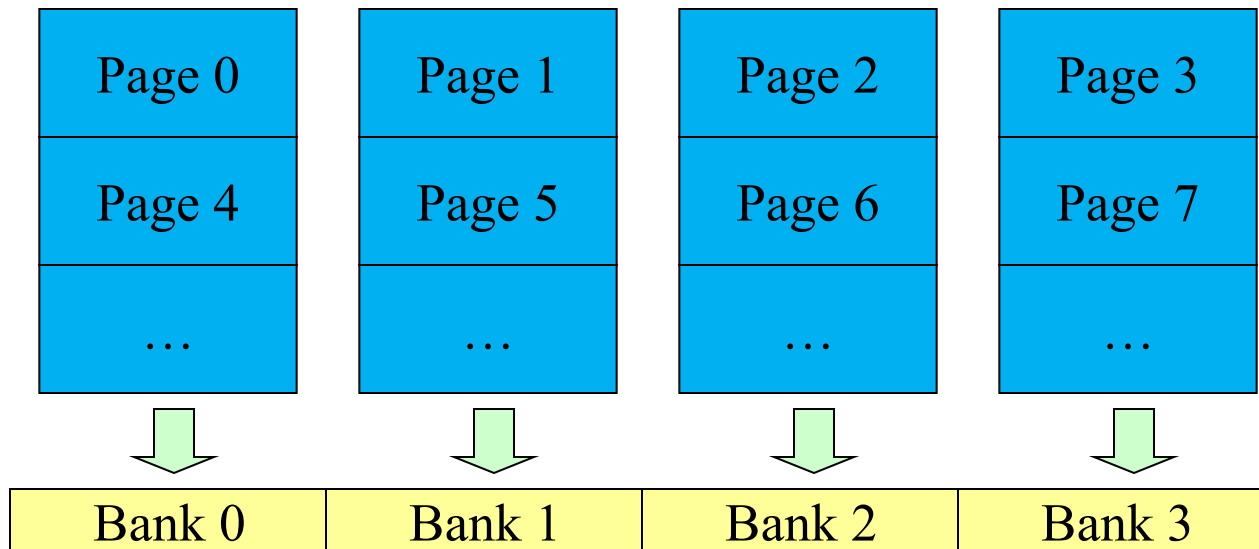


Address format




Spatial locality is not well preserved!

Page Interleaving




Address format




Cache Line Mapping

- Bank index is a subset of set index


Cache line interleaving



Page interleaving

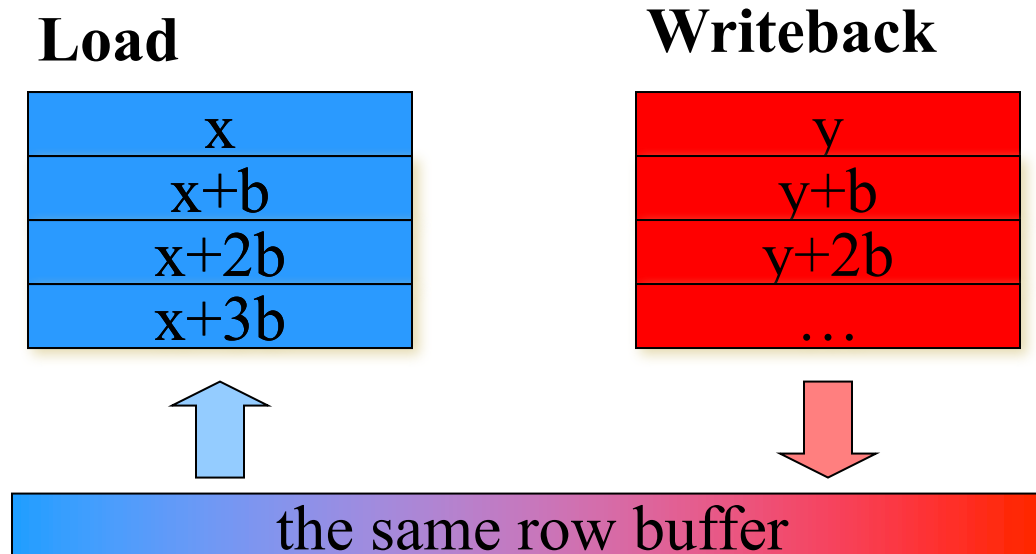


Cache-related representation




Row Buffer Conflict

- Problem: interleaving load and writeback streams with the same access pattern to the banks may result in row buffer misses




Key Issues

- To exploit spatial locality, use maximal interleaving granularity (or row-buffer size)
- To reduce row buffer conflicts, use only those bits in cache set index for “bank bits”

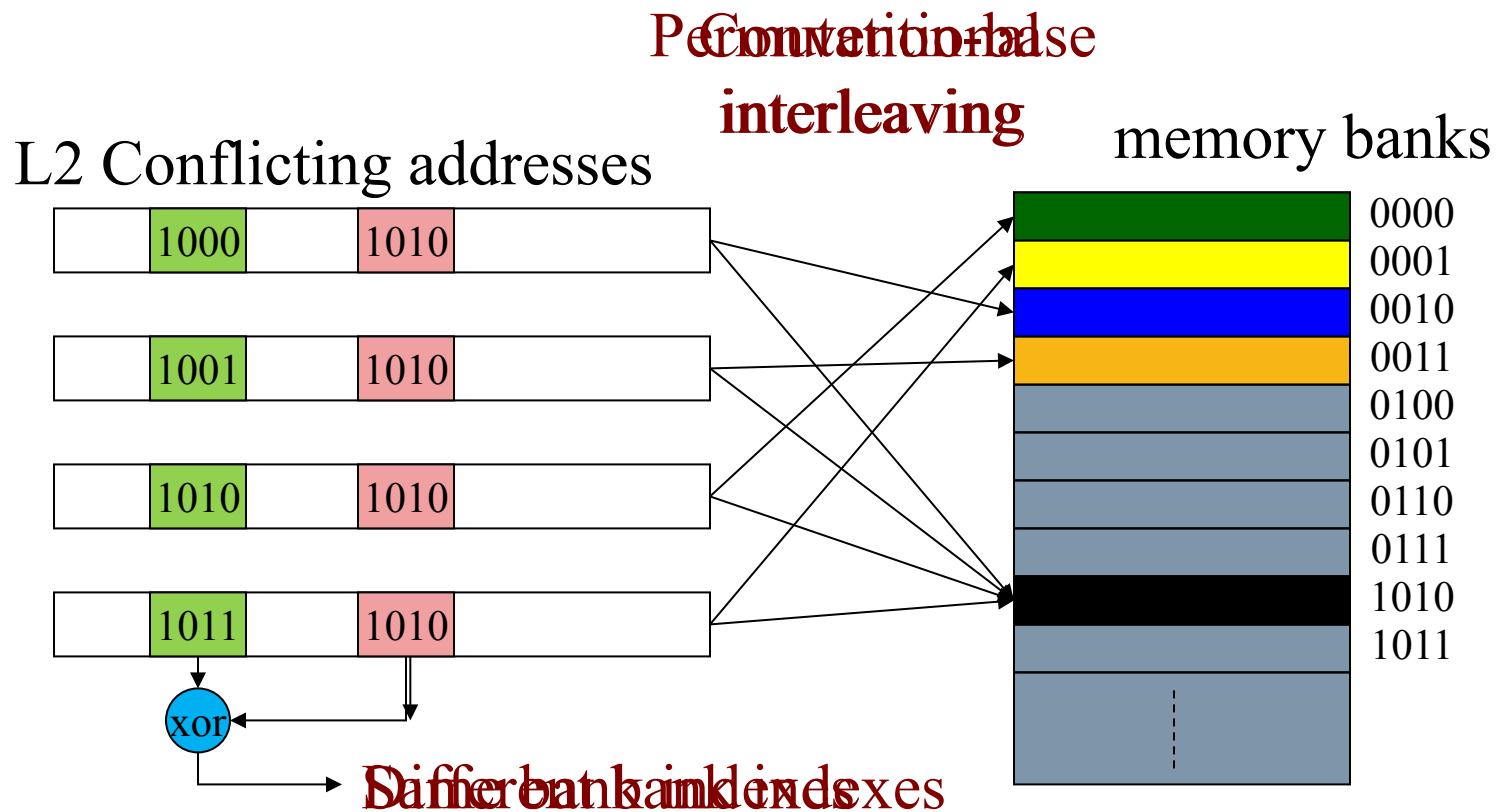


Permutation-based Interleaving




Permutation-based Interleaving

- New bank index





Permutation-based Interleaving



DRAM Refresh

- DRAM cells lose charge over time
- Periodic refresh operations are required to avoid data loss
- Two main strategies for refreshing DRAM cells
 - ▣ **Burst refresh:** refresh all of the cells each time
 - Simple control mechanism (e.g., LPDDR_x)
 - ▣ **Distributed refresh:** a group of cells are refreshed
 - Avoid blocking memory for a long time

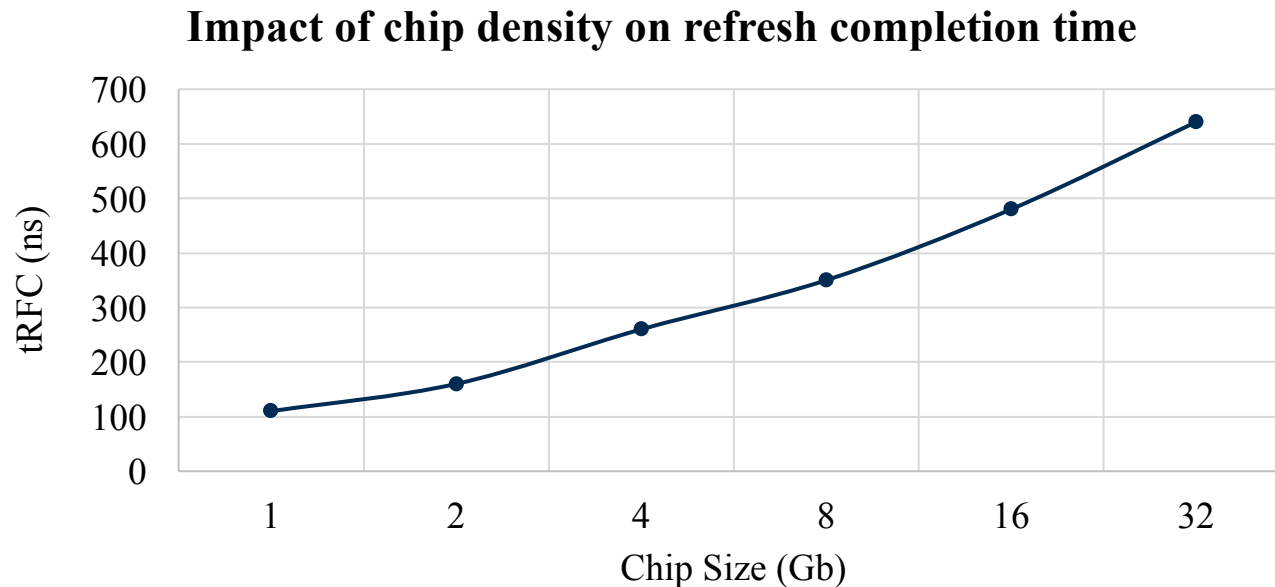


Refresh Basics

- **tRET**: the retention time of DRAM leaky cells (64ms)
 - ▣ All cells must be refreshed within tRET to avoid data loss
- **tREFI**: refresh interval, which is the gap between two refresh commands issues by the memory controller
 - ▣ MC sends 8192 auto-refresh commands to refresh one bin at a time
 - $tREFI = tRET / 8192 = 7.8\mu s$
- **tRFC**: the time to finish refreshing a bin (refresh completion)
- What is the bin size?

Refresh Basics

- tRFC increases with chip capacity




Controlling Refresh Operations

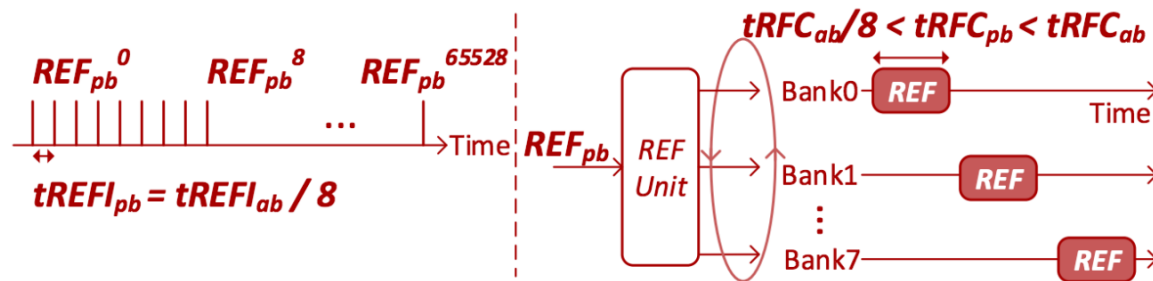
- CAS before RAS (CBR)
 - ▣ DRAM memory keeps track of the addresses using an internal counter
- RAS only refresh (ROR)
 - ▣ Row address is specified by the controller; similar to a pair of activate and precharge
- Auto-refresh vs. self refresh
 - ▣ Every 7.8us a REF command is sent to DRAM ($t_{RAS} + t_{RP}$)
 - ▣ LPDDR turns off IO for saving power while refreshing multiple rows

Refresh Granularity

- All bank vs. per bank refresh



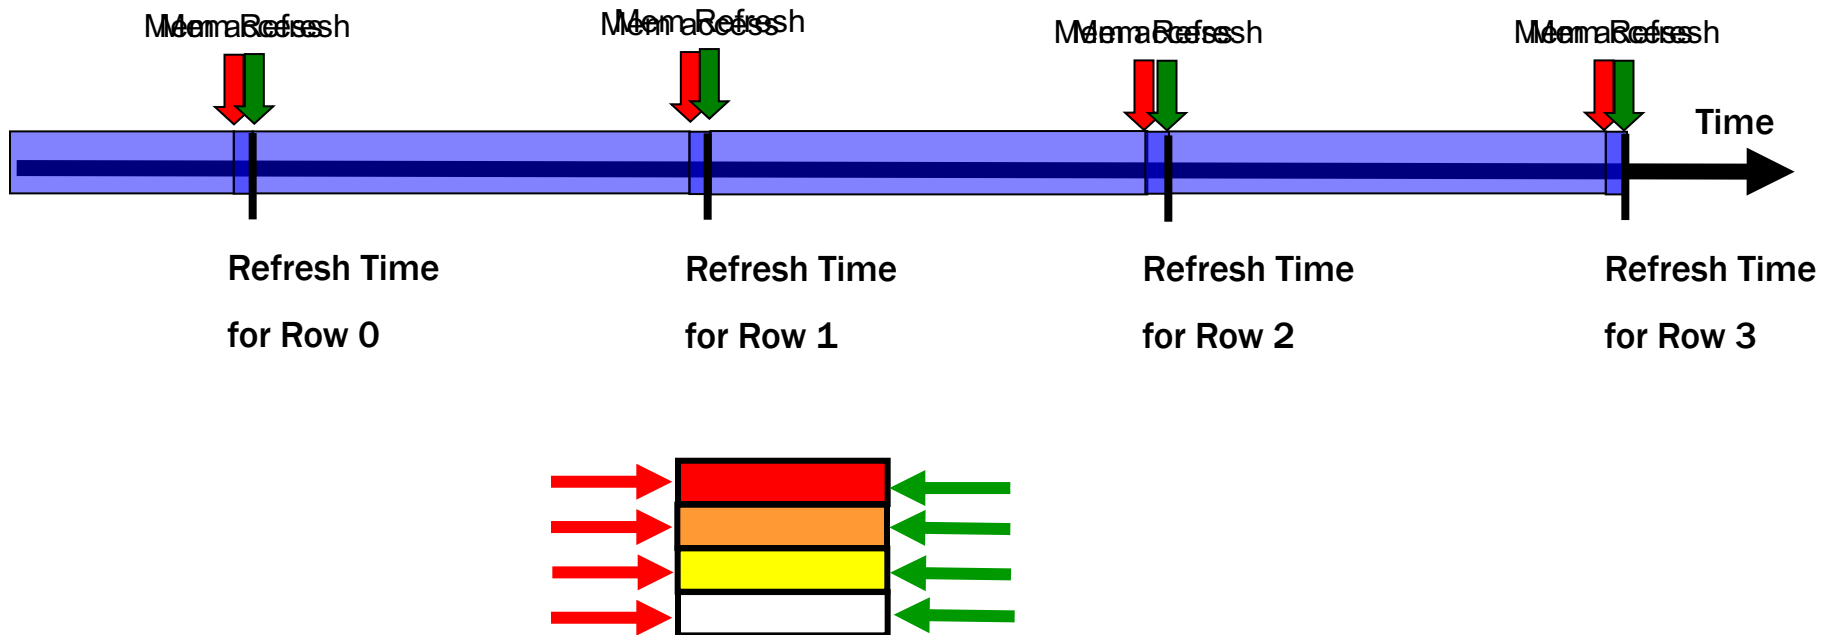
(a) All-bank refresh (REF_{ab}) frequency and granularity.



(b) Per-bank refresh (REF_{pb}) frequency and granularity.

Optimizing DRAM Refresh

- Observation: each row may be accessed as soon as it is to be refreshed



Smart Refresh

- Idea: avoid refreshing recently accessed rows

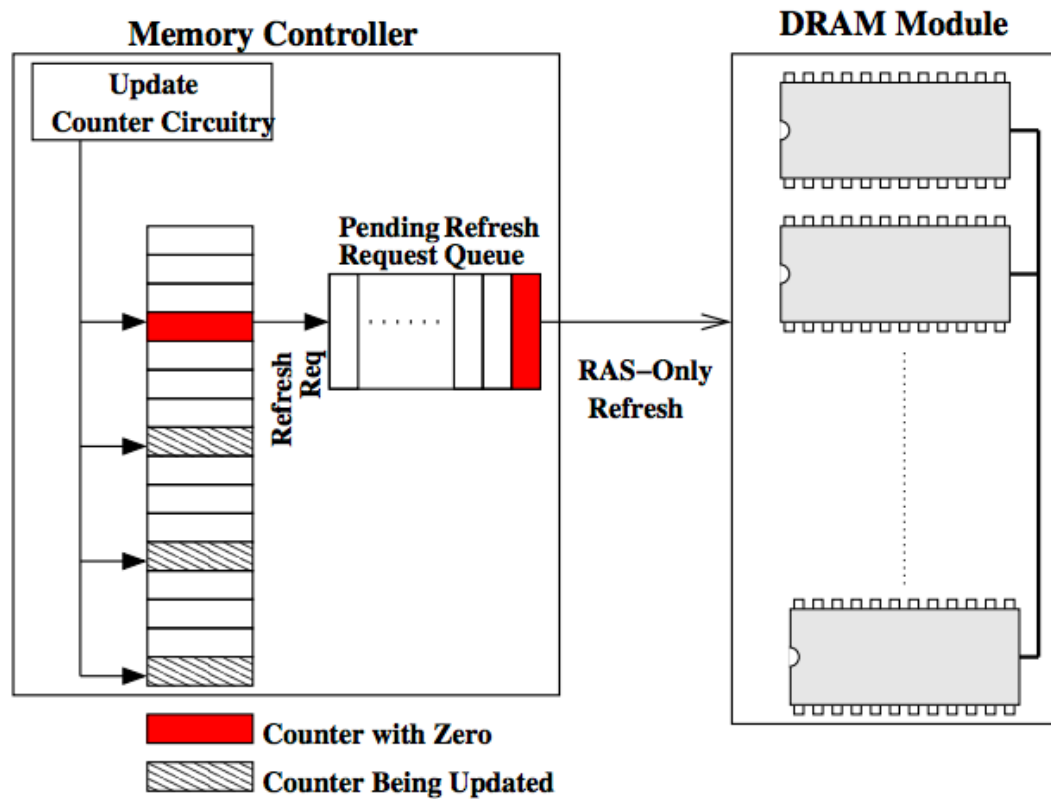




Figure 5: Smart Refresh Control Schematic

[Ghosh'07]

Diverse Impacts of Refresh



Worst Case Refresh Hit DRAM Read




DRAM capacity	tRFC	bandwidth overhead (95°C per Rank)	latency overhead (95°C)
512Mb	90ns	2.7%	1.4ns
1Gb	110ns	3.3%	2.1ns
2Gb	160ns	5.0%	4.9ns
4Gb	300ns	7.7%	11.5ns
8Gb	350ns	9.0%	15.7ns

Elastic Refresh

- Send refreshes during periods of inactivity
- Non-uniform request distribution
- Refresh overhead just has to fit in free cycles
- Initially not aggressive, converges with delay until empty (DUE) as refresh backlog grows
- Latency sensitive workloads are often lower bandwidth
- Decrease the probability of reads conflicting with refreshes

Elastic Refresh


- Introduce refresh backlog dependent idle threshold
- With a log backlog, there is no reason to send refresh command
- With a bursty request stream, the probability of a future request decreases with time
- As backlog grows, decrease this delay threshold



Key: to reduce REF and READ conflicts

[Stuecheli'10]


DRAM Refresh vs. ERROR Rate



If software is able to tolerate errors, we can lower DRAM refresh rates to achieve considerable power savings

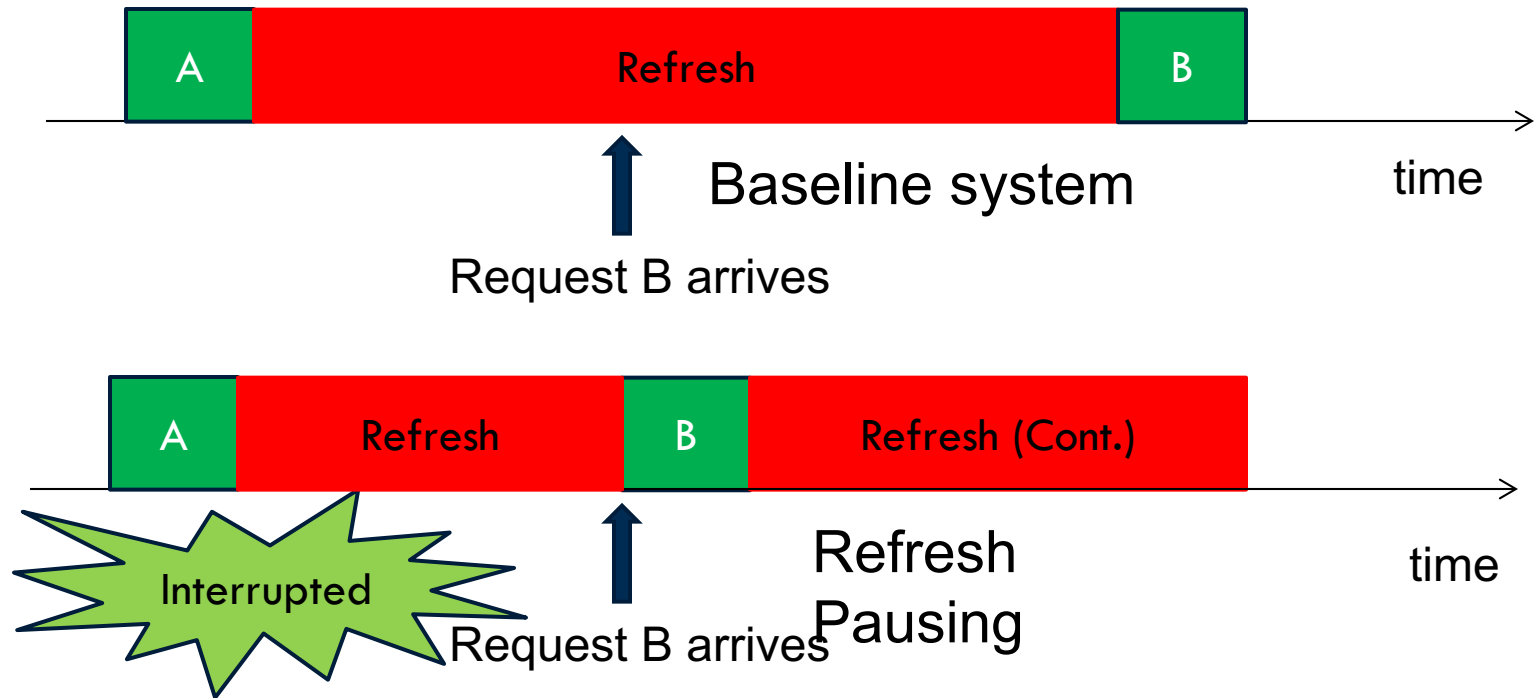
Flicker

- Divide memory bank into high refresh part and low refresh parts
- Size of high-refresh portion can be configured at runtime
- Small modification of the Partial Array Self-Refresh (PASR) mode



[Song'14]

Refresh Pausing



Pausing at arbitrary point can cause data loss

Pausing Refresh reduces wait time for Reads

Performance Results

