

RESISTIVE MEMORY TECHNOLOGY

Mahdi Nazm Bojnordi

Assistant Professor

School of Computing

University of Utah

Overview


- Upcoming deadlines
 - ▣ Tonight: homework assignment is due
 - ▣ March 27th: Sign up for your student paper presentation

- This lecture
 - ▣ Resistive memory technology
 - ▣ Write optimization techniques
 - ▣ Wear leveling
 - ▣ MLC technologies

Resistive Memory Technology

□ Main benefits


- ▣ Non-volatile memory
- ▣ Multi-level storage
- ▣ Denser cells
- ▣ Better scalability



□ Shortcomings

- ▣ Limited endurance
- ▣ High switching delay and energy

What can we do?



Comparison of Technologies

- Compared to NAND Flash, PCM is byte-addressable, has orders of magnitude lower latency and higher endurance.

	DRAM	PCM	NAND Flash
Page size	64B	64B	4KB
Page read latency	20-50ns	~ 50ns	~ 25 μ s
Page write latency	20-50ns	~ 1 μ s	~ 500 μ s
Write bandwidth	~GB/s per die	50-100 MB/s per die	5-40 MB/s per die
Erase latency	N/A	N/A	~ 2 ms
Endurance	∞	$10^6 - 10^8$	$10^4 - 10^5$
Read energy	0.8 J/GB	1 J/GB	1.5 J/GB [28]
Write energy	1.2 J/GB	6 J/GB	17.5 J/GB [28]
Idle power	~100 mW/GB	~1 mW/GB	1-10 mW/GB
Density	1×	2 - 4×	4×

Sources: [Doller '09] [Lee et al. '09] [Qureshi et al. '09]

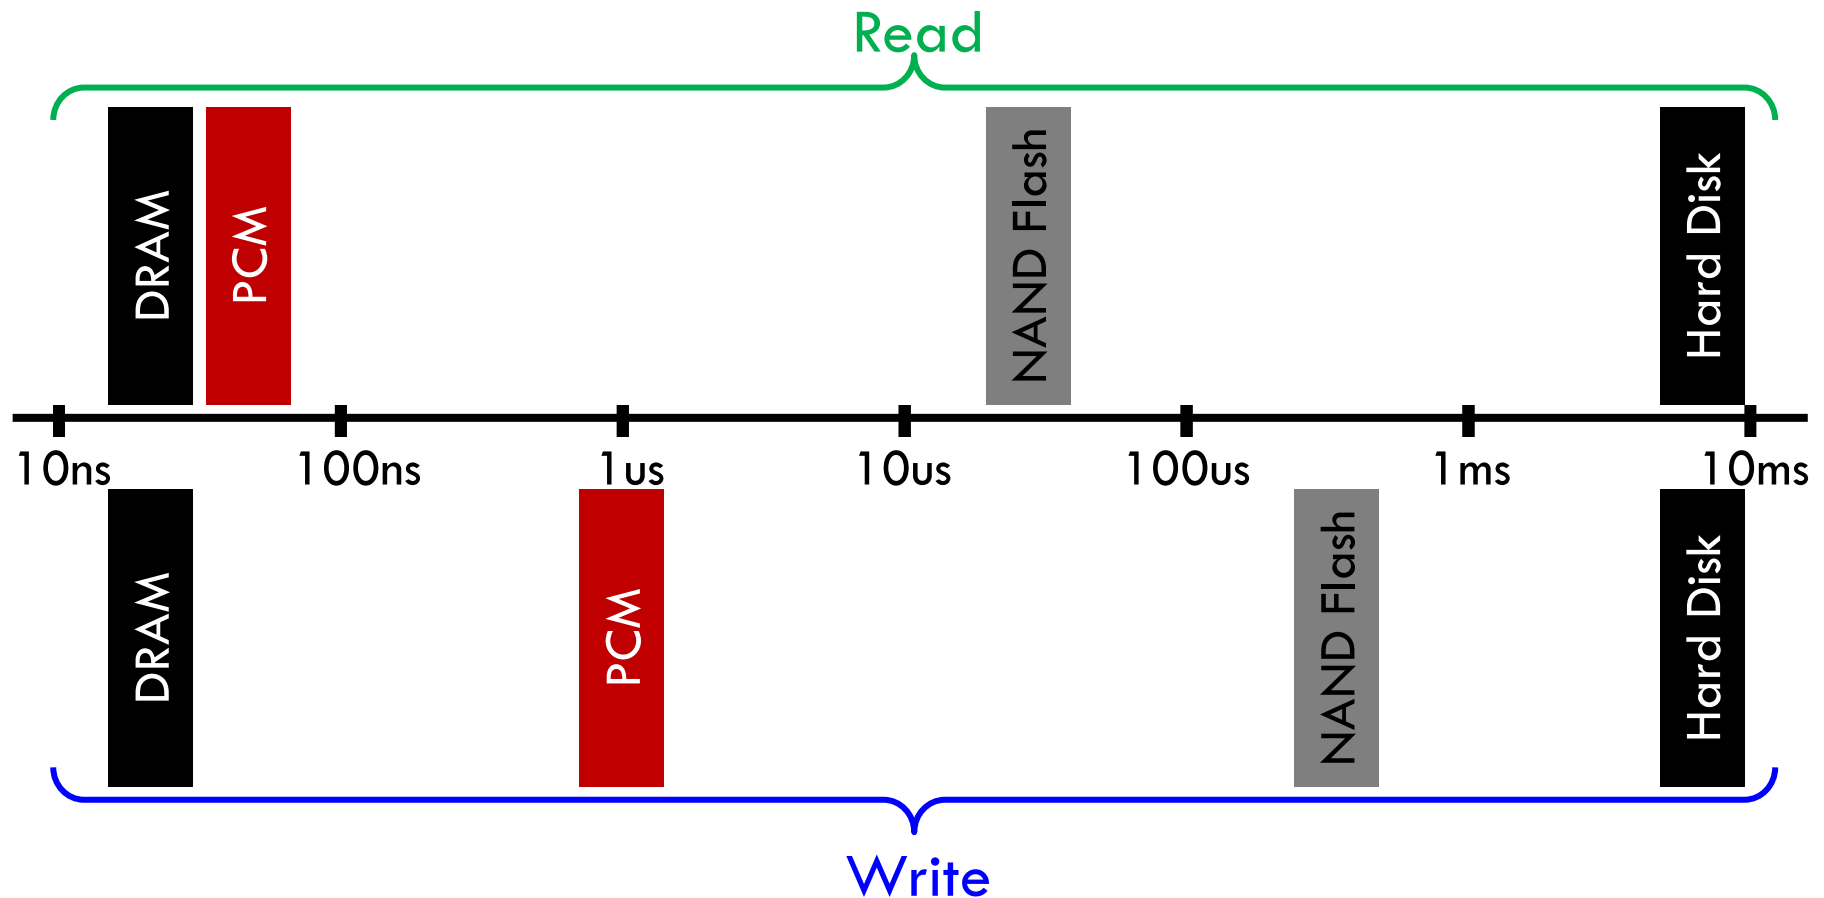
Comparison of Technologies

- Compared to DRAM, PCM has better density and scalability and similar read but longer write latencies

	DRAM	PCM	NAND Flash
Page size	64B	64B	4KB
Page read latency	20-50ns	~ 50ns	~ 25 μ s
Page write latency	20-50ns	~ 1 μ s	~ 500 μ s
Write bandwidth	~GB/s per die	50-100 MB/s per die	5-40 MB/s per die
Erase latency	N/A	N/A	~ 2 ms
Endurance	∞	$10^6 - 10^8$	$10^4 - 10^5$
Read energy	0.8 J/GB	1 J/GB	1.5 J/GB [28]
Write energy	1.2 J/GB	6 J/GB	17.5 J/GB [28]
Idle power	~100 mW/GB	~1 mW/GB	1-10 mW/GB
Density	1×	2 – 4×	4×

Sources: [Doller '09] [Lee et al. '09] [Qureshi et al. '09]

Latency Comparison



Read Compare Write

- A cache line is written in several cycles
- Read-compare-write (differential write)
 - Write only modified bits rather than entire cache line
- Skipping parts with no modified bits

Cache line

0	1	0	1	1	0	0	1	0	1	1	0	0	0	0	1
0	1	0	1	1	0	1	1	0	1	1	0	1	1	1	0

PCM

0	1	0	1	1	0	0	1	0	1	1	0	0	0	0	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

Reducing Bit Flips

- ❑ Encode write data into either its regular or inverted form and then pick the encoding that yields in less flips in comparison against old data.

Flip-N-Write [MICRO'09]

Old	0	0	1	0	1	1
New(Regular)	0	1	0	1	0	1
New (Inverted)	1	0	1	0	1	0

Saves 4 bit flips

- ❑ Encode write data into a set of data vectors and then pick the vector that yields in less flips in comparison against old data.


Flip-Min [HPCA'13]

Old	0	0	1	0	1	1
New ₁	0	1	0	1	0	1
New ₂	1	0	1	0	1	0
New ₃	1	0	1	0	1	1

Saves 5 bit flips

Limited Lifetime


Challenge : Each cell can endure **10-100 Million** writes



With uniform write traffic, system lifetime ranges from 4-20 years

Non-Uniform Writes


- Even with 64K spare lines, baseline gets 5% lifetime of ideal



Impact of Non-Uniformity

- Even with 64K spare lines, baseline gets 5% lifetime of ideal

$$\text{Norm. Endurance} = \frac{\text{Num. writes before system failure}}{\text{Num. writes before failure with uniform writes}} \times 100\%$$



[Qureshi'09]


Making Writes Uniform

- Wear Leveling: make writes uniform by remapping frequently written lines

Line Addr.	Lifetime Count	Period Count
A	99K (Low)	1K (Low)
B	100K (Med)	3K (High)
C	101K (High)	2K (Med)



Line	Remap Addr
A	C
B	A
C	B




How to Remap

- Tables
 - ▣ Area of several (tens of) megabytes
 - ▣ Indirection latency (table in EDRAM/DRAM)
- Area overhead can be reduced with more lines per region
 - ▣ Reduced effectiveness (e.g. Line0 always written)
 - ▣ Support for swapping large memory regions (complex)

Start-Gap Wear Leveling

- Two registers (Start & Gap) + 1 line (GapLine) to support movement
- Move GapLine every 100 writes to memory.



$\text{PCMAAddr} = (\text{Start} + \text{Addr}); (\text{PCMAAddr} \geq \text{Gap}) \text{PCMAAddr}++$


Storage overhead: less than 8 bytes (GapLine taken from spares)

Latency: Two additions (no table lookup)


Write overhead: One extra write every 100 writes → 1%

Start-Gap Results


- On average, Start-Gap gets 53% normalized endurance




Multi-Level Cells



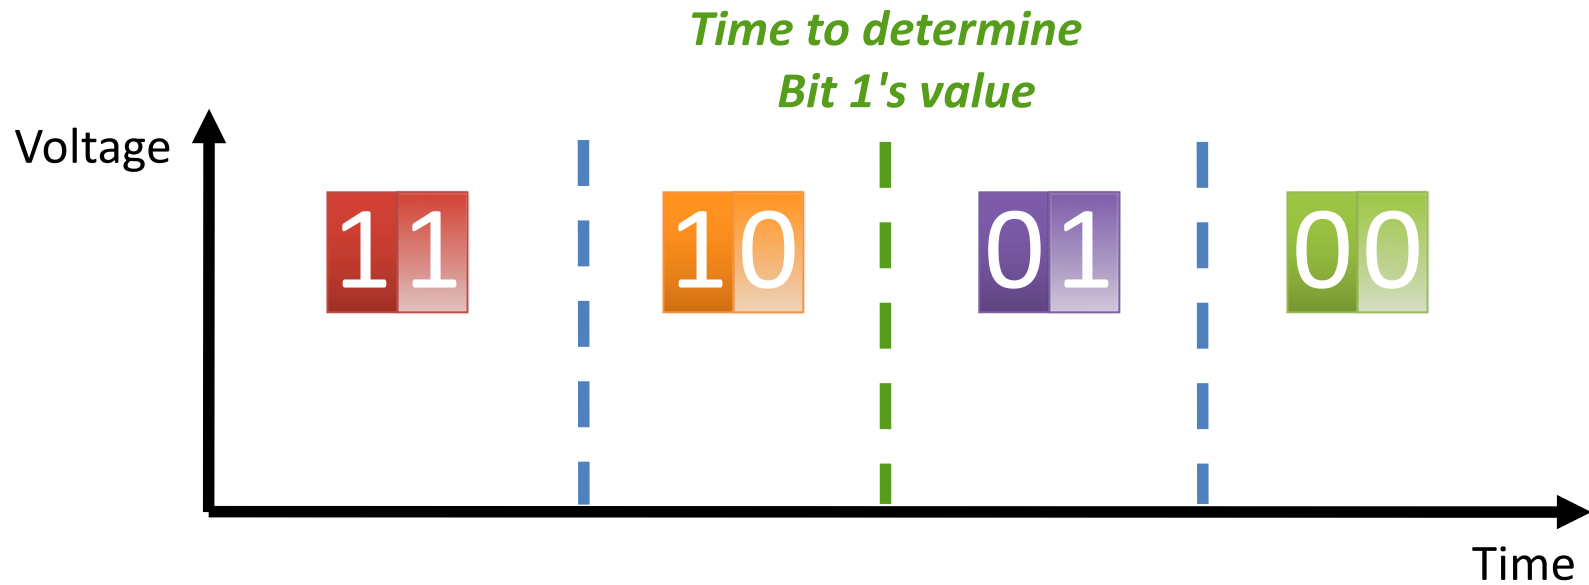
Sensing Multi-level Cells




Multi-Level Cells



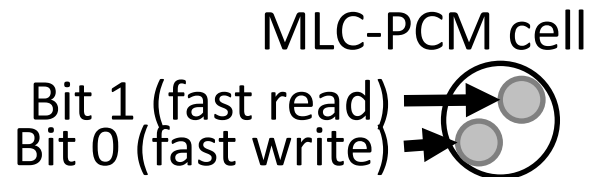
Multi-Level Cells



Multi-Level Cells



Decoupled Bit Mapping



Coupled (baseline): Contiguous ***bits*** alternate between FR and FW




Decoupled: Contiguous ***regions*** alternate between FR and FW



Decoupled Bit Mapping

- By decoupling, we've created regions with distinct characteristics
 - We examine the use of 4KB regions (e.g., OS page size)



- *Fast read page* *Fast write page*
- Want to match *frequently read data to FR pages* and vice versa
- Toward this end, we propose a new **OS page allocation** scheme

Performance Results

